



数据库系统原理

Database System Principle

邵莹侠

Email: shaoyx@bupt.edu.cn

北京邮电大学计算机学院

计算机应用技术中心



PART 2

RELATIONAL DATABASES



Chapter 5

Advanced SQL



Three Parts in Chapter 5

- Accessing SQL From a Programming Language
 - dynamic SQL, e.g. JDBC, ODBC, ADO, ...
 - Embedded SQL
- Functions and Procedural (过程) Constructs
- Triggers
- (Recursive Queries, Advanced Aggregation Features, OLAP)

5.1 Accessing SQL From a Programming Language

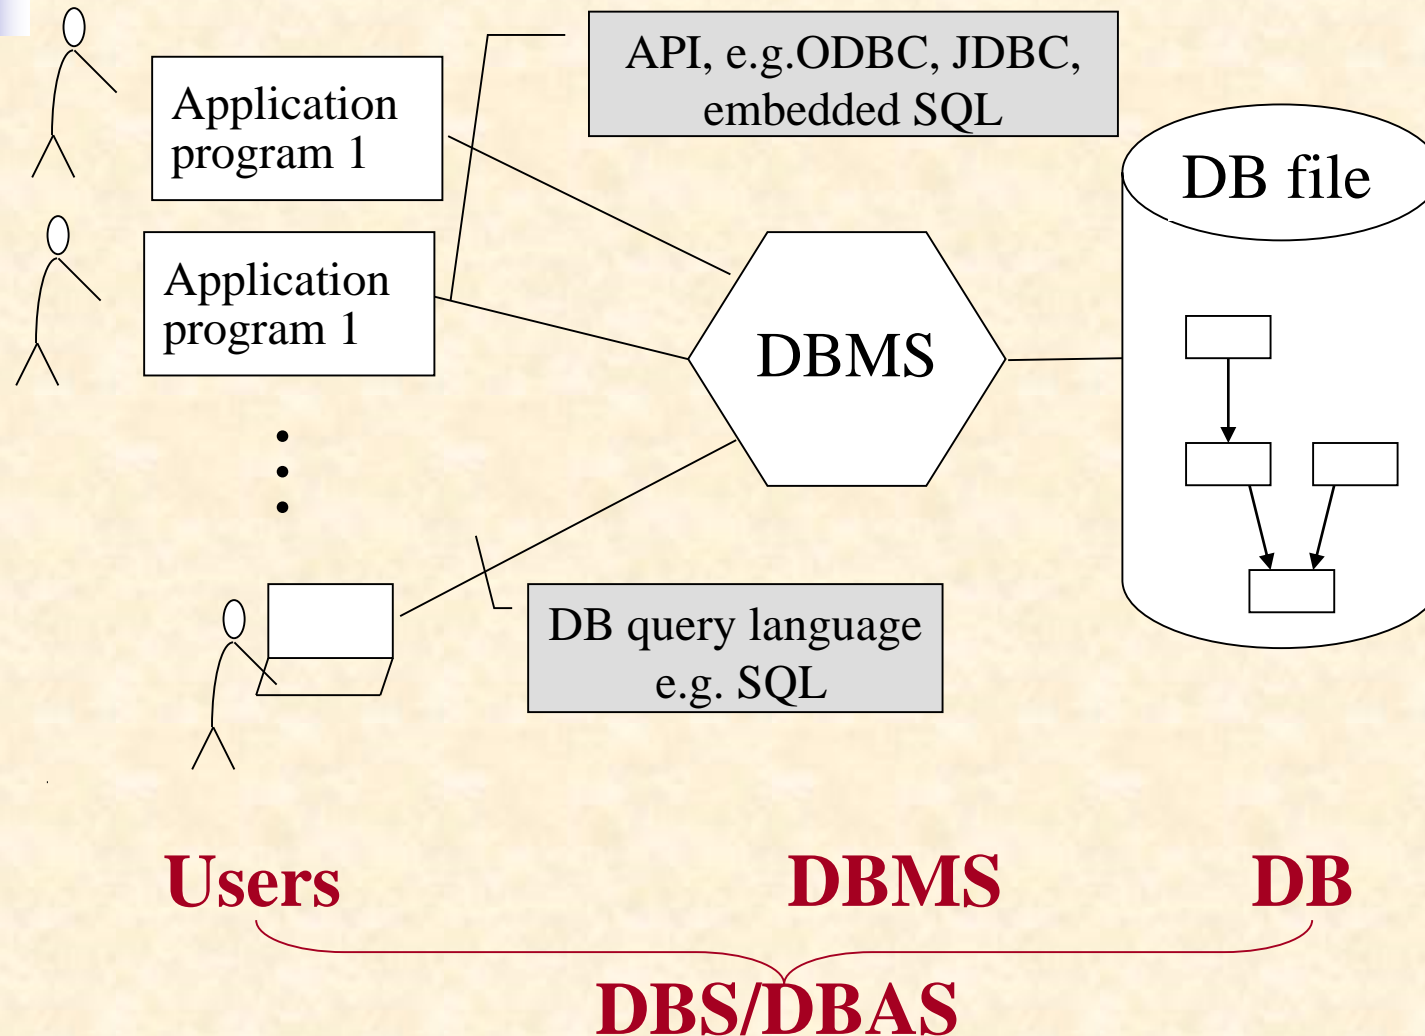


Fig.5.0.1 DBS and DBAS

Accessing SQL From a Programming Language

- API (application-program interface) for a program to interact with a database server
- Application conducts complex data processing, and makes calls to
 - Connect with the database server
 - Send SQL commands to the database server
 - Fetch tuples of result one-by-one into program variables
- Various tools:
 - JDBC (Java Database Connectivity) works with Java
 - ODBC (Open Database Connectivity) works with C, C++, C#, and Visual Basic
 - Other API's such as ADO.NET sit on top of ODBC
 - Embedded SQL

5.1.1 JDBC

- **JDBC** is a Java API for communicating with database systems supporting SQL.
- JDBC supports a variety of features for querying and updating data, and for retrieving query results.
- JDBC also supports metadata retrieval, such as querying about relations present in the database and the names and types of relation attributes.
- Model for communicating with the database:
 - Open a connection
 - Create a “SQL statement” object
 - Execute queries using the Statement object to send queries and fetch results
 - Exception mechanism to handle errors


```
public static void JDBCexample(String userid, String passwd)
{
    try
    {
        Class.forName ("oracle.jdbc.driver.OracleDriver");
        Connection conn = DriverManager.getConnection(
            "jdbc:oracle:thin:@db.yale.edu:1521:univdb",
            userid, passwd);
        Statement stmt = conn.createStatement();
        try {
            stmt.executeUpdate(
                "insert into instructor values('77987', 'Kim', 'Physics', 98000)");
        } catch (SQLException sqle)
        {
            System.out.println("Could not insert tuple. " + sqle);
        }
        ResultSet rset = stmt.executeQuery(
            "select dept_name, avg (salary) "+
            " from instructor "+
            " group by dept_name");
        while (rset.next()) {
            System.out.println(rset.getString("dept_name") + " " +
                rset.getFloat(2));
        }
        stmt.close();
        conn.close();
    }
    catch (Exception sqle)
    {
        System.out.println("Exception : " + sqle);
    }
}
```

Fig.5.1 A JDBC Example

JDBC

- Open a connection, named as *conn*
 - using the *getConnection* method of the *DriverManager* class
 - parameters
 - URL/machine to be connected;
 - user identifier, password
- Create a SQL “statement” object on the connection *conn*
 - using the *createStatement* method
 - its statement handle is named as *stmt*
 - allows the Java program to invoke (调用) methods that ship an SQL statement given as an argument for execution by the DBMS

JDBC

- Execute queries using the Statement object *stmt* to send queries and fetch results
- using *execute.query* or *execute.update* such as *insert/delete/update/createtable*
- parameters
 - the SQL statement to be executed, represented as a string
- using the *try {...}/catch {...}* construct to catch any exceptions or error conditions



JDBC

- Fetch the query result
 - retrieve the set of tuples in the result into a ***ResultSet*** object *rset*, and fetch them one tuple at a time
 - the *next()* method tests whether or not the result set has at least one tuple and if so, fetches it
- At the end of the query, the statement *stmt* and the connection *conn* are closed



5.1.2 ODBC

- Open DataBase Connectivity (ODBC) standard
 - standard for application program to communicate with a database server.
 - application program interface (API) to
 - open a connection with a database,
 - send queries and updates,
 - get back results.
- Applications such as GUI, spreadsheets, etc. can use ODBC
- More details can be found in Appendix A.

§ 5.4 Embedded SQL

- Approaches to take SQL as DB query tools
 - interactive SQL
 - SQL is used directly as DML and DDL through DBS human-machine interfaces
 - dynamic SQL, e.g JDBC, ODBC
 - embedded SQL
 - SQL is embedded in general-purpose programming languages, e.g. *C* language
 - executing of general-purpose programming language programs with SQL statement embedded results in DB access

Embedded SQL

- The SQL standard defines embeddings of SQL in a variety of programming languages such as C, C++, Pascal, Fortran, and Cobol
- A language to which SQL queries are embedded is referred to as a *host language*(宿主语言), and the SQL structures permitted in the host language comprise *embedded SQL*
- Merits of embedded SQL
 - 交互式SQL只能进行DB的访问操作，不能对DB访问结果进行进一步的数据处理
 - Embedded SQL将SQL的数据库访问功能与C语言等宿主语言的数据处理能力相结合，提高了数据应用系统的能力

Embedded SQL

- **EXEC SQL** statement is used to identify embedded SQL request to the preprocessor

EXEC SQL <embedded SQL statement >;

- Note: this varies by language:
 - In some languages, like COBOL, the semicolon(分号) is replaced with END-EXEC
 - In Java embedding uses # SQL { };

Embedded SQL (Cont.)

- Before executing any SQL statements, the program must first connect to the database. This is done using:

EXEC-SQL connect to *server* user *user-name* using *password*;

Here, *server* identifies the server to which a connection is to be established.

Embedded SQL (Cont.)

- Variables of the host language can be used within embedded SQL statements. They are preceded by a colon (:) to distinguish from SQL variables
 - e.g., *:credit_amount*
- Variables used as above must be declared within DECLARE section, as illustrated below. The syntax (语法) for declaring the variables, however, follows the usual host language syntax.

EXEC-SQL BEGIN DECLARE SECTION;

int *credit-amount* ;

EXEC-SQL END DECLARE SECTION;



Embedded SQL (Cont.)

- E.g.1 From within a host language, find the ID and name of students who have completed more than the number of credits stored in variable *credit_amount* in the host language
 - *credit_amount* is the shared variable defined in the declaration part

C program

EXEC SQL BEGIN DECLARE SECTION

int *credit_amount*;

EXEC SQL END DECLARE SECTION

credit_amount := e.g. input-from-screen-by-users

EXEC SQL

select *ID, name*

from *student*

where tot_cred > *:credit_amount*

END-EXEC

shared variable defined
in host language

Cursor (游标) in Embedded SQL

- To write an embedded SQL query, we use the
declare *c* cursor for <SQL query>
statement.

The variable *c* is used to identify the query

EXEC SQL

```
declare c cursor for  
select ID, name  
from student  
where tot_cred > :credit_amount
```

END_EXEC

why using Cursor in Embedded SQL

- *利用Embedded SQL进行查询时，查询结果有可能包括多个元组，此时无法直接将多个元组通过共享变量赋值传递给宿主程序
- *系统开辟专门working区域存放SQL查询的结果关系（也称作中间关系），并利用查询游标c指向此区域。宿主程序根据c指向的查询结果关系集合，使用**open, fetch, close**依次获取结果关系中的各元组

Example 2

EXEC SQL BEGIN DECLARE SECTION

int *credit_amount*;

char *si*, *sn*;

EXEC SQL END DECLARE SECTION

credit_amount := input-from-screen-by-users

EXEC SQL

declare *c* cursor for

select *ID*, *name*

from *student*

where tot_cred > *:credit_amount*

END-EXEC

*EXEC SQL open *c* END-EXEC*

*EXEC SQL fetch *c* into *:si*, *:sn* END-EXEC*

*EXEC SQL close *c* END-EXEC*

可以用循环语句依次
取走全部查询结果

shared variable

Cursor in Embedded SQL (cont.)

- Usage of cursor in embedded SQL

declare cursor – open – fetch - close

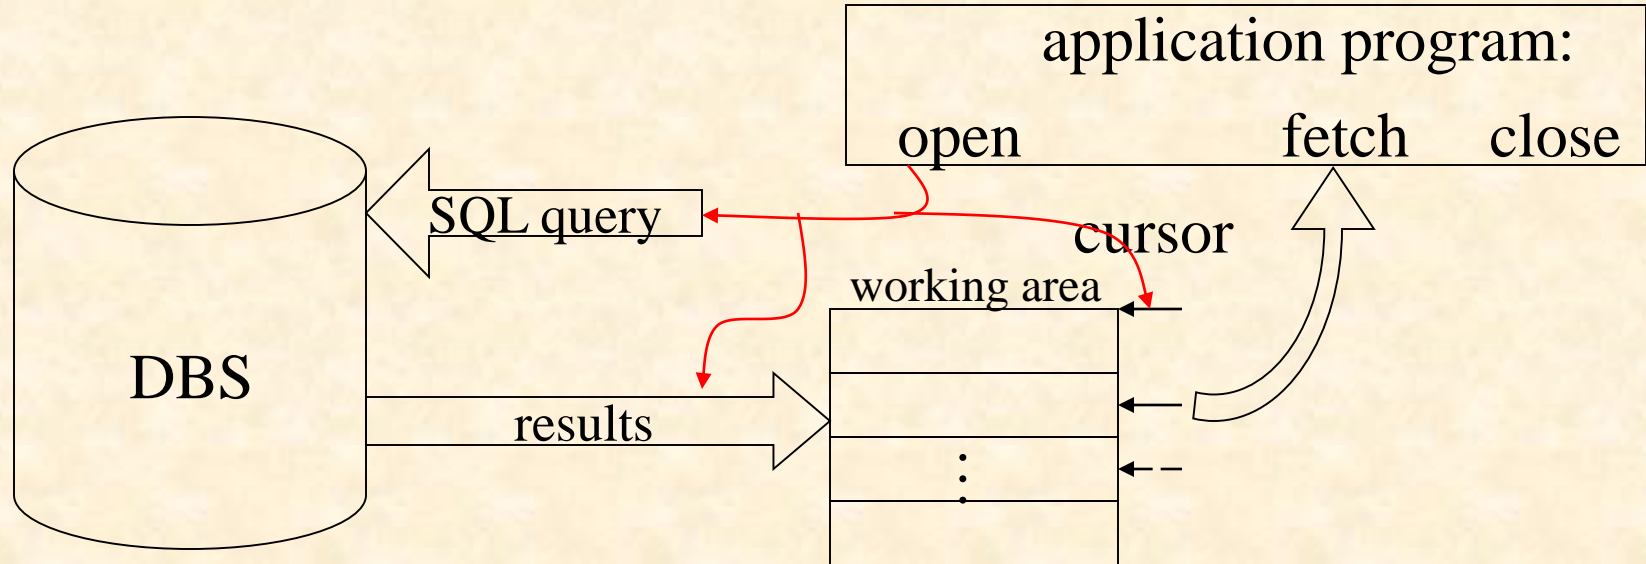


Fig. 5.0.3 Cursor in Embedded SQL

Cursor in Embedded SQL (cont.)

- The **open** statement for our example is as follows:

EXEC SQL open c ;

This statement causes the database system to execute the query and to save the results within a temporary relation. The query uses the value of the host-language variable *credit-amount* at the time the **open** statement is executed.

- The **fetch** statement causes the values of one tuple in the query result to be placed on host language variables.

EXEC SQL fetch c into :si, :sn END_EXEC

Repeated calls to fetch get successive tuples in the query result

Cursor in Embedded SQL (cont.)

- A variable called SQLSTATE in the SQL communication area (SQLCA) gets set to '02000' to indicate no more data is available
- The **close** statement causes the database system to delete the temporary relation that holds the result of the query.

EXEC SQL close *c* ;

- Note
 - above details vary with language
 - for example, the Java embedding defines Java iterators to step through result tuples.

SQL Procedures

- The *dept_count* function could instead be written as procedure:

```
create procedure dept_count_proc (in dept_name varchar(20),  
                                     out d_count integer)  
  
begin  
    select count(*) into d_count  
    from instructor  
    where instructor.dept_name = dept_count_proc.dept_name  
end
```

SQL Procedures

- Procedures can be invoked either from an SQL procedure or from embedded SQL, using the **call** statement.

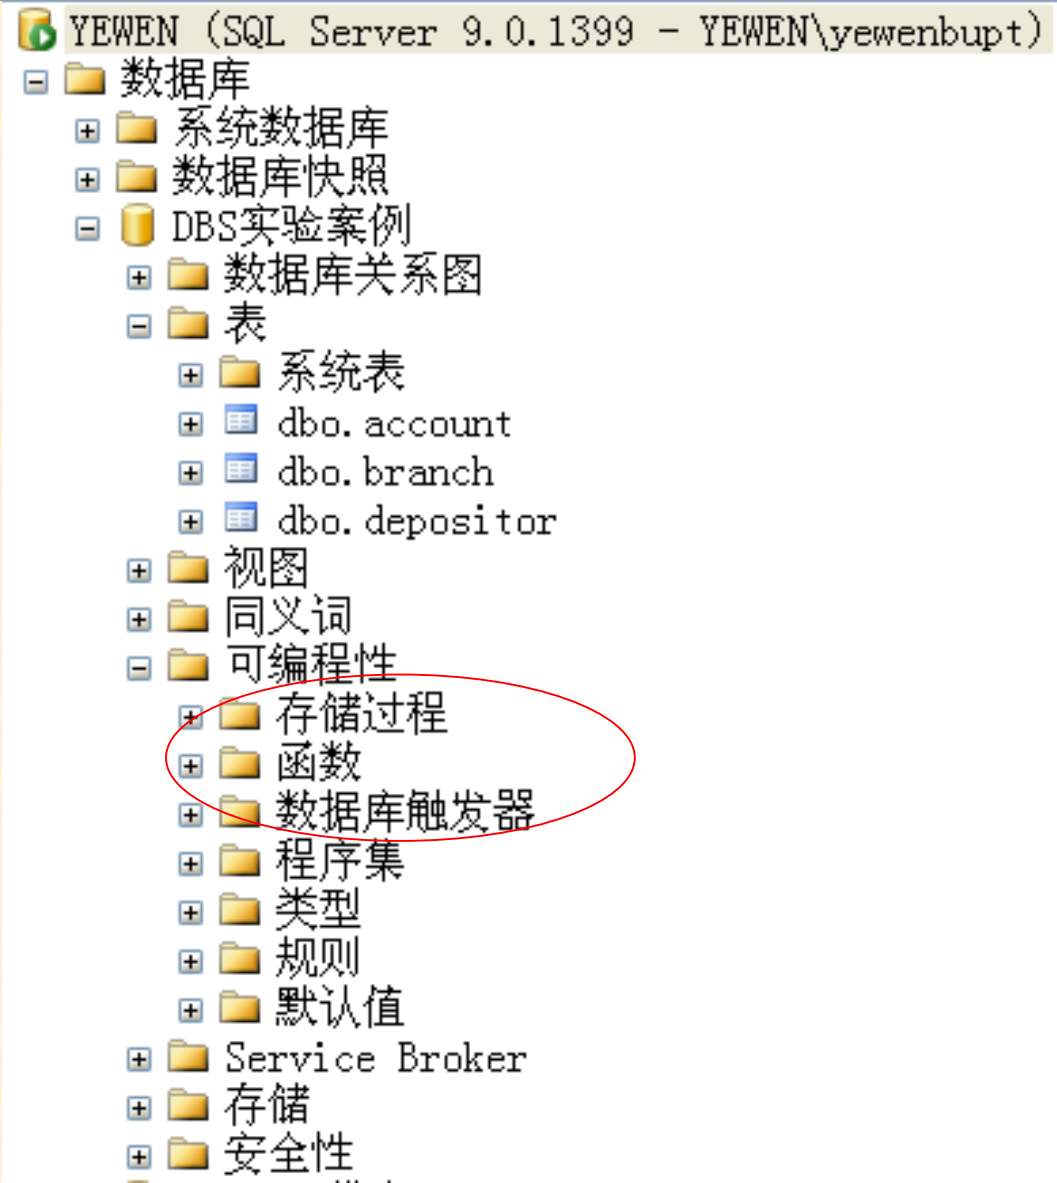
```
declare d_count integer;  
call dept_count_proc( 'Physics', d_count);
```

Procedures and functions can be invoked also from dynamic SQL

函数、存储过程预先生成对应的查询执行计划（类似于目标代码），存储在**DBMS**中，应用程序直接调用，不需再进行查询处理和优化（**i.e.**编译）

Functions and Procedures (cont.)

- The functions and procedures defined are stored in DBS, and can be called by SQL statements or programs



§ 5.3 Trigger (触发器)

- **Trigger**

- a statement that is executed automatically by DBMS as a side effect of a modification to the database
- As an **integrity control mechanism**, trigger is introduced to SQL standard in SQL:1999, but supported even earlier using non-standard syntax by most databases.
 - Syntax illustrated here may not work exactly on your database system;
 - check the system manuals

5.3 Trigger (cont.)

- Trigger is an *event-condition-action model* based integrity definition, checking, remedy (更正) mechanism
 - specify what **events** cause the trigger to be executed (e.g. **insert**, **delete**, **update**), and under which **conditions** the trigger execution will proceed
 - integrity constraints checking
 - specify the **actions** to be taken when the trigger executes
 - if constraints is violated, remedy actions are taken



Triggering Events and Actions in SQL

- Triggers on update can be restricted to specific attributes
 - for example, **after update of** *takes on grade*
- Values of attributes before and after an update can be referenced
 - **referencing old row as** : for deletes and updates
 - **referencing new row as** : for inserts and updates

Triggering Events and Actions in SQL

- Triggers can be activated before an event, which can serve as extra constraints.
- For example, convert blank grades to null.

```
create trigger setnull_trigger before update of takes  
referencing new row as nrow  
for each row  
when (nrow.grade = ' ')  
begin atomic  
    set nrow.grade = null;  
end;
```

优点：
预先编译（查询处理
优化），存储在
DBMS中，直接调用

Example: Trigger to Maintain `credits_earned` value

```
■ create trigger credits_earned after update of takes on (grade)  
  referencing new row as nrow  
  referencing old row as orow  
  for each row  
  when nrow.grade <> 'F' and nrow.grade is not null  
    and (orow.grade = 'F' or orow.grade is null)  
  begin atomic  
    update student  
    set tot_cred= tot_cred +  
      (select credits  
       from course  
       where course.course_id= nrow.course_id)  
    where student.id = nrow.id;  
  end;
```



Conclusion

- Accessing SQL From a Programming Language
 - dynamic SQL, e.g. JDBC, ODBC, ADO, ...
 - Embedded SQL
- Procedure
- Triggers



Appendix A: ODBC

- Architecture of ODBC
 - refer to Fig. 4.0.6
 - ODBC利用 ***driver*** 来适应不同厂家数据库的**物理结构**，如分块大小、访问方式

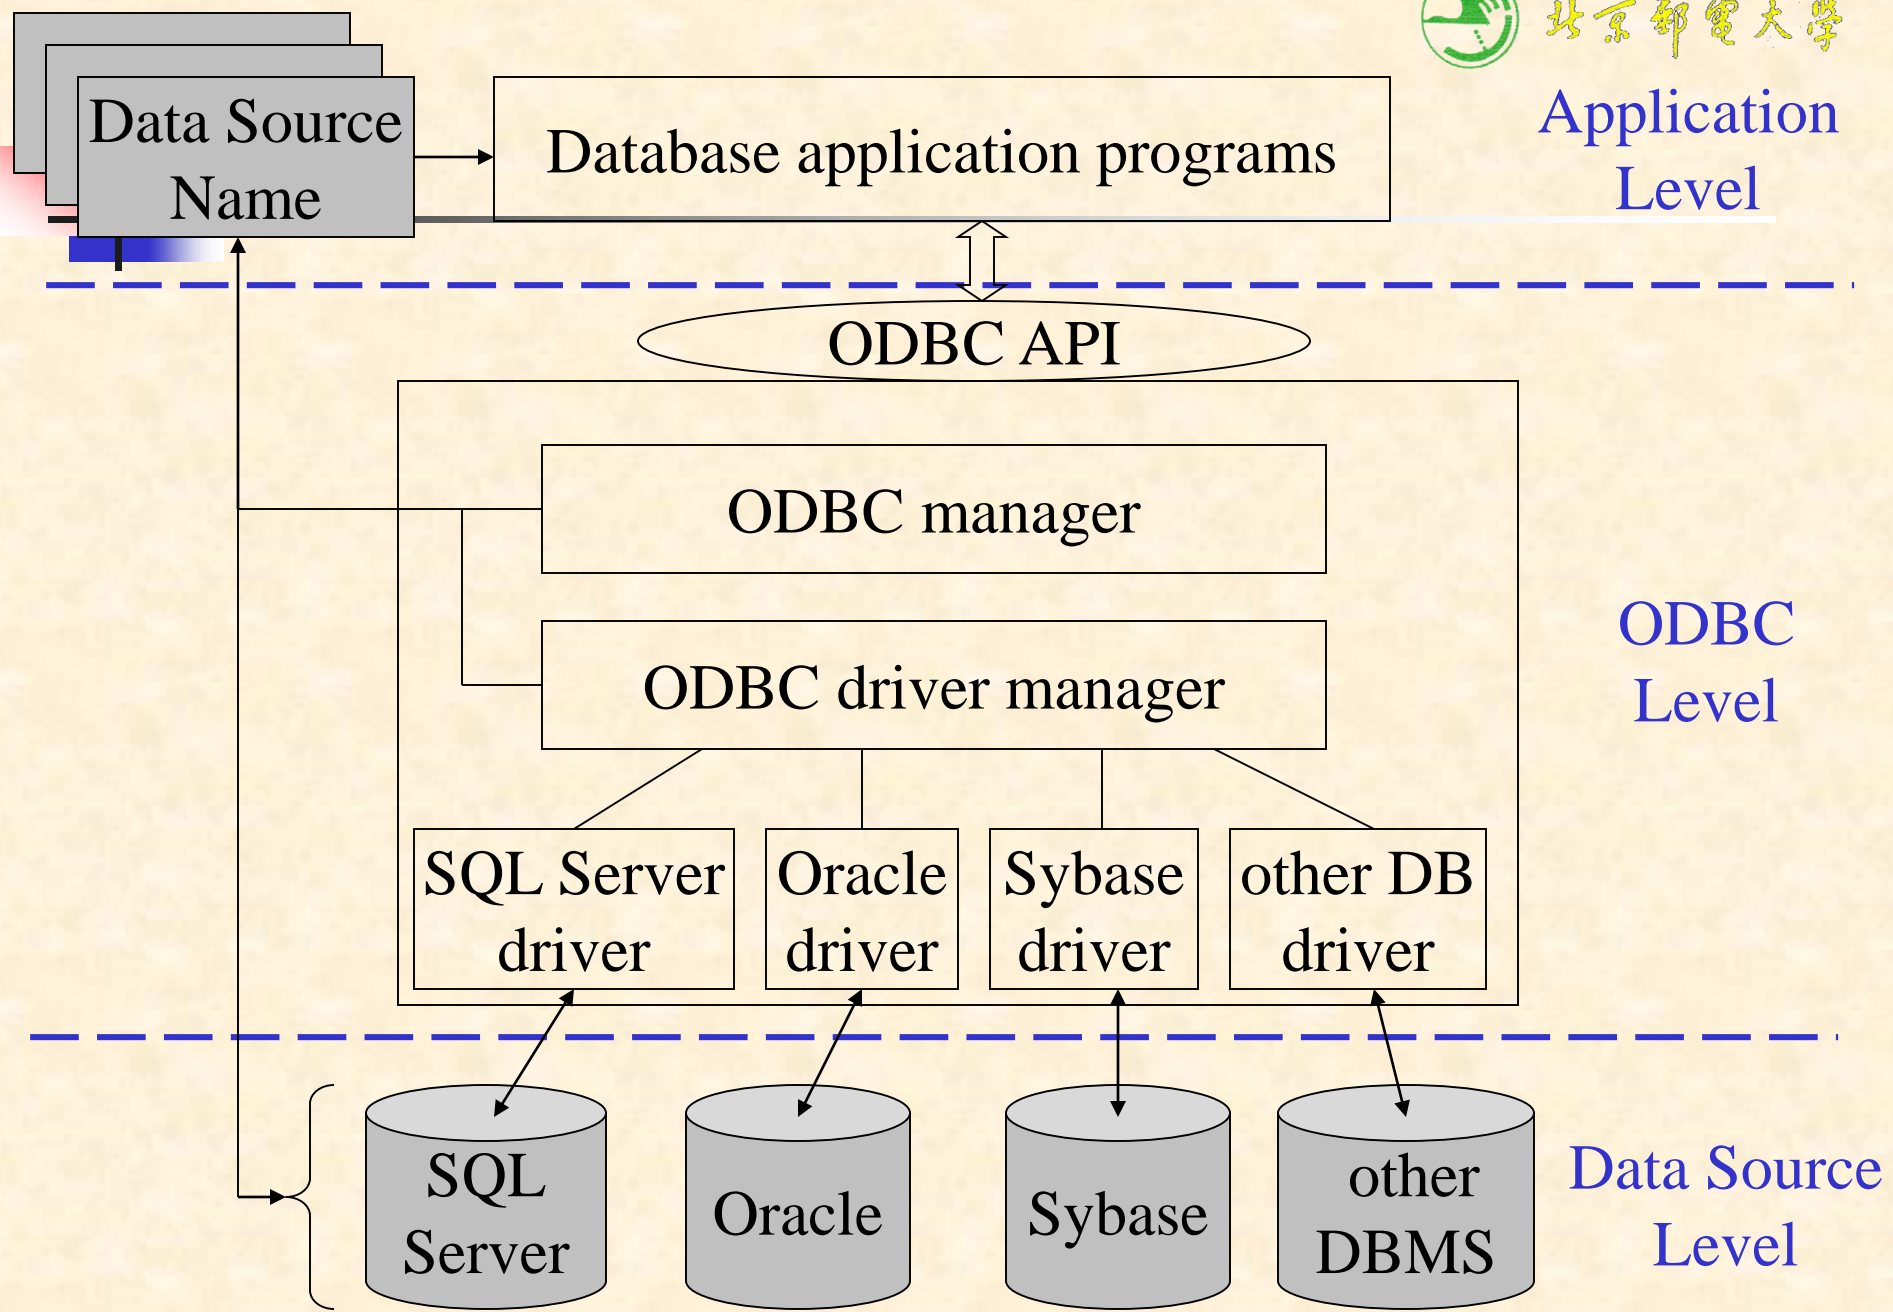


Fig.5.0.2 ODBC Architecture


```
void ODBCexample()
```

```
{
    RETCODE error;
    HENV env; /* environment */
    HDBC conn; /* database connection */

    SQLAllocEnv(&env);
    SQLAllocConnect(env, &conn);
    SQLConnect(conn, "db.yale.edu", SQL_NTS, "avi", SQL_NTS,
               "avipasswd", SQL_NTS);
    {
        char deptname[80];
        float salary;
        int lenOut1, lenOut2;
        HSTMT stmt;

        char * sqlquery = "select dept_name, sum (salary)
                           from instructor
                           group by dept_name";
        SQLAllocStmt(conn, &stmt);
        error = SQLExecDirect(stmt, sqlquery, SQL_NTS);
        if (error == SQL_SUCCESS) {
            SQLBindCol(stmt, 1, SQL_C_CHAR, deptname, 80, &lenOut1);
            SQLBindCol(stmt, 2, SQL_C_FLOAT, &salary, 0, &lenOut2);
            while (SQLFetch(stmt) == SQL_SUCCESS) {
                printf (" %s %g\n", deptname, salary);
            }
        }
        SQLFreeStmt(stmt, SQL_DROP);
    }
    SQLDisconnect(conn);
    SQLFreeConnect(conn);
    SQLFreeEnv(env);
}
```

Appendix B: Updates Through Embedded SQL

Embedded SQL expressions for database modification
(**update**, **insert**, and **delete**)

Can update tuples fetched by cursor by declaring that the cursor is for update

EXEC SQL

```
declare c cursor for  
select *  
from instructor  
where dept_name = 'Music'  
for update
```

Appendix B: Updates Through Embedded SQL

We then iterate through the tuples by performing **fetch** operations on the cursor (as illustrated earlier), and after fetching each tuple we execute the following code:

```
update instructor  
set salary = salary + 1000  
where current of c
```