



数据库系统原理

Database System Principle

邵莹侠

Email: shaoyx@bupt.edu.cn

北京邮电大学计算机学院

计算机应用技术中心



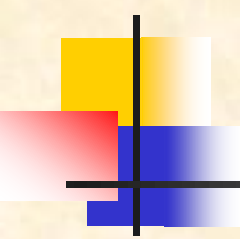
Part Two

Database Design



Chapter 6

Database Design and E-R Model

- 
-
- Design Process
 - Modeling
 - Constraints
 - E-R Diagram
 - Design Issues
 - Weak Entity Sets
 - Extended E-R Features
 - Design of the Bank Database
 - Reduction to Relation Schemas
 - Database Design
 - UML

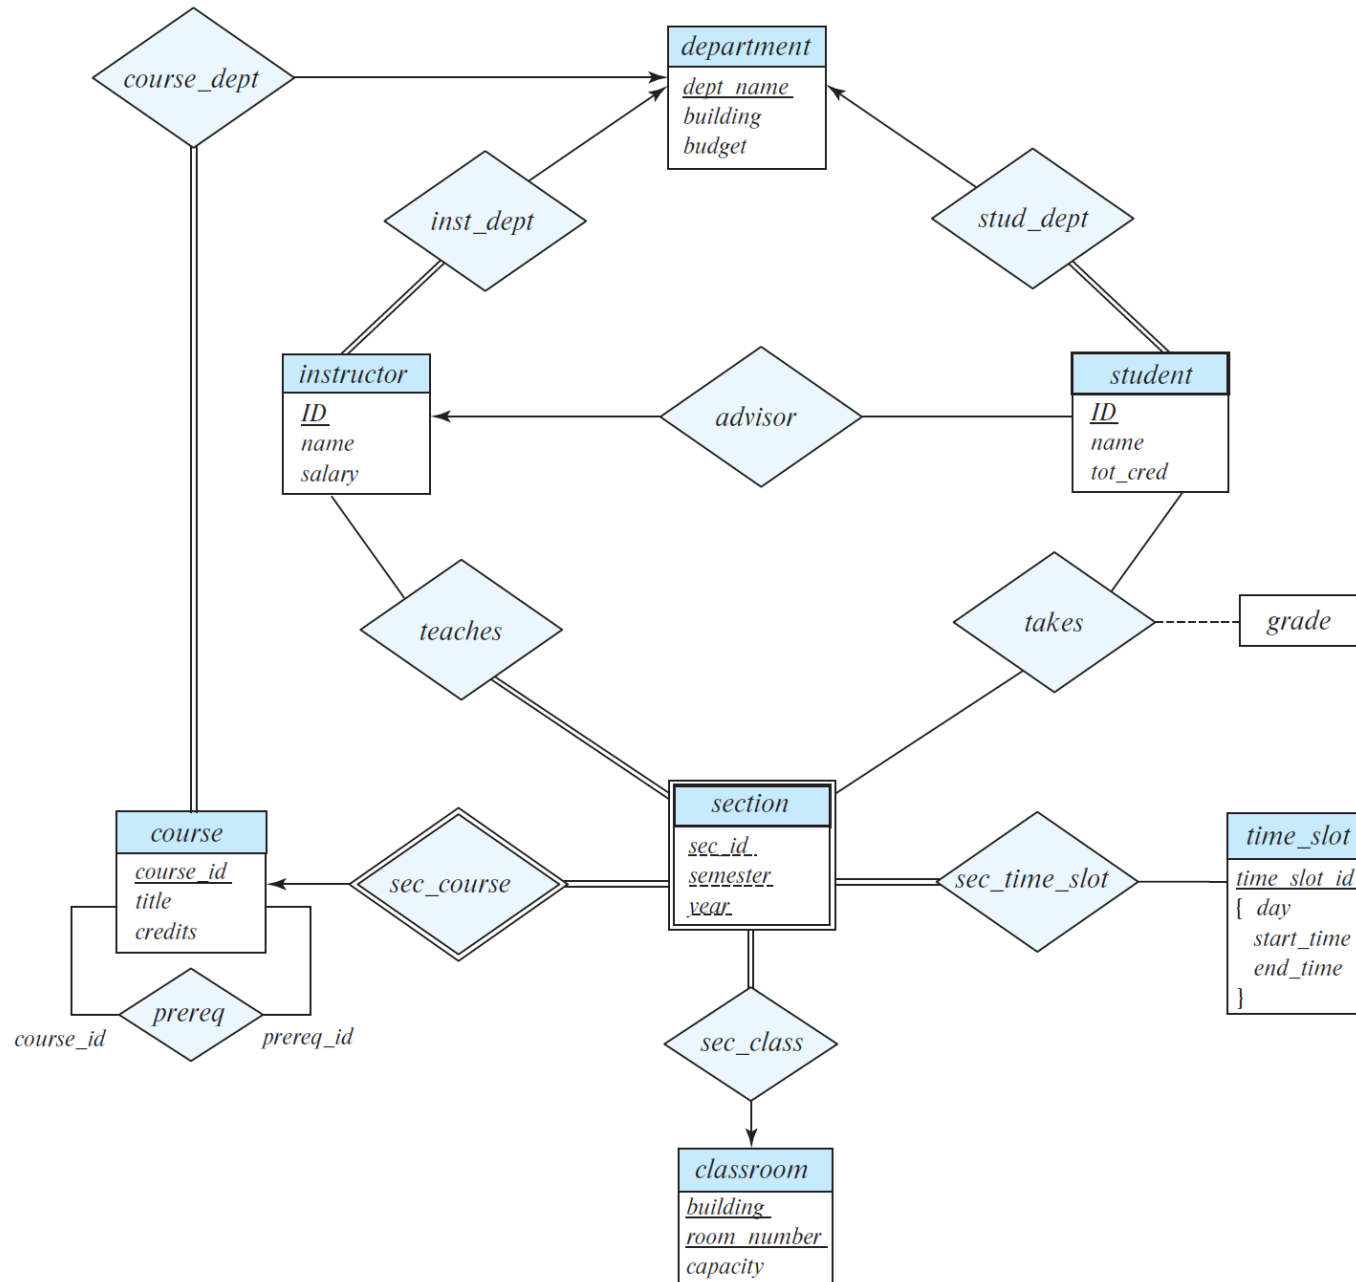
Parts in Chapter 6

- Part 1. DB/DBS/DBAS design process (§ 6.1/6.10 + Supplementary)
 - DB design phases (§ 6.1)
 requirement analysis, conceptual design, logical design,
 physical design
 - DBAS life-cycle model (生命周期模型)
- Part 2. The Entity-Relationship Model
 - *basic* E-R model
 - modeling elements (§ 6.2.1-6.2.3): entity sets, relationship sets, attributes
 - constraints (§ 6.3.1-6.3.3): mapping cardinality, participation constraint, keys
 - weak entity sets(§ 6.5.3)
 - removing redundant attributes in entity sets (§ 6.6)
 - E-R diagram(§ 6.2) and alternative notations(§ 6.10)

Parts in Chapter 6

- *extended* E-R Features (§ 6.8)
 - (§ 6.8.1-6.8.4) OO features in E-R model, i.e. specialization, generalization, attributes inheritance, constraints on generalization
 - (§ 6.8.5) aggregation: relationship among relationships
- Part 3. Reduction to Relational Schemas(§ 6.6, 6.8.6)
 - mapping elements in E-R model to that in relational models, i.e. conceptual schema → initial logical schema
- Part 4. E-R design issues (§ 6.9)
 - when applying E-R model to model the objects in real worlds, some issues (considerations and principles) should be addressed, to guarantee DBS effectiveness and efficiency for data

E-R Diagram for a University Enterprise



§ 6.1 Overview of the Design Process

- Database design consists of two sequential phases
 - analyzing of user requirements
 - *what data* should be stored in the database
 - what operations/transaction, such as *insert*, *delete*, *update* and *retrieve* are needed to conducted on these data
 - designing of DB schemas , in accordance with the *three-level of data abstract* (refer to Fig. 1.1)
 - conceptual design
 - logical design, at the logical level and view level
 - physical design, at the physical level
- refer to Fig.1.0.1

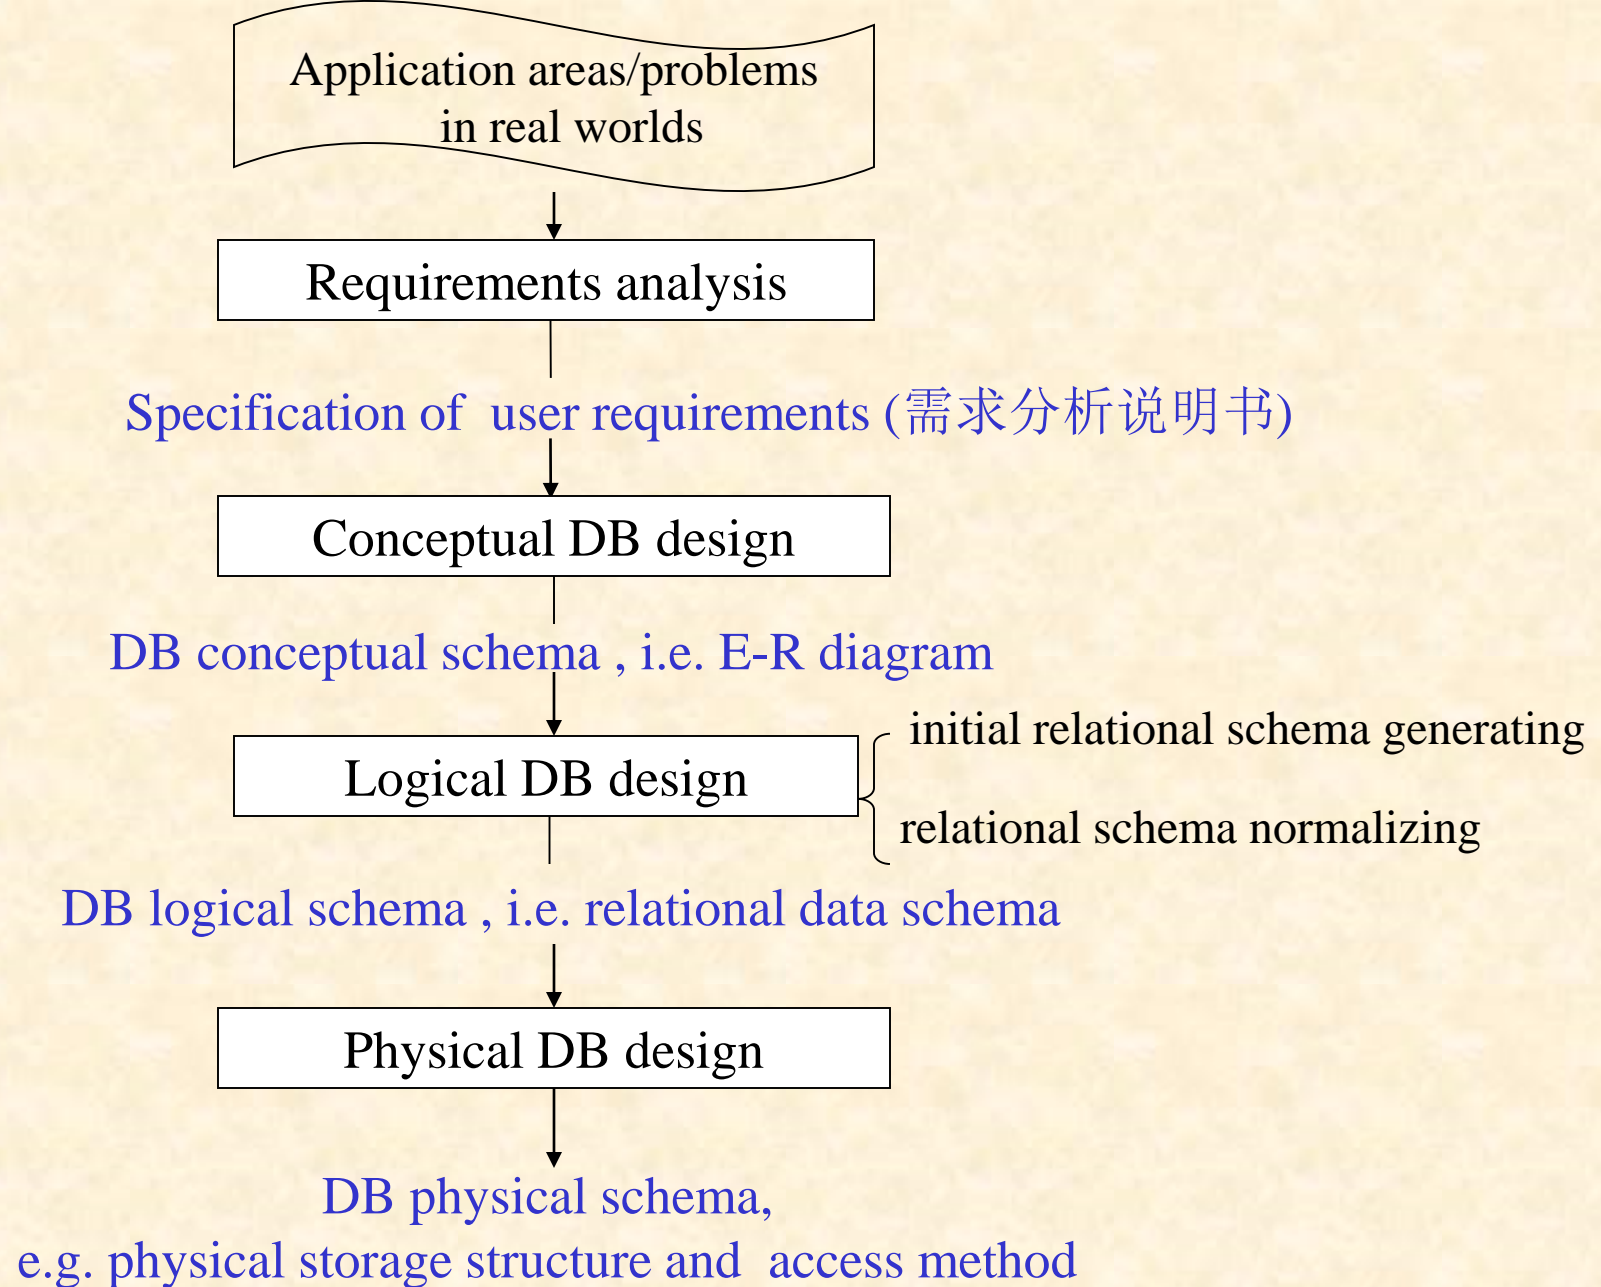
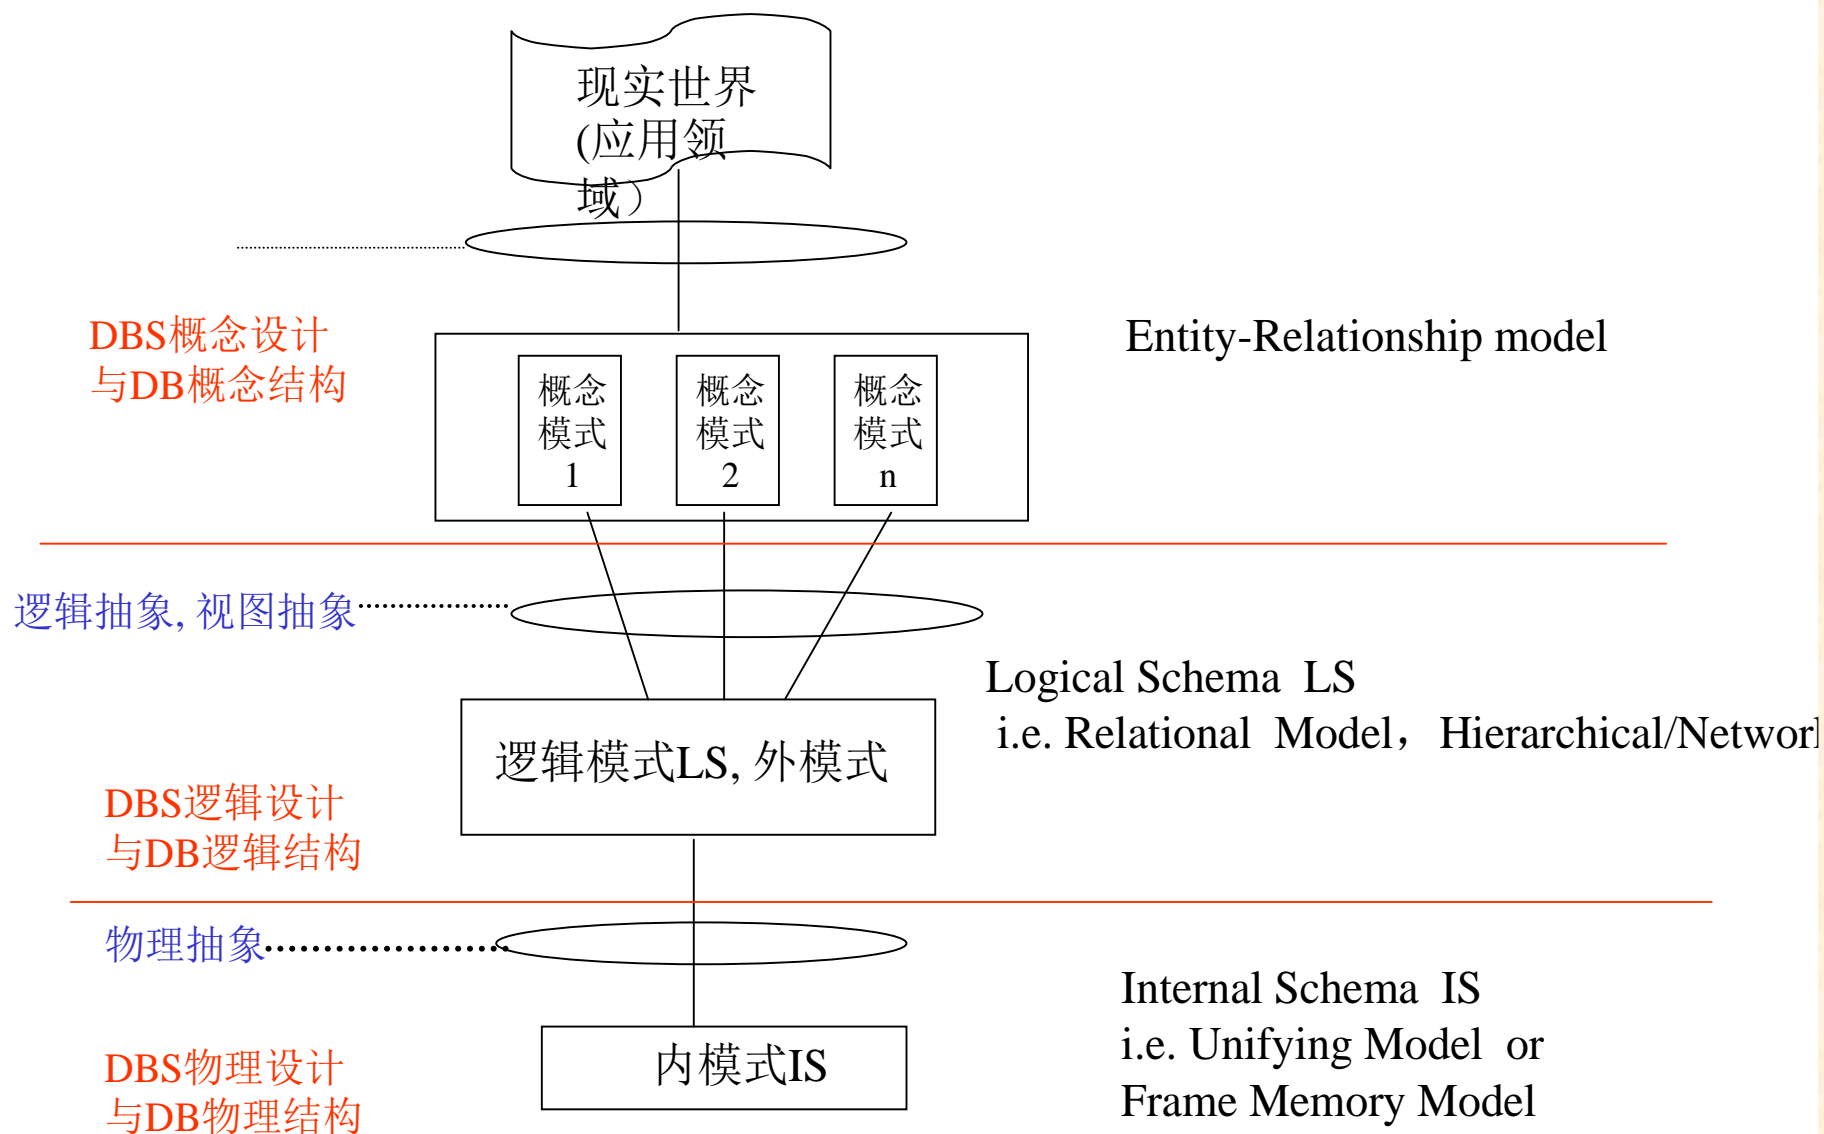


Fig. 6.0.1 DB design phases



DBS设计与设计阶段： 数据模型—数据抽象—数据模式

Overview of the Design Process (cont.)

■ 数据库应用系统DBAS设计

- DB, DBMS, users, **application programs**
- refer to Fig. 7.0.2

■ DBAS设计

- refer to Fig.7.0.3
- 参照软件工程中软件开发瀑布模型原理，DBAS的生命周期由**项目规划、需求分析、系统设计、实现与部署、运行管理与维护**等5个基本活动组成
- 根据DBAS的软件组成和各自功能，分为**数据组织与存储设计、数据访问与处理设计、应用设计**三条设计主线，分别用于设计数据库、数据库事务和应用程序



Overview of the Design Process (cont.)

- 根据数据库系统三级模式结构，DBAS设计阶段分为概念设计、逻辑设计、物理设计三个步骤，每一步设计内容涵盖了三条设计主线

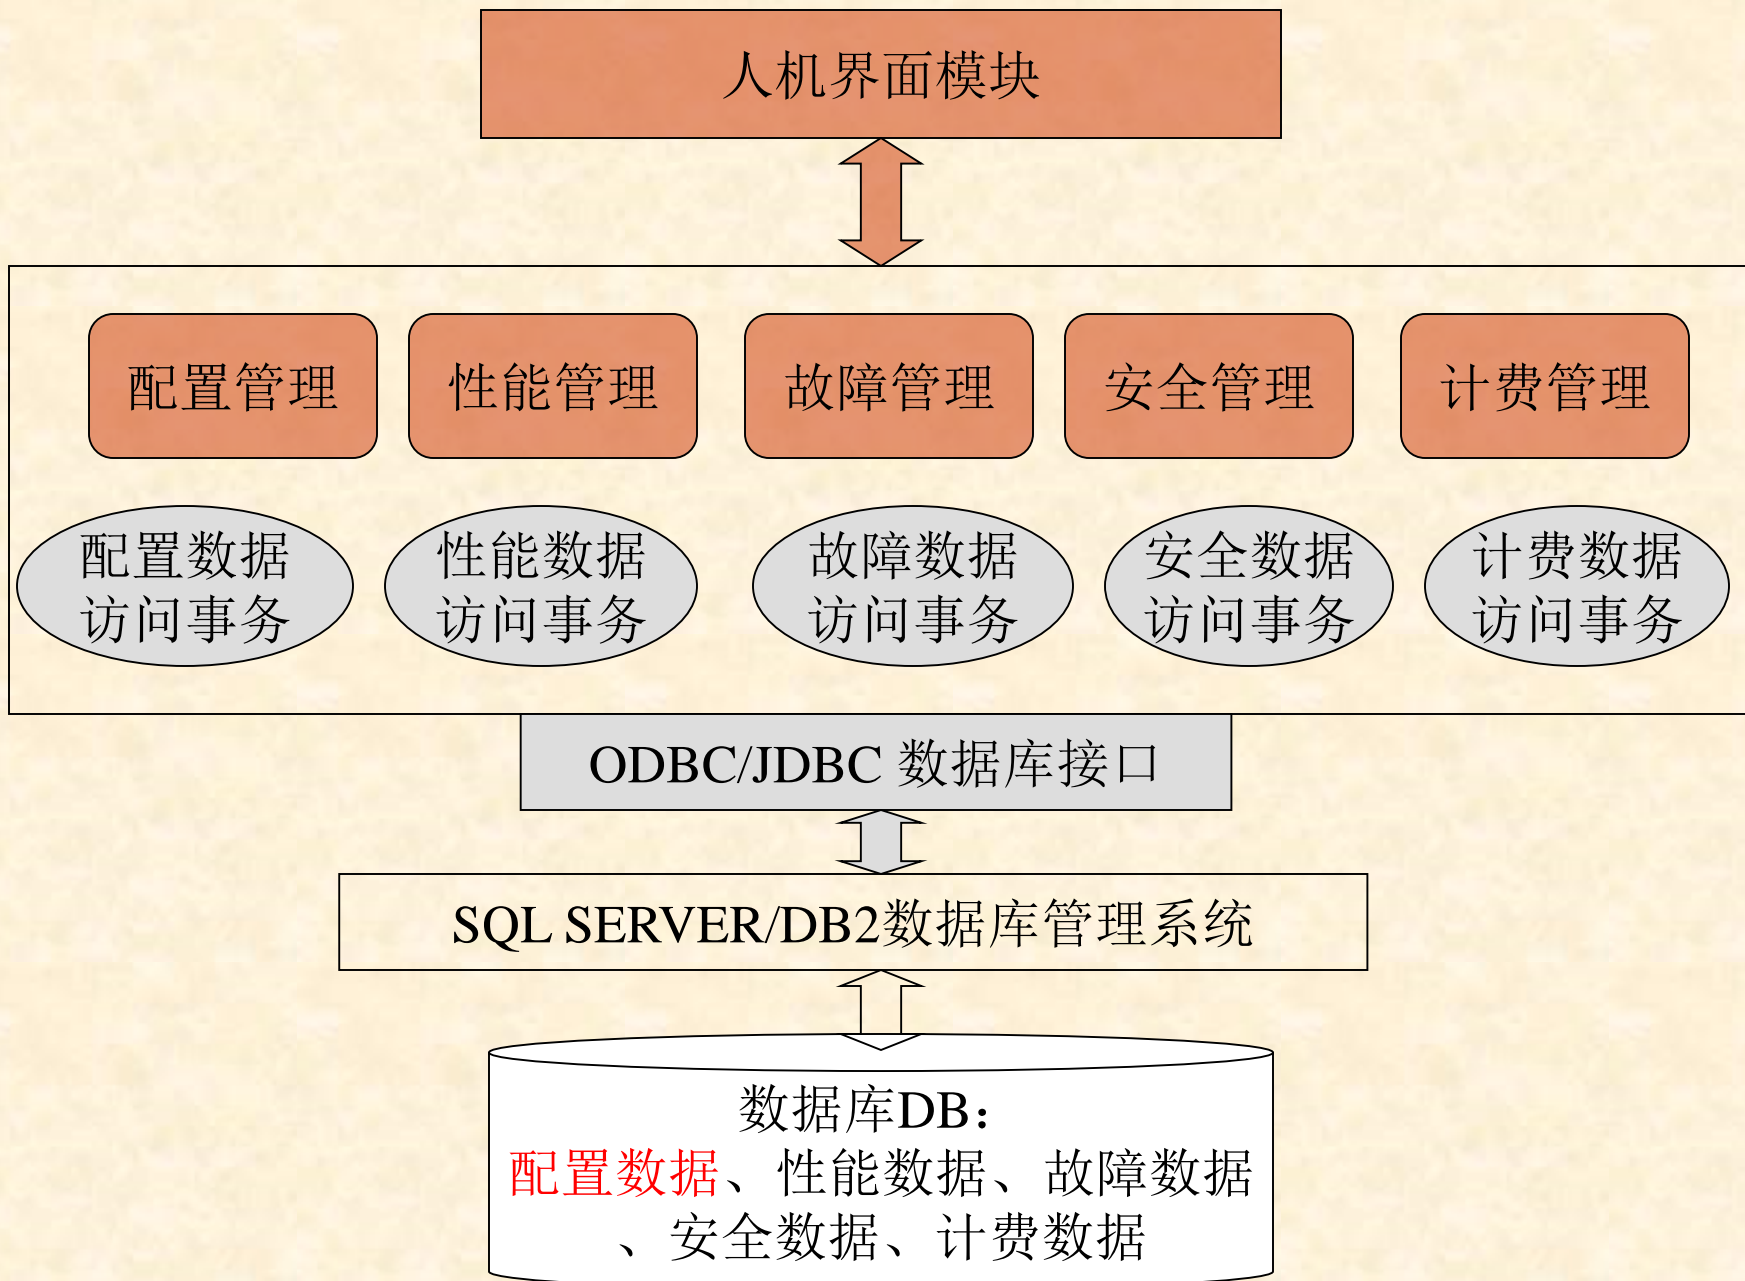


Fig. 7.0.2 电信网管系统示意图

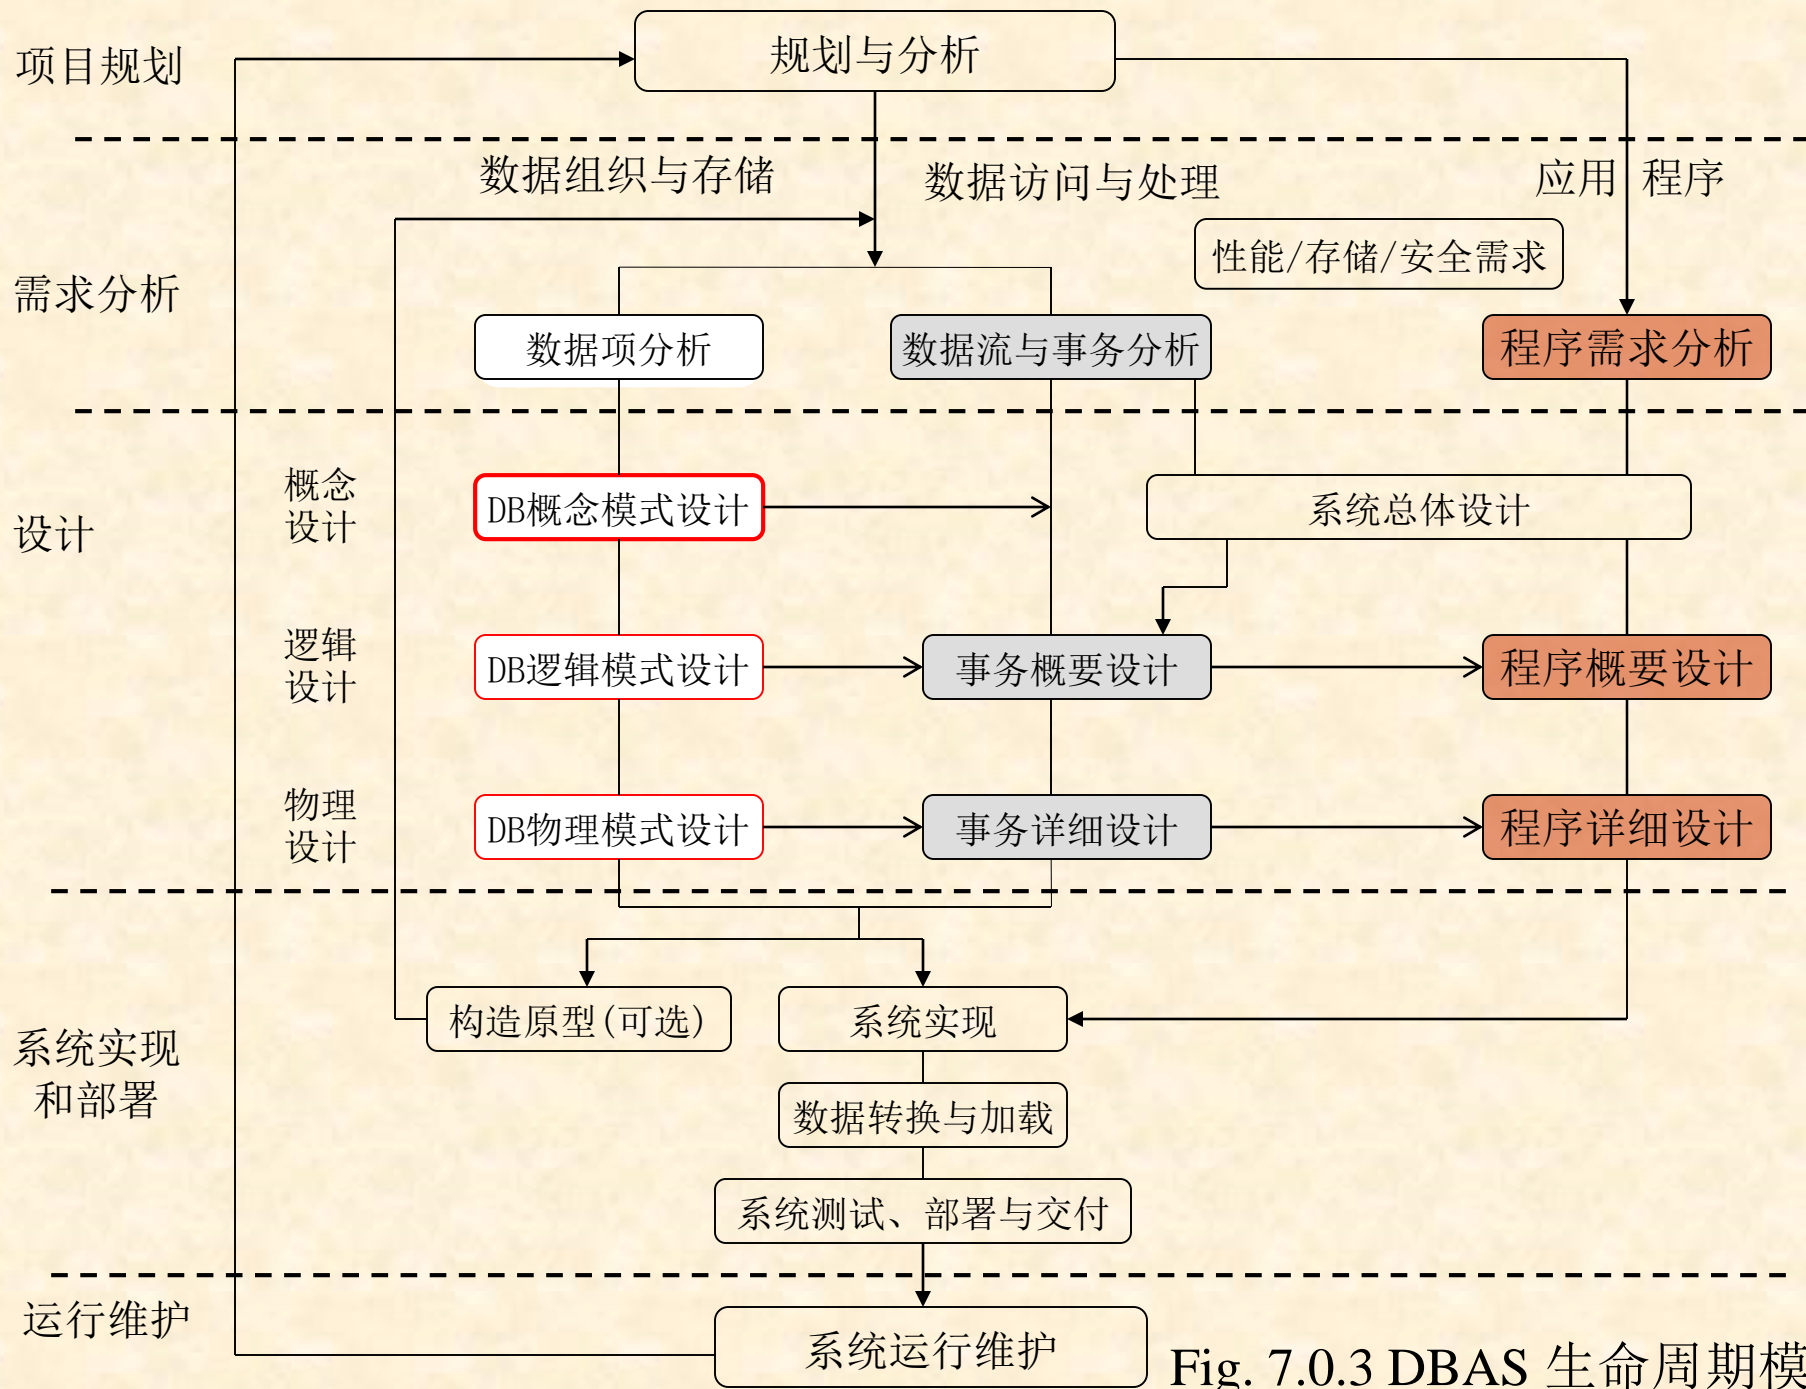


Fig. 7.0.3 DBAS 生命周期模型

中国移动网管支撑系统规范CMOSS2.0——性能指标要求

系统运行基本指标要求如下（仅供参考）：

No	分类	指标	参考值
1	可用性指标	系统可用性	99.9%（7×24）
2		故障平均恢复时间	2 小时
3		故障发生频次	3 次/年
4	应用性能指标	操作响应时延	≤5 秒
5		复杂报表生成时延	≤20 秒
6	系统容量指标	总用户数	按需
7		并发用户数	总用户数×10%
8		存储容量	按需
9		处理能力	按需
10		系统预留容量（存储和处理能力）	≥3 年，年增长≥20%
11		服务器 CPU 峰值利用率	≤70%
12		服务器内存峰值利用率	≤90%
13		有效存储空间利用率	≤80%

6.2 The Entity-Relationship Model

- The ER data model was developed to facilitate database design by allowing specification of an **enterprise schema** that represents the overall logical structure of a database.
- The ER model is very useful in mapping the meanings and interactions of real-world enterprises onto a conceptual schema. Because of this usefulness, many database-design tools draw on concepts from the ER model.
- The ER data model employs three basic concepts:
 - entity sets,
 - relationship sets,
 - attributes.
- The ER model also has an associated diagrammatic representation, the ER diagram, which can express the overall logical structure of a database graphically.

6.2.1 Entity Sets (实体集, 实体型)

- An **entity** is an object that exists and is distinguishable from other objects.
 - Example: specific person, company, event, plant
- An **entity set** is a set of entities of the same type that share the same properties.
 - Example: set of all persons, companies, trees, holidays
- An entity is represented by a set of attributes; i.e., descriptive properties possessed by all members of an entity set.
 - Example:
$$\text{instructor} = (ID, name, street, city, salary)$$
$$\text{course} = (course_id, title, credits)$$
- A subset of the attributes form a **primary key** of the entity set; i.e., uniquely identifying each member of the set.

Entity Sets (cont.)

- An entity consists of all values of its all attributes
 - e.g. Fig 6.2

- Entity-set

$customer\text{-}set = \{ \langle ID\text{-}value, name\text{-}value, street\text{-}value, city\text{-}value \rangle$
| $ID\text{-}value \in D_1, name\text{-}value \in D_2,$
 $street\text{-}value \in D_3, city\text{-}value \in D_4 \}$
 $\subseteq D_1 \times D_2 \times D_3 \times D_4$

Entity Sets

instructor_ID instructor_name

76766	Crick
45565	Katz
10101	Srinivasan
98345	Kim
76543	Singh
22222	Einstein

instructor

student-ID student_name

98988	Tanaka
12345	Shankar
00128	Zhang
76543	Brown
76653	Aoi
23121	Chavez
44553	Peltier

student

§ 6.2.2 Relationship Sets

- A **relationship** is an association among several entities

- e.g. in Fig. 6.2

Jone
customer
entity set

borrower
relationship

L-17
loan
entity set

- A **Relationship set** (联系集, 联系型) is a set of relationship of the same type
 - *note*: more than one relationship set among the same entity sets

§ 6.2.2 Relationship Sets

- A **relationship** is an association among several entities

Example:

44553 (Peltier) advisor 22222 (Einstein)
student entity relationship set *instructor* entity

- A **relationship set** is a mathematical relation among $n \geq 2$ entities, each taken from entity sets

$$\{(e_1, e_2, \dots, e_n) \mid e_1 \in E_1, e_2 \in E_2, \dots, e_n \in E_n\}$$

where (e_1, e_2, \dots, e_n) is a relationship

- Example:

$$(44553, 22222) \in \text{advisor}$$

Relationship Sets (cont.)

- The entity sets E_1, E_2, \dots, E_n **participate in** relationship set R
- A **relationship instance** in an E-R schema represents an association between the named entities in the real-world enterprise that is being modeled
 - e.g. in Fig. 6.2, the **relationship instance** between *Crick* and the *Tanaka*

Relationship Set *advisor*

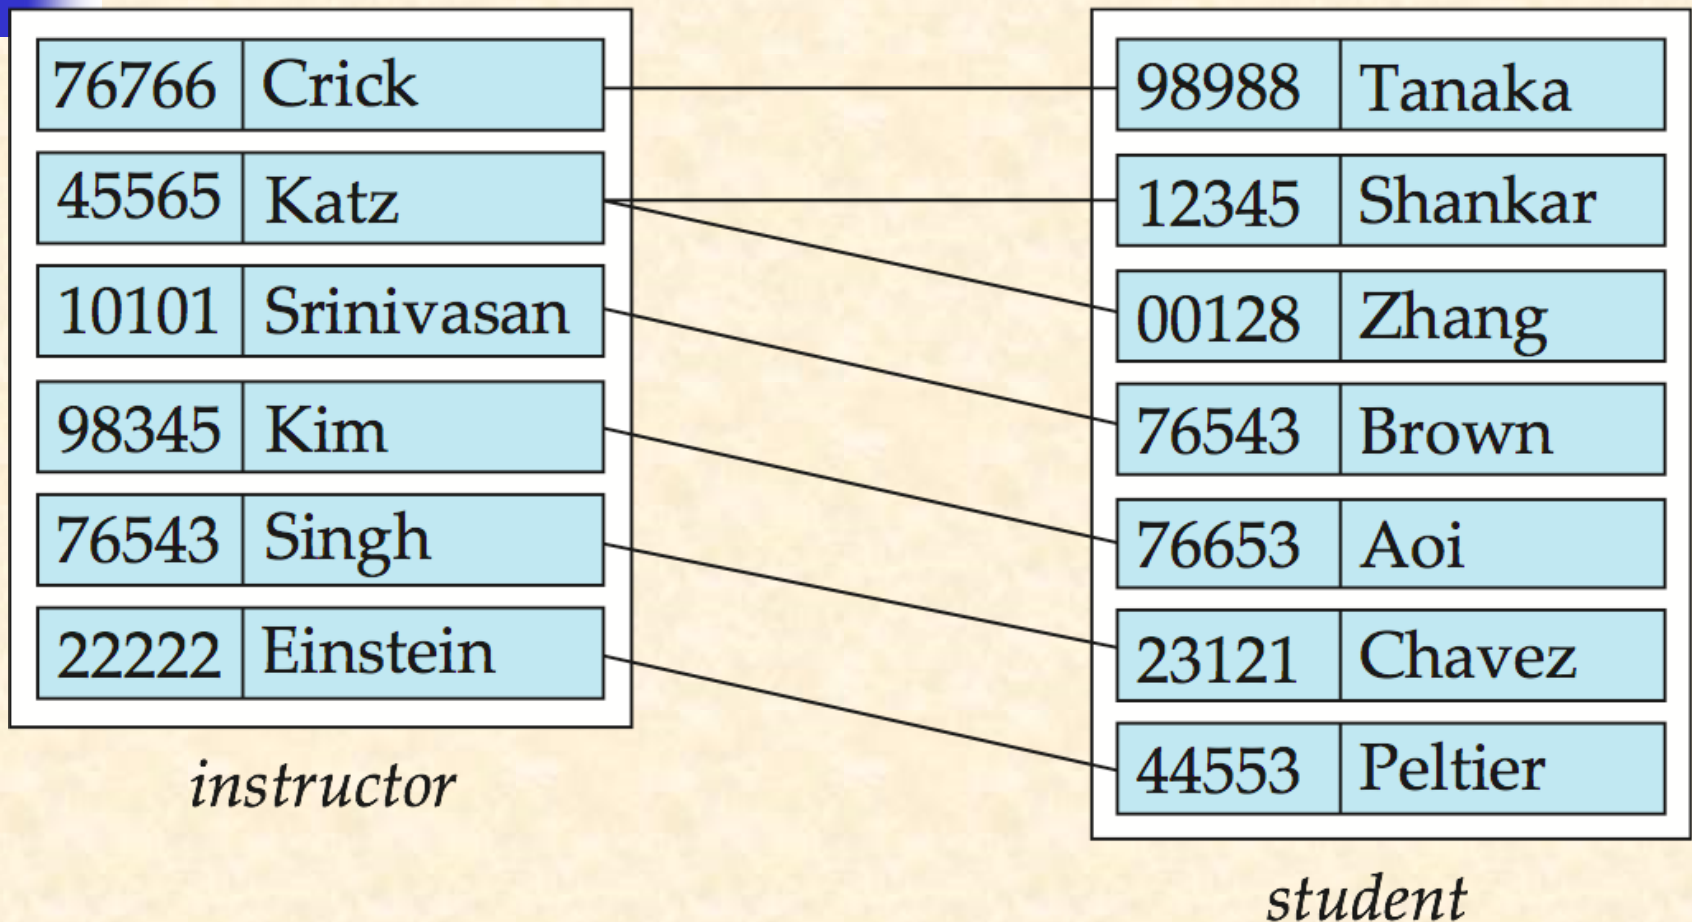
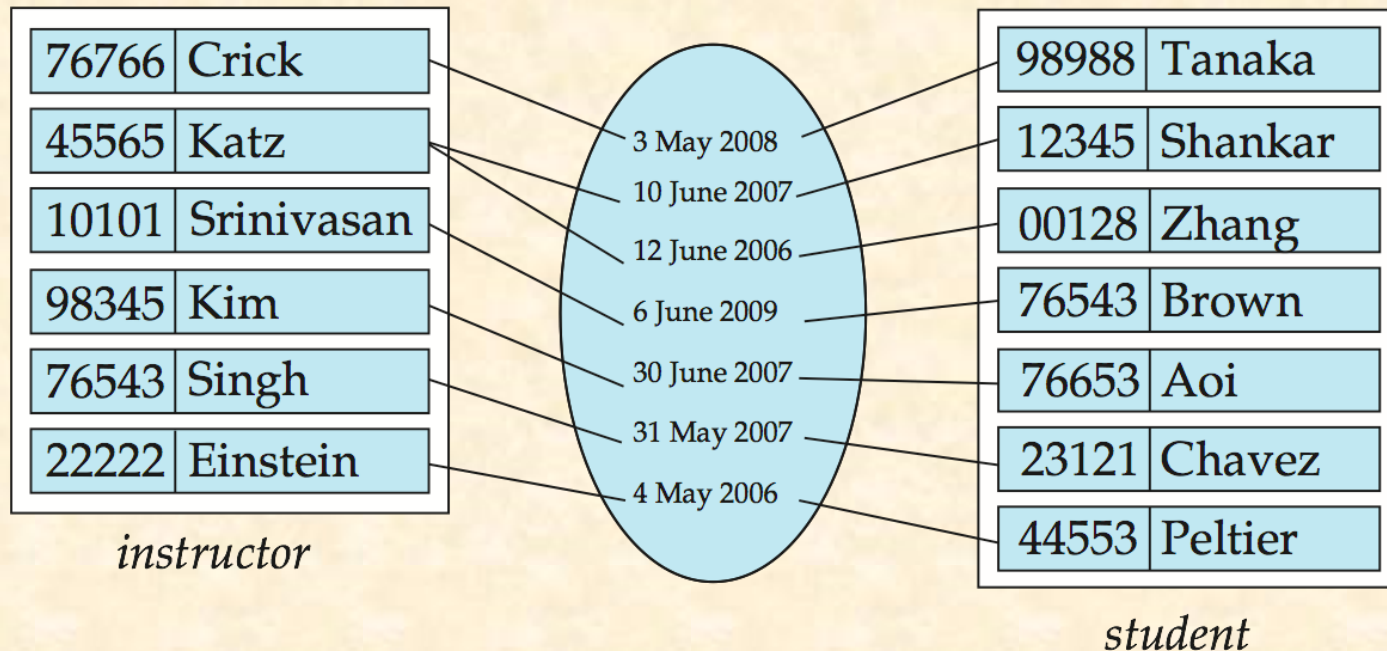


Fig.6.2

Relationship Sets (Cont.)

- An attribute can also be associated with a relationship set.
- For instance, the *advisor* relationship set between entity sets *instructor* and *student* may have the attribute *date* which tracks when the student started being associated with the advisor






Degree of a Relationship Set

- binary relationship
 - involve two entity sets (or degree two).
 - most relationship sets in a database system are binary.
- Relationships between more than two entity sets are rare. Most relationships are binary. (More on this later.)
 - ▶ Example: *students* work on research *projects* under the guidance of an *instructor*.
 - ▶ relationship *proj_guide* is a ternary relationship between *instructor*, *student*, and *project*

Relationship Sets (cont.)

■ Role

- the functions that an entity plays in a relationship is called that entity's role




The Entity-Relationship Model

1) 用长方形表示实体集，长方形内写明实体集名。

学生


教师

2) 用椭圆形表示实体集的属性，并用线段将其与相应的实体集连接起来。




The Entity-Relationship Model


3) 用菱形表示实体集间的联系，菱形内写上联系名，用线段分别与有关实体集连接起来，在线段旁标出联系的类型（1:1、1:n或m:n）。



1:1联系



1:n联系



m:n联系


§ 6.3 Attributes

- The *domain* or *value set* of the attribute
 - the set of permitted values for the attribute


- Attribute types:
 - **Simple** and **composite** attributes.
 - **Single-valued** and **multivalued** attributes
 - Example: multivalued attribute: *phone_numbers*
 - **Derived** attributes
 - Can be computed from other attributes
 - Example: age, given date_of_birth

Composite Attributes

composite
attributes



component
attributes





Attributes (cont.)

- Null value for an attribute means
 - the attribute “not applicable” for the entity, not existing
 - the value for the attribute exists, but is “unknown”



Redundant Attributes

- Suppose we have entity sets:
 - *instructor*, with attributes: *ID*, *name*, *dept_name*, *salary*
 - *department*, with attributes: *dept_name*, *building*, *budget*
- We model the fact that each instructor has an associated department using a relationship set *inst_dept*
- The attribute *dept_name* appears in both entity sets. Since it is the primary key for the entity set *department*, it replicates information present in the relationship and is therefore redundant in the entity set *instructor* and needs to be removed.
- BUT: when converting back to tables, in some cases the attribute gets reintroduced, as we will see later.

6.4 Mapping Cardinalities

- Mapping cardinalities
 - *semi-quantitatively* expressing the number of entities to which another entity can be associated via a relationship set
- For a binary relationship set ***R***, the mapping cardinality must be one of the following types: Fig.6.9, Fig.6.10, *from A to B*
 - *one to one*
 - *one to many*
 - *many to one*
 - *many to many*

6.4 Mapping Cardinalities

- Mapping cardinalities
 - *semi-quantitatively* expressing the number of entities to which another entity can be associated via a relationship set
- For a binary relationship set ***R***, the mapping cardinality must be one of the following types: Fig.6.9, Fig.6.10, *from A to B*
 - *one to one*
 - An entity in A is associated with *at most one* entity in B, and an entity in B is associated with *at most one* entity in A.
 - *one to many*
 - An entity in A is associated with *any number* (zero or more) of entities in B. An entity in B, however, can be associated with *at most one* entity in A.

—注意many-to-one和one-to-many 定义和方向！




Figure 6.9 Mapping cardinalities. (a) One-to-one. (b) One-to-many.

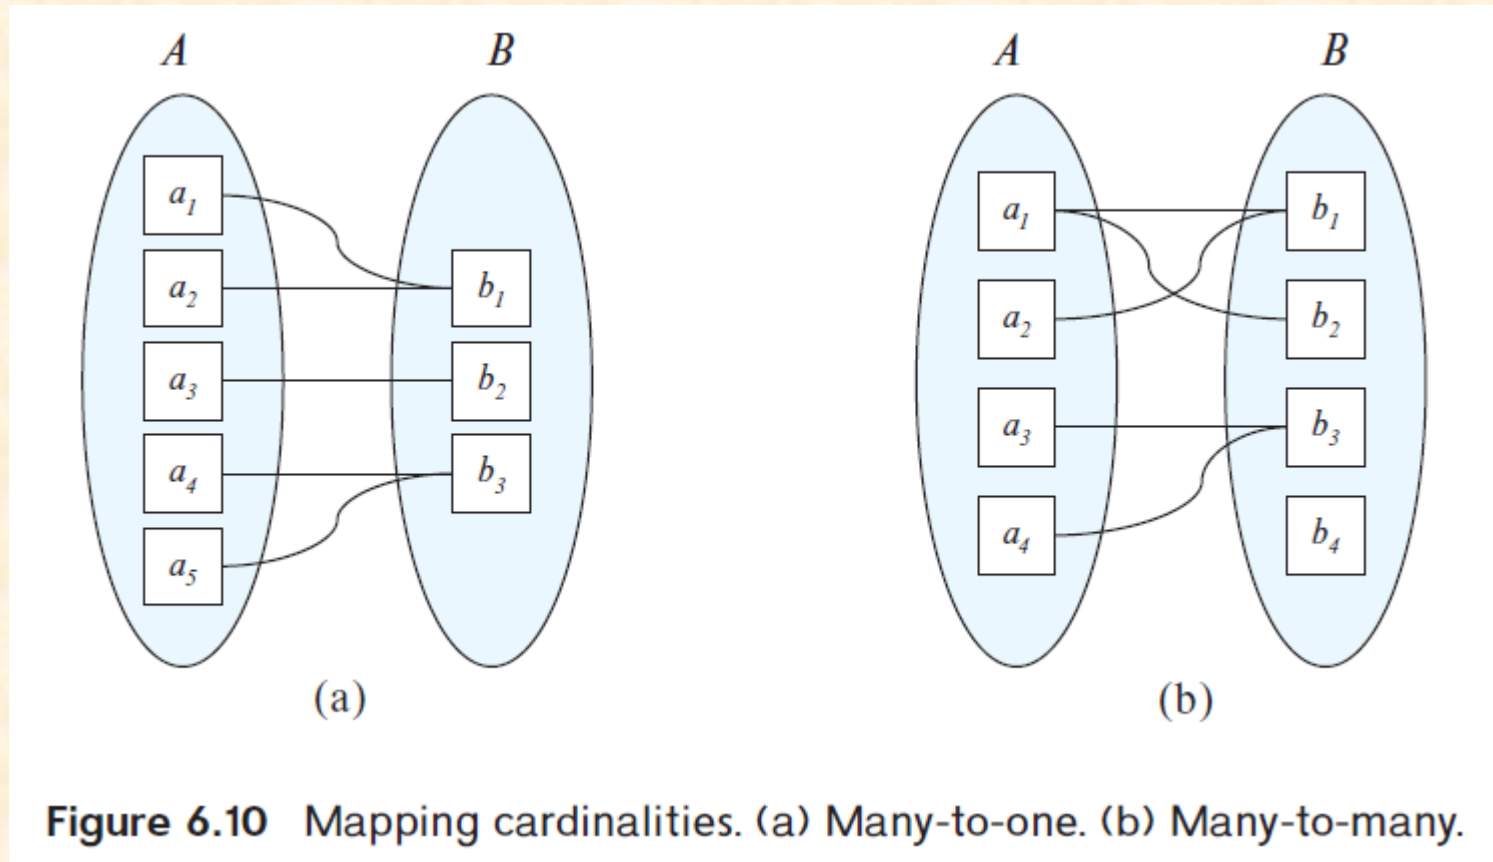
Note: some elements in A or B may not be mapped to any elements in the other set

Fig. 6.9 Mapping cardinalities

6.4 Mapping Cardinalities

- Mapping cardinalities
 - *semi-quantitatively* expressing the number of entities to which another entity can be associated via a relationship set
- For a binary relationship set ***R***, the mapping cardinality must be one of the following types: Fig.6.9, Fig.6.10, *from A to B*
 - *many to one*
 - An entity in A is associated with *at most one* entity in B. An entity in B, however, can be associated with *any number (zero or more)* of entities in A.
 - *many to many*
 - An entity in A is associated with *any number (zero or more)* of entities in B, and an entity in B is associated with *any number (zero or more)* of entities in A.

—注意many-to-one和one-to-many 定义和方向！




Note: Some elements in A or B may not be mapped to any elements in the other set

Fig. 6.10 Mapping cardinalities


Cardinality Constraints

- We express cardinality constraints by drawing either a directed line (\rightarrow), signifying “one,” or an undirected line ($—$), signifying “many,” between the relationship set and the entity set.
- One-to-one relationship between an *instructor* and a *student* :
 - A student is associated with at most one *instructor* via the relationship *advisor*
 - A *student* is associated with at most one *department* via *stud_dept*



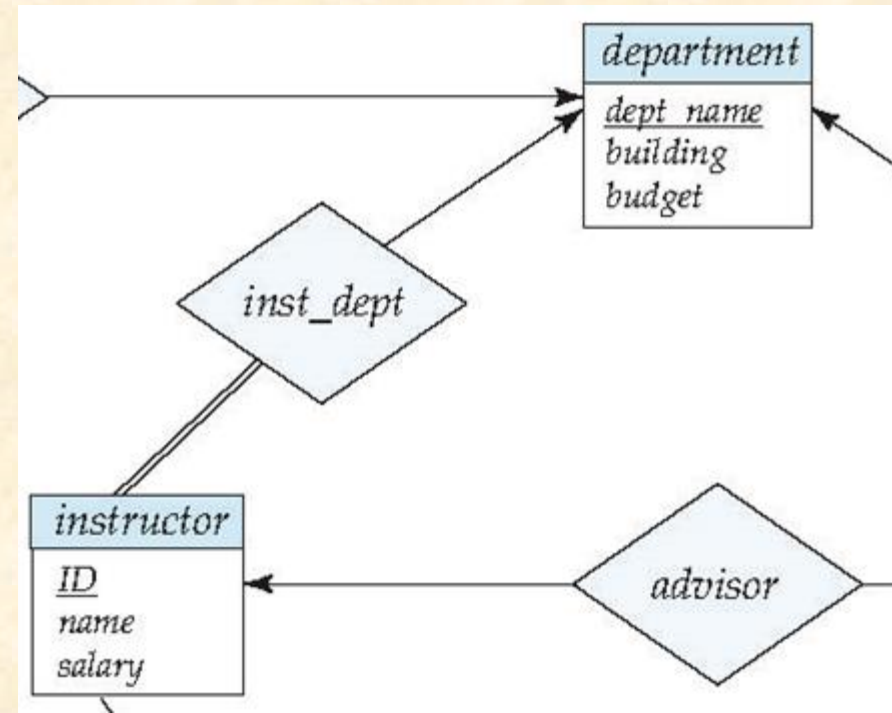
One-to-Many Relationship

- one-to-many relationship between an *instructor* and a *student*
 - an instructor is associated with several (including 0) students via *advisor*
 - a student is associated with at most one instructor via *advisor*,




Many-to-One Relationships

- In a many-to-one relationship between an *instructor* and a *department*,
 - an instructor is associated with at most one department via *inst_dept*,
 - and a department is associated with several (at least one) instructor via *inst_dept*





Many-to-Many Relationship


- A *section* is associated with several (possibly 0) timeslots via *sec_time_slot*
- A *timeslot* is associated with several (possibly 0) sections via *sec_time_slot*




Participate Constraints

- The participation of an entity E in a relationship R is *total*, if
 - every entity in E participates in at least one relationship in R
 - e.g. participation of *instructor* in *inst_dept* is *total*
 - refer to Fig.6.15 
 - every *instructor* **must** have a *department* associated to it via *inst_dept*
- The participation of an entity E in a relationship R is *partial*, if
 - some entities in E may not participate in any relationship in R
 - e.g. participation of *student* in *advisor* is *partial*, because some students maybe have no advisors
 - refer to Fig.6.15 

Participate Constraints (cont.)



(a) partial participation



(b) total participation

Fig. 7.0.4 Total/participation participation

Cardinality Limits for Participation

- **Cardinality limits** (参与的基数界限) are used to express *quantitative* constraints on participation
- *E.g. instructor, student, advisor*
 - 每个 *student* 最少有1个指导 *instructor*, 最多也只有1个指导 *instructor*
 - 每个 *instructor* 最多可以指导多个 *student*, 最少可以指导0个 *student*




Fig.6.13

Instructor can advise 0 or more students. A student must have 1 advisor; cannot have multiple advisors

Cardinality Limits for Participation

- /* 设联系 R 关联了 entity sets A 和 B , 为定量地描述 A 参与 R 的 total/partial participation 和 A 中的 entity 与 B 中的 entity 的 mapping cardinality, 引入实体参与联系的 cardinality limits
- A 参与 R 的基数下界 l_A 和上界 h_A , refer to Fig.7.0.5
 - A 中的每个实体 a 通过 R 关联了最少 l_A 个、最多 h_A 个 B 中实体 b
 - l_A : 对 A 中的每个实体 a , B 中至少有 l_A 个实体 b 通过 R 与其对应/关联
 - h_A : 对 A 中的每个实体 a , B 中至多有 h_A 个实体 b 通过 R 与其对应/关联




Fig.7.0.5 Illustration for cardinality limits

Cardinality Limits for Participation

(cont.)

- Note
 - A maximum value of * indicates no limit
 - *in some textbooks*, $\langle l_A, h_A \rangle$ is put at the side of the entity **B**
- Cardinality limits vs total/partial participation
 - A minimum value $l_A = 0$: **A** is partial participation of **R**
 - A minimum value $l_A > 0$: **A** is total participation of **R**, equivalent to *double line*

Cardinality Limits for Participation


(cont.)

- Cardinality limits vs mapping cardinality
 - 设联系 R 关联了实体集 A 和 B , 利用 A 的基数界限 $\langle l_A, h_A \rangle$ 中的 h_A 、 B 的基数界限 $\langle l_B, h_B \rangle$ 中的 h_B , 可推导出联系 R 的映射基约束
 - $\langle h_B, h_A \rangle$ 表示了联系 R 的从 A 到 B 的映射基约束 !!!!
 - e.g. in Fig.7.10 , considering mapping cardinality form *instructor* to *student*
 - mapping cardinality form *instructor* to *student* depends on $\langle h_{\text{student}}, h_{\text{instructor}} \rangle$
 - for student, $h_{\text{student}}=1$, for instructor, $h_{\text{instructor}}=*$
 - so, $\langle h_{\text{student}}, h_{\text{instructor}} \rangle = \langle 1, * \rangle$, and is *one to many*


Cardinality Limits for Participation

(cont.)

- The associations between cardinality limits and the mapping cardinality are classified as follows, and illustrated by Fig.7.0.6



(a)



(b)

From A to B , **one** to **one** (1:1),

$$\langle l_A, h_A \rangle = \langle 1, 1 \rangle,$$


$$\langle l_B, h_B \rangle = \langle 0, 1 \rangle;$$

From A to B , **one** to **many** (1:2),

$$\langle l_A, h_A \rangle = \langle 0, 2 \rangle,$$

$$\langle l_B, h_B \rangle = \langle 1, 1 \rangle$$

Fig.7.0.6-1 Associations between cardinality
limits and mapping cardinality




(a)

Form A to B , many to one (2:1),

$$\langle l_A, h_A \rangle = \langle 1, 1 \rangle,$$

$$\langle l_B, h_B \rangle = \langle 1, 2 \rangle$$



(b)

Form A to B , many to many (2:2)

$$\langle l_A, h_A \rangle = \langle 1, 2 \rangle$$

$$\langle l_B, h_B \rangle = \langle 1, 2 \rangle$$

6.5 Keys

- **Key** is a set of attributes (of a *entity set* or *relationship set*)
 , in which there are one or more attributes
 - the values of these attributes in one *entity* can be used to uniquely distinguish this entity from others , or
 - the values of these attributes in one *relationship* are used to uniquely identify the relationship
- Keys include
 - superkey (超键), candidate key (候选键), primary key (主键)

Keys For Entity Sets

- A *super key* of an entity set is a set of one or more attributes, whose values uniquely determine each entity in the entity set
 - e.g. $\{instructor_id, instructor_name\}$
 - the super key may contain extraneous attributes
 - e.g. *instructor_name*
- A candidate key is the minimal super key
 - *non-redundant* super key
 - e.g. *instructor-id* is the candidate key of *instructor*

Keys For Entity Sets (cont.)

- The *primary key* is a candidate key chosen by the database designer as the principal means of identifying entities within an entity set
 - although several candidate keys may exist, one of the candidate keys is selected to be the *primary key*
 - need to consider semantics of relationship set in selecting the *primary key* in case of more than one candidate key

构造关系表时，如果有多个候选键，最好选取数值型（`int`，`float`）候选键作为关系表主键，便于提高基于主键的查询速度

—不要选字符串型属性，如 `varchar`、`datetime`

—e.g. `studentname`, `instructorName`

Keys for Relationship Sets

- Keys for relationship sets **R** on entities E_1, E_2, \dots, E_n
 - $R = \{(e_1, e_2, \dots, e_n) \mid e_1 \in E_1, \dots, e_n \in E_n\}$
 $\subseteq E_1 \times E_2 \times \dots \times E_n$
 , how to uniquely distinguish each relationship **instances** $\{(e_1, e_2, \dots, e_n) ?$
 - **R** is the combination of E_1, E_2, \dots, E_n , each E_i can be uniquely distinguished by $\text{primary_key}(E_i)$, $1 \leq i \leq n$, so the set of all attributes in $\text{primary_key}(E_1), \text{primary_key}(E_2), \dots, \text{primary_key}(E_n)$ can be used to recognize (e_1, e_2, \dots, e_n)

Keys for Relationship Sets (cont.)

- The **super**_key for R
 - $\text{primary_key}(E_1) \cup \text{primary_key}(E_2) \dots \cup \text{primary_key}(E_n)$
 - e.g. in Fig.6.3, $(\text{InstructorID}, \text{StudentIDr})$ is the super key of *advisor* ▶
 - *note*
 - if the attribute names of primary-keys are not unique, the **attributes with the same names** should be renamed
- The candidate keys for R
 - minimal, non-redundant super keys

Keys for Relationship Sets (cont.)

- The *candidate* keys or the *primary* key for a **binary** relationship set R among entity sets A and B can be decided as follows, in accordance with the mapping cardinality of R

- R is *many-to-many*, ►

$$\text{primary_key}(R) = \text{primary_key}(A) \cup \text{primary_key}(B)$$


- R is *many-to-one from A to B*, !!

$$\text{primary_key}(R) = \text{primary_key}(A)$$

- R is *one-to-many*, !! ►

$$\text{primary_key}(R) = \text{primary_key}(B)$$

Keys for Relationship Sets (cont.)

- R is *one-to-one*, 

$$\text{primary_key}(R) = \text{primary_key}(A)$$

$$\text{or: primary_key}(R) = \text{primary_key}(B)$$