



# 数据库系统原理

---

## Database System Principle

**邵莹侠**

**Email: [shaoyx@bupt.edu.cn](mailto:shaoyx@bupt.edu.cn)**

**北京邮电大学计算机学院**

**计算机应用技术中心**



## PART 3

---

# QUERY PROCESSING AND OPTIMIZATION



# Chapter 15

---

## Query Processing



# Introduction

---

## ■ Query

- one or more operations on a database, or requests for DB access
- a high-level database language (e.g. SQL, QUEL, declarative ) statement, or a sequence of statements

## ■ Query processing

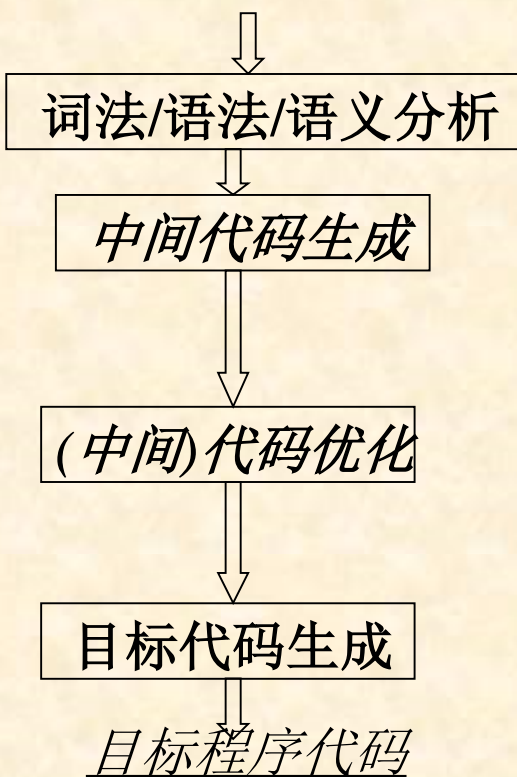
- (DBMS' ) activities involved in extracting data from a database, including *parsing and translation*, *query optimization*, and *evaluation* (执行)

## ■ Programs compiling and executing vs. Query processing

- refer to Fig. 15.0.1

## C, Pascal programs

程序编译 / 编译器



process / thread

程序执行 / OS

进程管理  
并发控制  
进程调度  
死锁处理

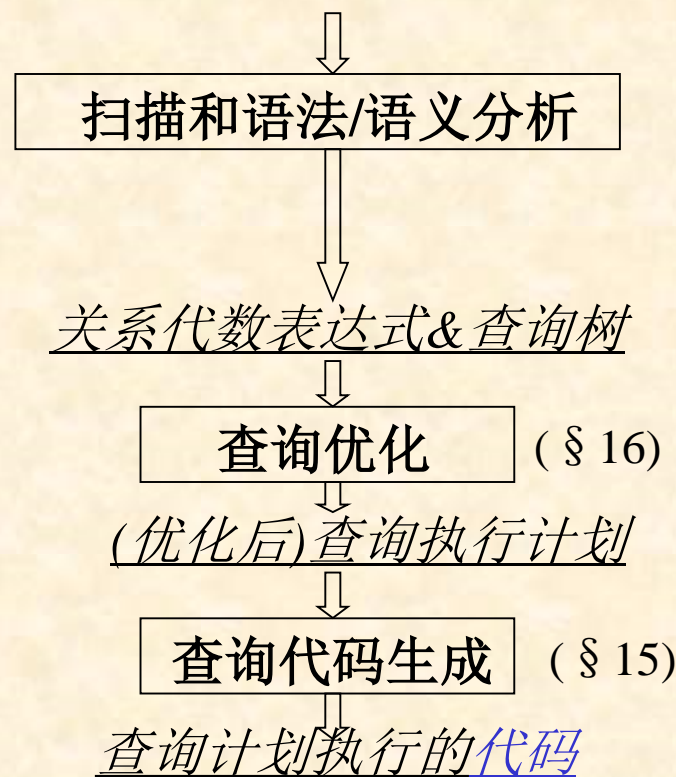
Chapter.  
15, 16

Chapter.  
17, 18, 19

Fig. 15.0.1

## query

Query processing / DBMS



transaction

事务处理 / DBMS

事务管理 ( § 17)

并发控制 ( § 18)

事务调度 ( § 18)

死锁处理 ( § 19)

恢复技术 ( § 19)



# Chapter 15: Query Processing

---

- Overview
- Measures of Query Cost
- Selection Operation
- Sorting
- Join Operation
- Other Operations
- Evaluation of Expressions



## Main Parts Chapter 15

---

- Basic steps in query processing, § 15.1 Overview
- Measures of query cost, § 15.2
- Evaluating of individual relational algebra operations, § 15.3- § 15.6
  - selection, sorting, join, project, set operations, ...
- Evaluating of expression, i.e. a sequence of relational algebra operations, § 15.7



## § 15.1 Overview

---

- Steps in query processing
  - Fig.15.1



1. Parsing and translation
2. Optimization
3. Evaluation

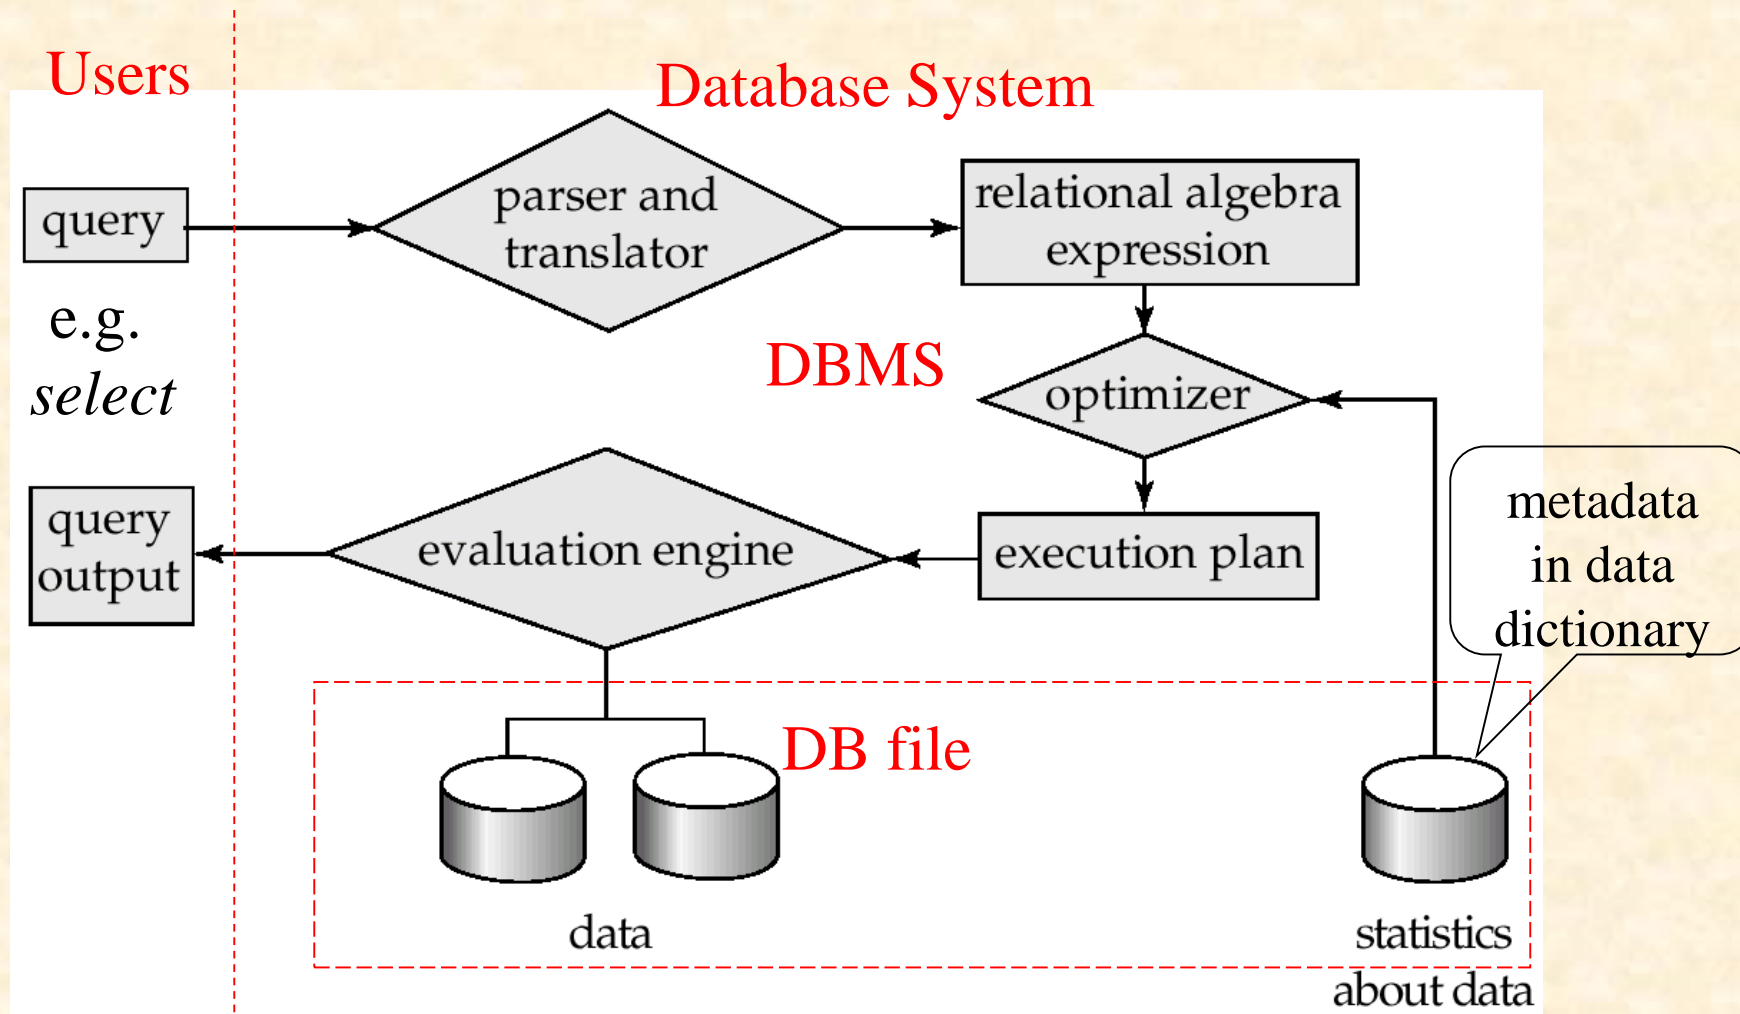


Fig.15.1 Steps in query processing



## Overview (cont.)

---

- **Step1.Parsing and translation**

- translate the query into a parser-tree representation
  - parser checks syntax, verifies the correctness of the relations in the query
  - parser also replaces the view in the query with the relations on which this view is built on
- the parser-tree is then translated into relational algebra expression



## Overview (cont.)

---

- **Step2. Query Optimization**

- amongst all equivalent query *evaluation plans*, choose the one with lowest cost
- the cost is estimated using statistical information in data dictionary
  - e.g. the number of tuples in each relation, size of tuples, etc

```
select salary  
from instructor  
where salary < 75000;
```

## Step2. Optimization

- A relational algebra expression may have many equivalent expressions
  - E.g.,  $\sigma_{salary < 75000}(\Pi_{salary}(instructor))$  is equivalent to  $\Pi_{salary}(\sigma_{salary < 75000}(instructor))$
- Each relational algebra operation can be evaluated using one of several different algorithms
  - e.g. the *select* operation can be evaluated using one of A1, A2, ..., A11 algorithms in § 15.3
  - e.g. merge *join*, hash *join* operation
- Correspondingly, a relational-algebra expression can be evaluated in many ways

## Step2. Optimization

- Annotated expression specifying detailed evaluation strategy is called an **evaluation-plan**, such as *index*, *evaluation algorithms*, etc.
- e.g., can use an index on *salary* to find *instructors* with *salary* < 75000,
- or can perform complete relation scan and discard instructors with *salary* ≥ 75000

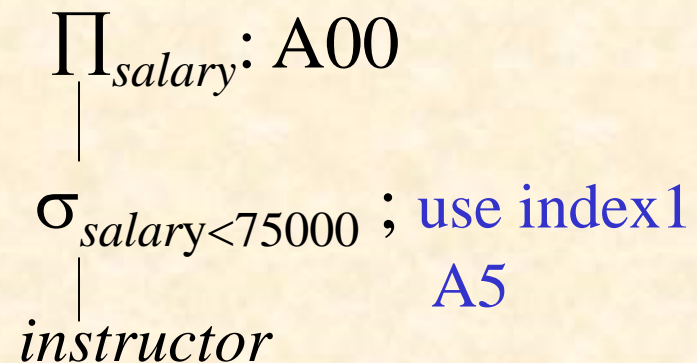


Fig. 15.2 A query-evaluation plan

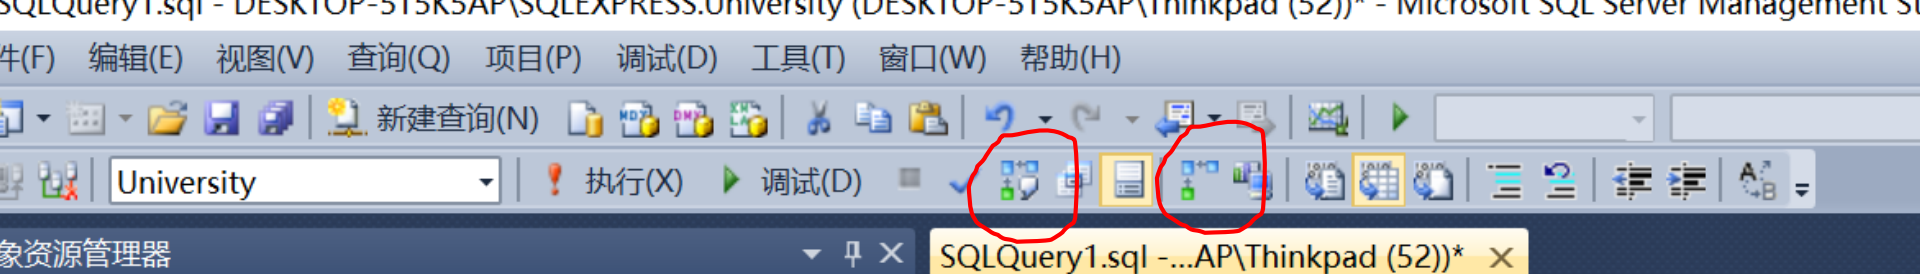


## Overview (cont.)

---

### ■ Step3. Evaluation plan execution

- the query-execution engine takes a *optimized* query-evaluation plan, executes that plan, and returns the answers to the query

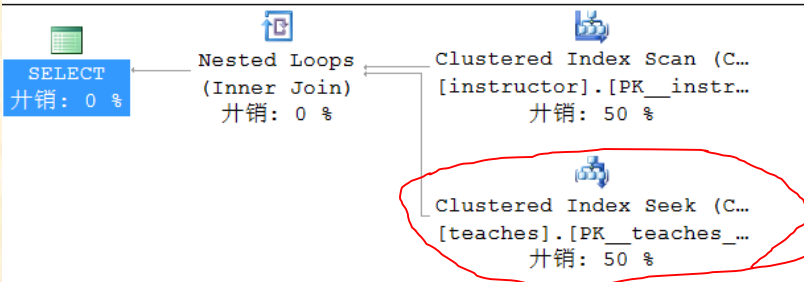


```
select name, course_id from instructor , teaches
where instructor.ID = teaches.ID
and instructor. dept_name = 'Art'

select dept_name, avg (salary) as avg_salary
from instructor group by dept_name;
```

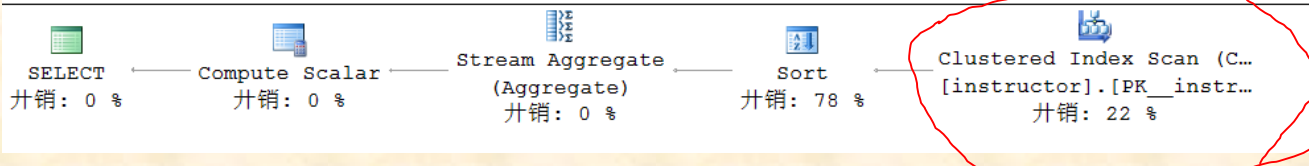
查询 1: (与该批有关的) 查询开销: 31%

select name, course\_id?from instructor , teaches? where instructor.ID = teaches.ID and instructor. dept\_name = 'Art'

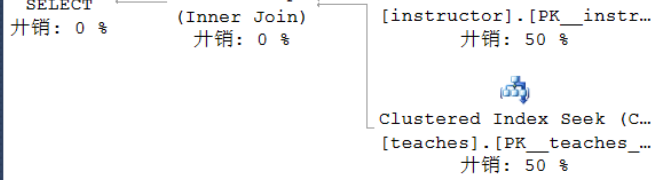


查询 2: (与该批有关的) 查询开销: 69%

select dept\_name, avg (salary) as avg\_salary?from instructor?group by dept\_name

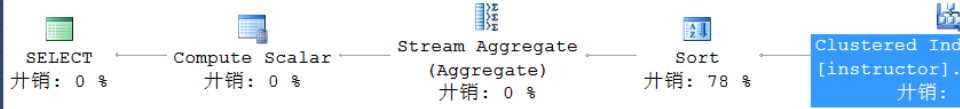






```
select name, course_id from instr  
where instructor.ID = teaches.ID  
and instructor. dept_name =  
  
select dept_name, avg (salary) as  
from instructor group by dept_name
```

查询 2: (与该批有关的) 查询开销: 69%  
select dept\_name, avg (salary) as avg\_salary?from instructor?group by



Clustered Index Scan (Clustered)	
整体扫描聚集索引或只扫描一定范围。	
物理运算	Clustered Index Scan
Logical Operation	Clustered Index Scan
估计的执行模式	Row
存储	RowStore
Estimated Operator Cost	0.0032831 (22%)
Estimated I/O Cost	0.003125
Estimated CPU Cost	0.0001581
Estimated Subtree Cost	0.0032831
Estimated Number of Executions	1
Estimated Number of Rows	1
Estimated Row Size	26 字节
Ordered	False
节点 ID	3
对象	
[University].[dbo].[instructor].	
[PK_instruct_3214EC277D1BCB90]	
Output List	
[University].[dbo].[instructor].dept_name, [University].	
[dbo].[instructor].salary	

就绪







## § 15.2 Measures of Query Costs

---

- Cost is generally measured as total elapsed time for answering query
  - many factors contribute to time cost
    - *disk accesses, CPU, or even network communication*
- Typically ***disk access*** is the predominant cost, and is also relatively easy to estimate. Measured by taking into account
  - number of seeks \* average-seek-cost
  - number of blocks read \* average-block-read-cost
  - number of blocks written \* average-block-write-cost
    - cost to write a block is greater than cost to read a block
      - data is read back after being written to ensure that the write was successful



## Measures of Query Cost (cont.)

---

- For simplicity we just use the **number of block transfers** *from disk and the number of seeks* as the cost measures
  - $t_T$  – time to transfer one block
  - $t_S$  – time for one seek
  - cost for  $b$  block transfers plus  $S$  seeks
$$b * t_T + S * t_S$$
- We ignore CPU costs for simplicity
  - real systems do take CPU cost into account
- We do not include cost to writing output to disk in our cost formulae



## Measures of Query Cost (Cont.)


---

- Several algorithms can reduce disk IO by using extra buffer space
  - Amount of real memory available to buffer depends on other concurrent queries and OS processes, known only during execution
    - We often use worst case estimates, assuming only the minimum amount of memory needed for the operation is available
- Required data may be buffer resident already, avoiding disk I/O
  - But hard to take into account for cost estimation



## § 15.3 Selection Operation

---

- Selection operation on a relation  $r$  by **file scan**
  - locating and scanning the file in which  $r$  is stored to retrieving the *file records* satisfying the selection conditions
- E.g.  $\Pi_{\text{salary}}(\sigma_{\text{salary} < 2500}(\text{instructor}))$  




## Selection Operation

---

- Types of query conditions (查询条件类型)
  - equality(等值), e.g. *balance* = 100
  - range (范围), e.g. *balance* between 50 and 400
  - comparison (比较), e.g. *balance* > 300
- Several file scan algorithms
  - linear search/scan – A1
  - selections using indices – A2, A3, A4
  - selections involving comparisons – A5, A6
  - complex selections – A7, A8, A9, A10

## A1: File scan by linear search

- Algorithm A1 (**linear search**).
- Scan each file block and test all records to see whether they satisfy the selection condition. 
  - cost estimate =  $b_r$  block transfers + 1 seek
    - $b_r$  denotes number of blocks containing records from relation  $r$
  - if selection is on a key attribute, can stop on finding record
    - cost =  $(b_r/2)$  block transfers + 1 seek

—必须扫描全部blocks，方能找到全部满足查询条件的数据





## A1: File scan by linear search

---

- Linear search can be applied regardless of
  - selection condition or
  - ordering of records in the file, or
  - availability of indices
- Note: binary search generally does not make sense since data is not stored consecutively
  - except when there is an index available,
  - and binary search requires more seeks than index search

## Selections Using Indices

- **Index scan** – search algorithms that use an index

—— e.g. in SQL Server, *index seek* 

- selection condition must be on search-key of index,  
i.e. *instructor*(*ID*, name, *dept-name*, salary)

- **A2 (primary index, equality on key such as *ID*)**. 

Retrieve a single record that satisfies the corresponding equality condition, using a  $B^+$ -tree as **the clustering/primary index**

- $Cost = (h_i + 1) * (t_T + t_S)$
- $h_i$  : height of the tree
- $t_T$  : time to transfer one block
- $t_S$  : *time for one seek*



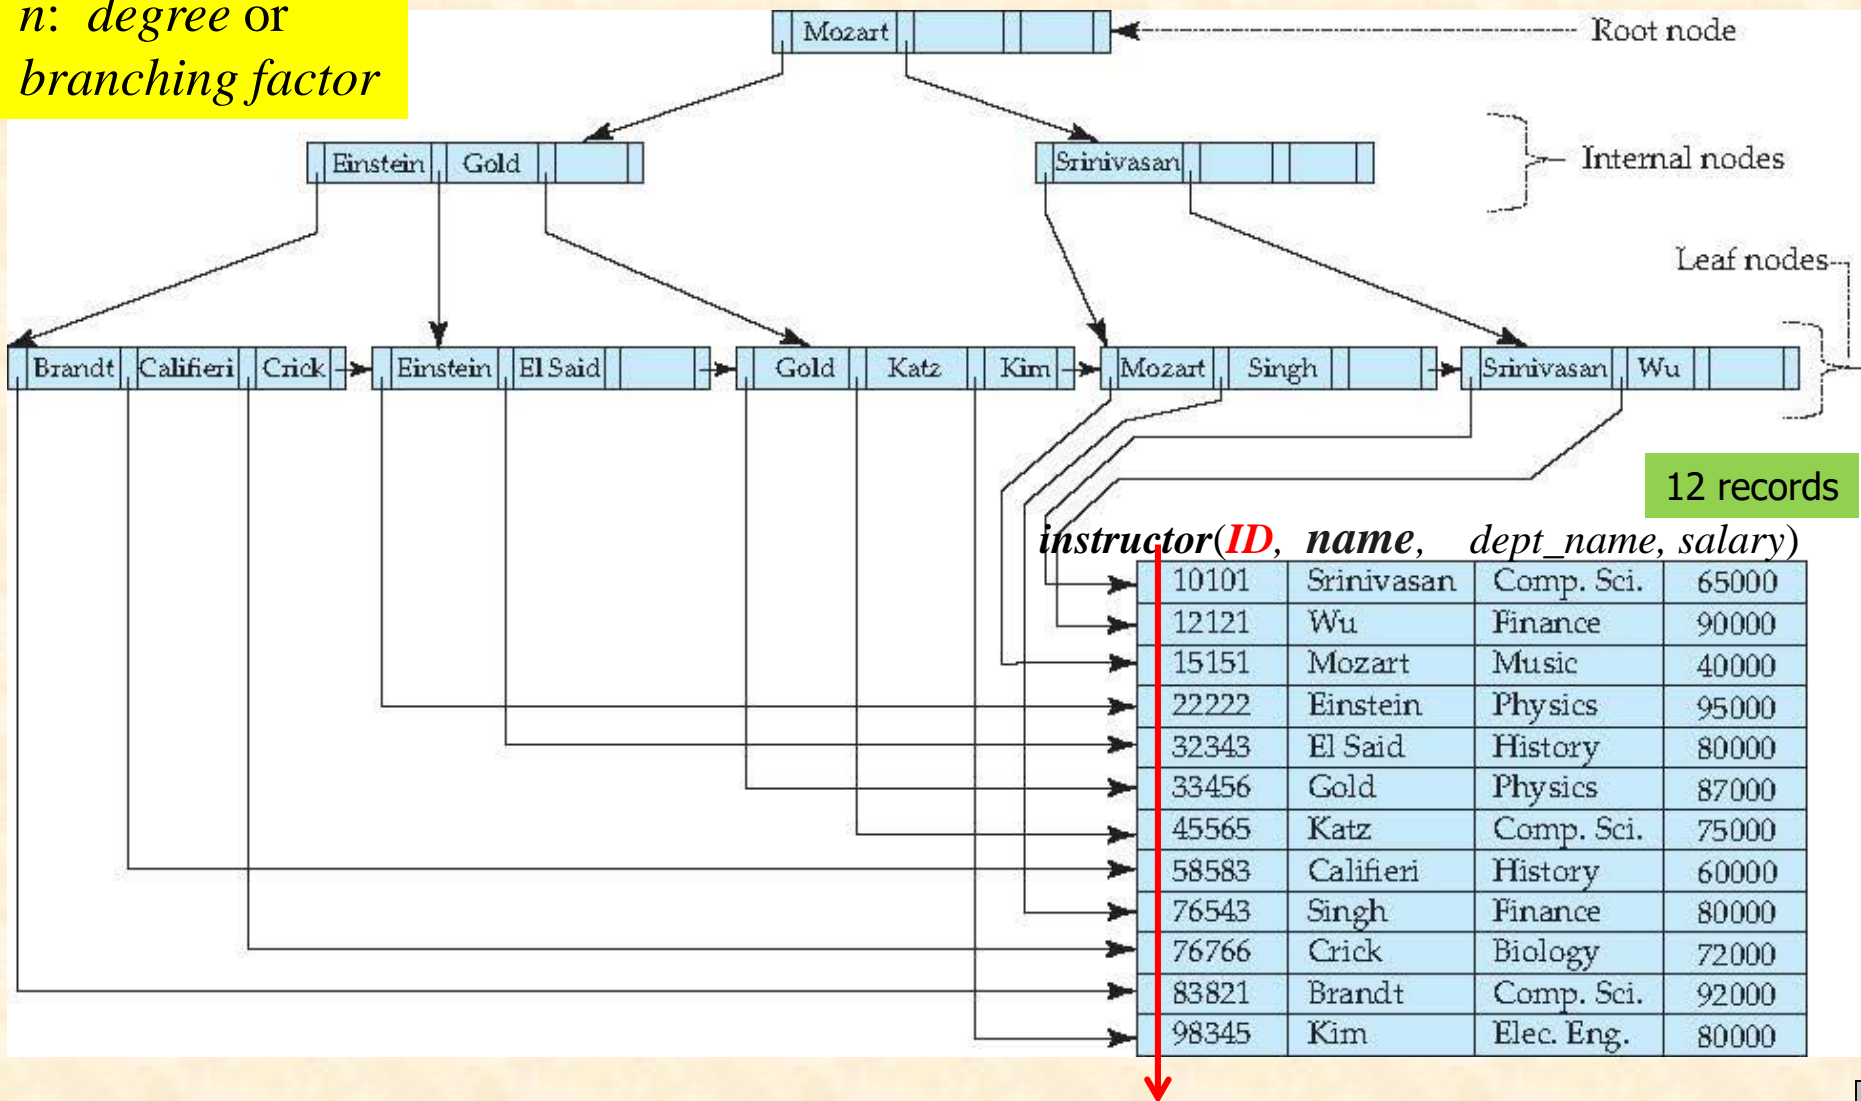
# Example of B<sup>+</sup>-Tree ( $n=4$ )

**Not a  
primary index!!**

非根、非叶结点的儿子  
结点个数  $k$ :

$$2 = \lceil n/2 \rceil \leq k \leq n = 4$$

$n$ : degree or  
branching factor






## Selections Using Indices

---

- **A3 (primary index, equality on nonkey)**

- Retrieve multiple records. 

- Records will be on consecutive blocks 


- Let  $b$  = number of blocks containing matching records

- $Cost = h_i * (t_T + t_S) + t_S + t_T * b$



## Selections Using Indices

---

- **A4 (secondary index, equality on nonkey).**
  - retrieve a single record if the search-key is a candidate key, e.g. *name* in *instructor* 
    - $cost = (h_i + 1) * (t_T + t_S)$
  - retrieve multiple records if search-key is not a candidate key
    - each of  $n$  matching records may be on a different block, e.g. *salary* in *instructor* - *next slide*
    - $cost = (h_i + n) * (t_T + t_S)$ 
      - can be very expensive!

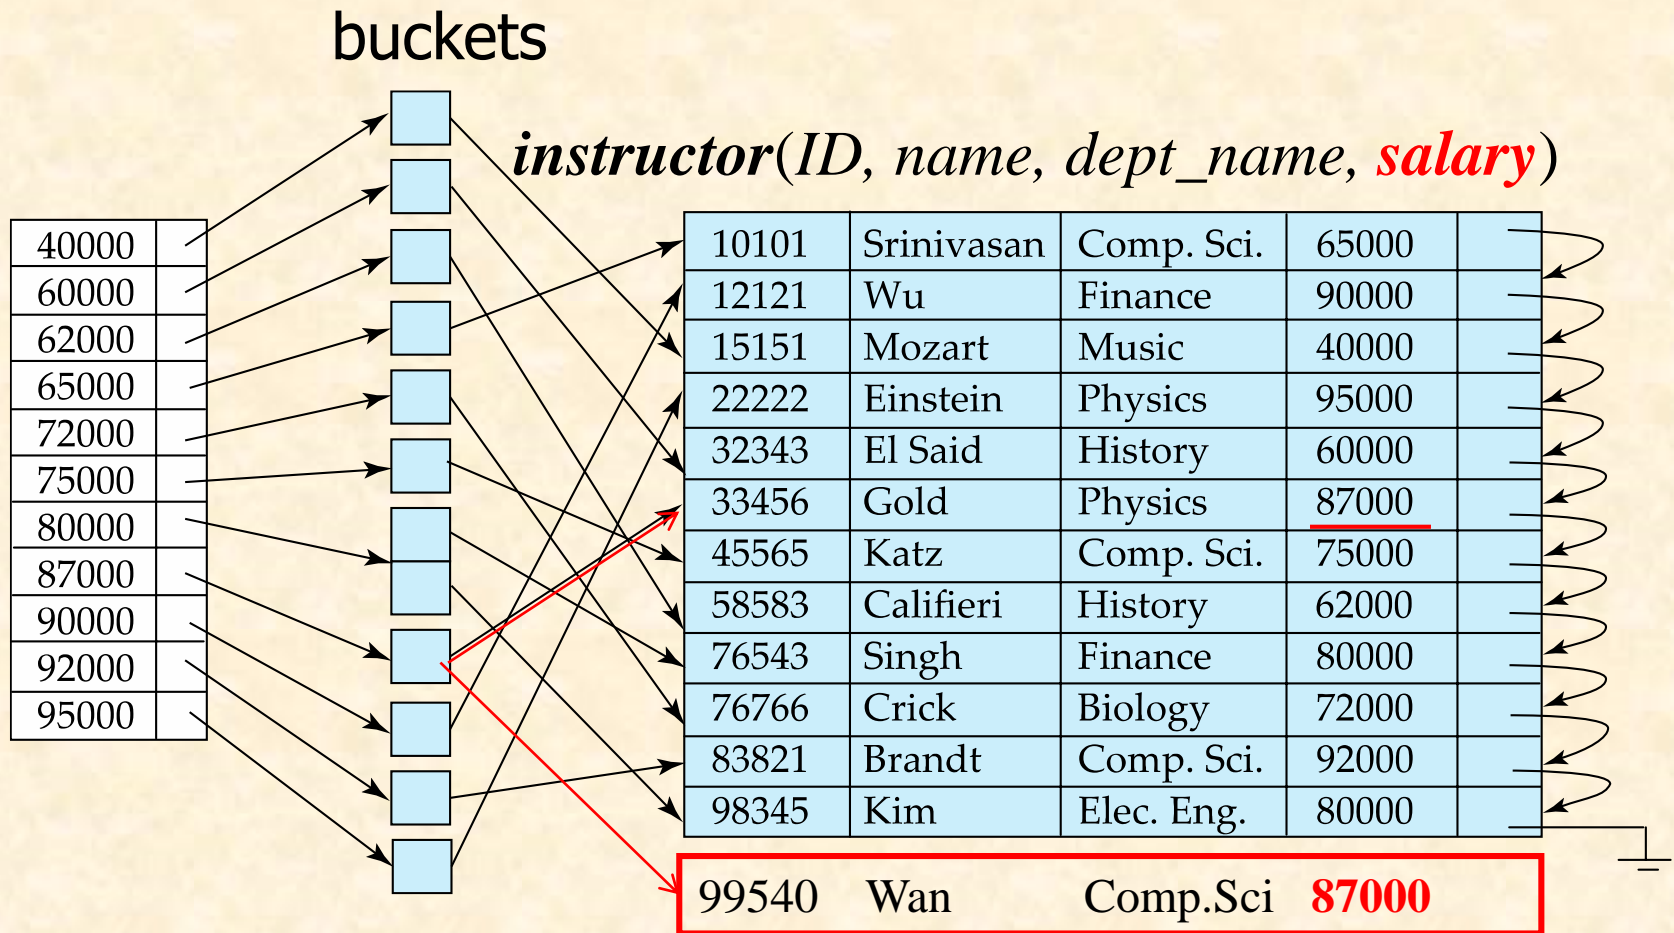
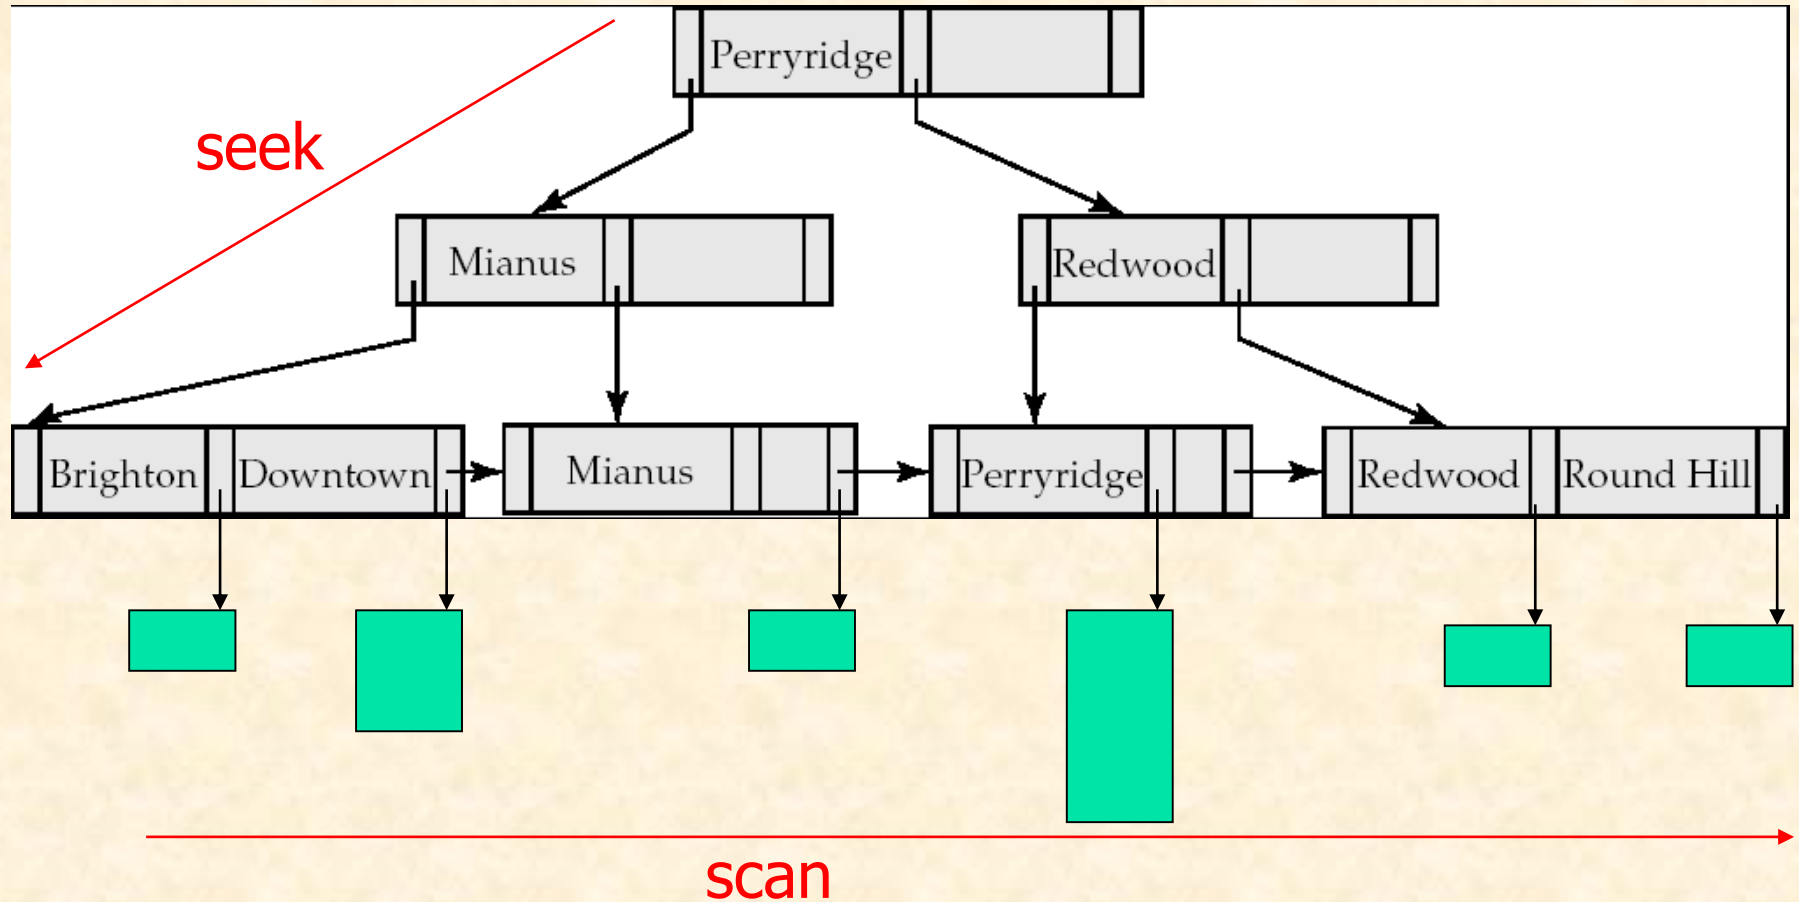


Fig. 11.6 Secondary index on *salary* field of *instructor*

# 聚集索引树：索引+数据

访问方式：

1. 聚集索引查找seek, e.g. branch-name, 二分查找
2. 聚集索引扫描scan, e.g. account-number, 线性扫描



未建立索引，堆文件  
访问方式：线性扫描、表扫描

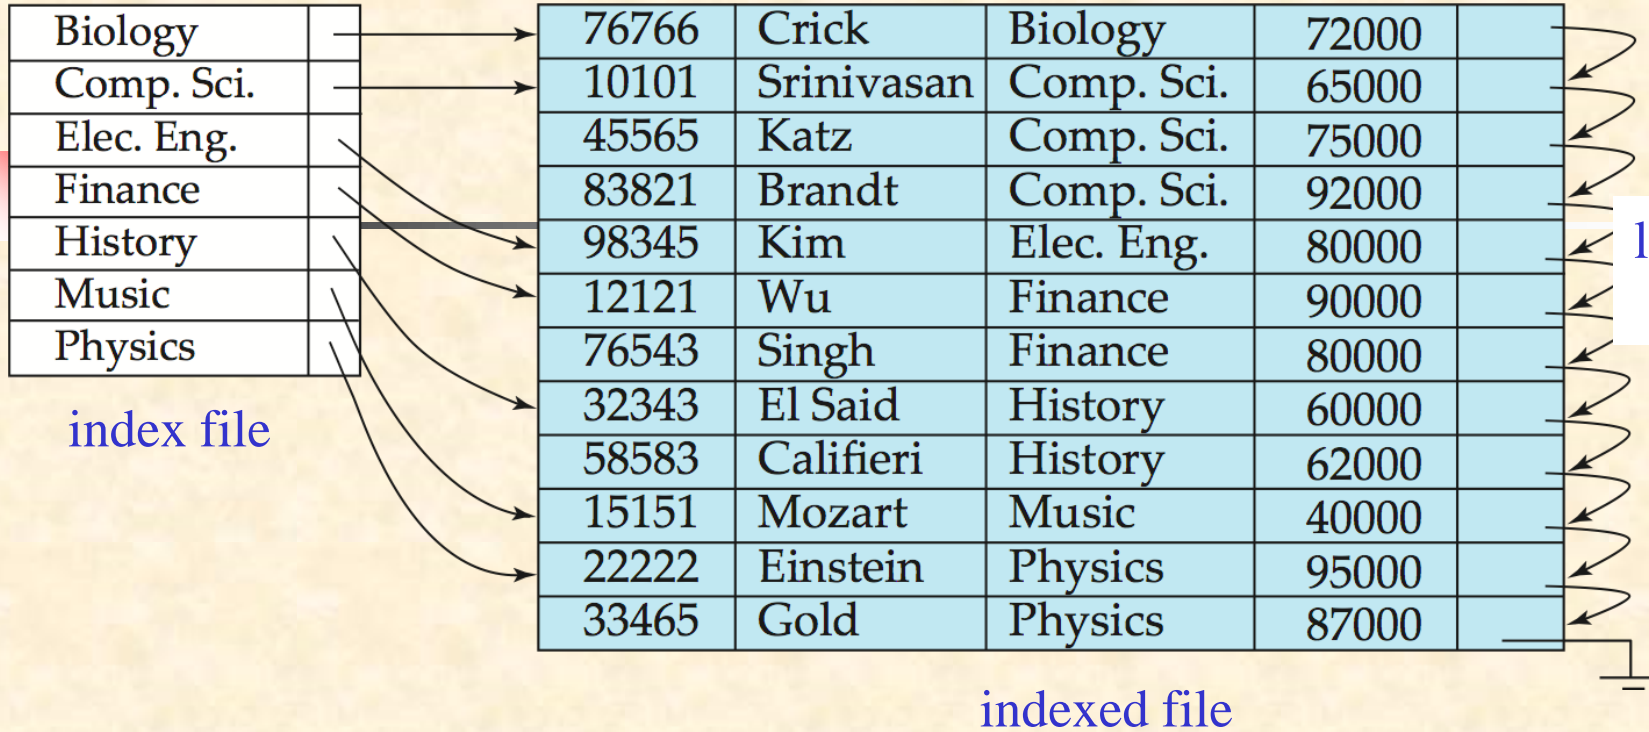
*Account(account number, **branch-name**, balance)*



A-217	Brighton	750	
A-101	Downtown	500	
A-110	Downtown	600	
A-215	Mianus	700	
A-102	Perryridge	400	
A-201	Perryridge	900	
A-218	Perryridge	700	
A-222	Redwood	700	
A-305	Round Hill	350	



*instructor*(ID, *name*, dept\_name, salary)



Note: the file *instructor* is logically a sequential file, but its records may be stored *non-contiguously* or *non-ordered* on the disk

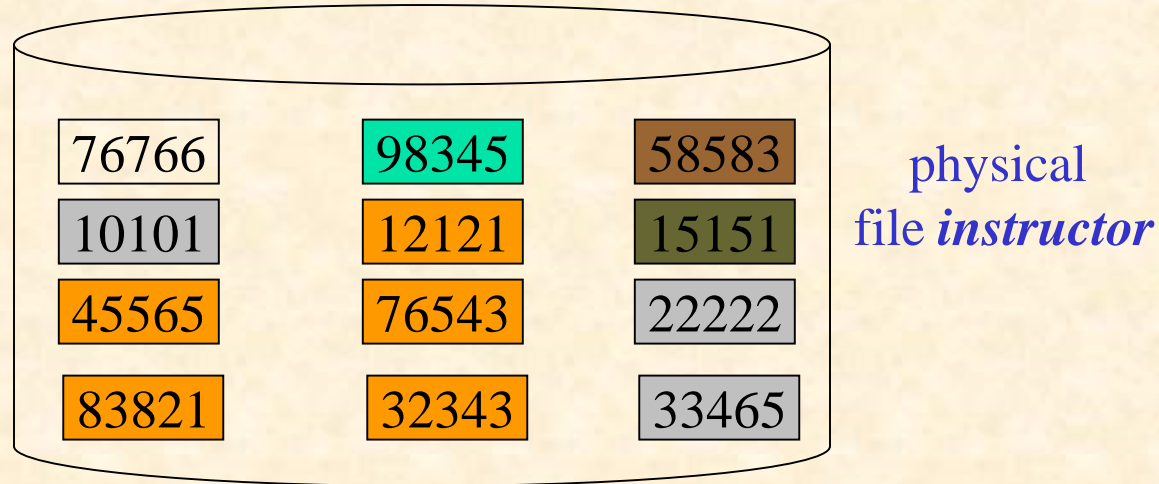


Fig. 11.1 DB *indexed* file *instructor* and its *index* file

# Selections Involving Comparisons

## - Range Search

- Can implement selections of the form  $\sigma_{A \leq V}(r)$  or  $\sigma_{A \geq V}(r)$  by using
  - a linear file scan,
  - or by using indices in the following ways:
- **A5 (primary index, comparison)**. (Relation is sorted on A)
  - For  $\sigma_{A \geq V}(r)$  use index to find first tuple  $\geq v$  and scan relation sequentially from there
  - For  $\sigma_{A \leq V}(r)$  just scan relation sequentially till first tuple  $> v$ ; do not use index
- e.g.  $\sigma_{ID \geq 22222}(instructor)$  in





## Selections Involving Comparisons

### ■ A6 (secondary index, comparison).

- For  $\sigma_{A \geq v}(r)$  use index to find first *index entry*  $\geq v$  and scan index sequentially from there, to *find pointers* to records.
- For  $\sigma_{A \leq v}(r)$  just scan leaf pages of index finding pointers to records, till first entry  $> v$
- In either case, retrieve records that are pointed to
  - requires an I/O for each record
  - Linear file scan may be cheaper

# Implementation of Complex Selections

## - Conjunction

- **Conjunction:**  $\sigma_{\theta_1 \wedge \theta_2 \wedge \dots \wedge \theta_n}(r)$
- **A7 (conjunctive selection using one index).**
  - Select a combination of  $\theta_i$  and algorithms A1 through A7 that results in the least cost for  $\sigma_{\theta_i}(r)$ .
  - Test other conditions on tuple after fetching it into memory buffer.
- **A8 (conjunctive selection using composite index).**
  - Use appropriate composite (multiple-key) index if available.

# Implementation of Complex Selections - Conjunction

- **A9 (conjunctive selection by intersection of identifiers).**
  - Requires indices with record pointers.
  - Use corresponding index for each condition, and take intersection of all the obtained sets of record pointers.
  - Then fetch records from file
  - If some conditions do not have appropriate indices, apply test in memory.



## Algorithms for Complex Selections

---

- **Disjunction:**  $\sigma_{\theta_1 \vee \theta_2 \vee \dots \vee \theta_n}(r)$ .
- **A10 (disjunctive selection by union of identifiers).**
  - Applicable if *all* conditions have available indices.
    - Otherwise use linear scan.
  - Use corresponding index for each condition, and take union of all the obtained sets of record pointers.
  - Then fetch records from file
- **Negation:**  $\sigma_{\neg\theta}(r)$ 
  - Use linear scan on file
  - If very few records satisfy  $\neg\theta$ , and an index is applicable to  $\theta$ 
    - Find satisfying records using index and fetch from file



## Other Operations

---

- *Sorting, join, projection, set operations* (e.g. union, intersection, set-difference), *aggregation*, and *duplicate elimination* are implemented on the basis of the operation *select*
  - sorting,  
e.g. merge sort
  - join,  
(block/indexed) nested-loop join, merge join, hash join
  - *project, distinct, order by, outer join, aggregation*



## § 15.7 Evaluation of Expressions

---

- How to evaluate an expression containing multiple evaluation primitives
  - key: how to process intermediate computing results
  - e.g.  $\Pi_{\text{salary}}(\sigma_{\text{salary} < 2500}(\text{instructor}))$
- Two strategies, that is **materialization**(实体化/物化) and **pipeline** are used for expression evaluation

# Materialization

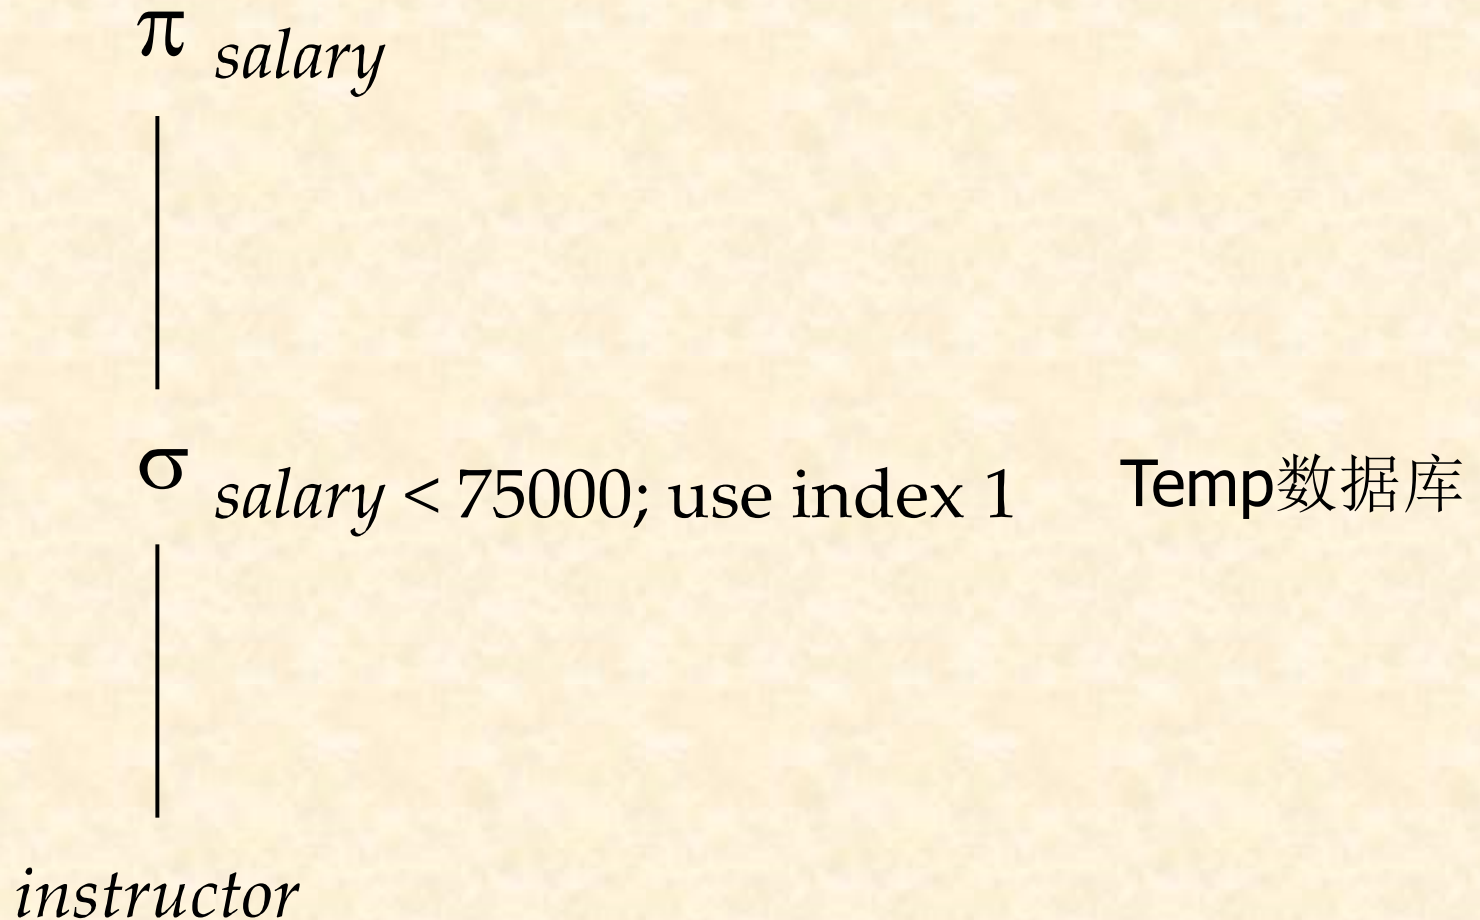


Fig. 12.0.2 Pictorial representation of an expression





## Materialization (cont.)

- Principles
  - start from the lowest-level, i.e. at the bottom of the tree, evaluate one operation at a time
  - the results of each evaluation (i.e. intermediate computing results) are stored in *temporal relations* on the disk for subsequent evaluation
    - serial evaluating
- E.g., in Fig.12.0.2, compute and store in a *temporal relation*

$$\sigma_{\text{salary} < 2500}(\text{instructor})$$

at first;





## Materialization (cont.)

---

- Cost of *writing results to disk and reading them back* can be quite high
- Approximately,  
overall cost =  
    sum of costs of individual operations +  
    cost of writing intermediate results to disk

## Pipelined (流水线)

### ■ Principles of pipelined evaluation

- evaluate several operations simultaneously in a pipeline, with the results of one operation passed to the next, without the need to store temporary relations in disk
  - parallel evaluating
- Much cheaper than materialization:
  - no need to store a temporary relation to disk
- E.g., in the expression tree in Fig.12.0.2, don't store result of

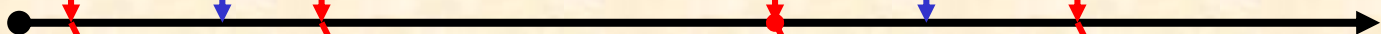
$$\sigma_{salary < 2500} (instructor)$$

- don't store result of selection, pass tuples directly to projection.

*instructor*




$\sigma_{\text{salary} < 2500}$




$\Pi_{\text{salary}} :$



$\Pi_{\text{salary}}(\sigma_{\text{salary} < 2500})$

  $\text{salary} < 2500$

  $\text{salary} \geq 2500$