

16.4 Choice of Evaluation Plans

— Query optimization

- How to generate the *evaluation plan* for a relational algebra expression with lower or the lowest cost
 - generating of the *optimized query trees*
 - selecting
 - implementation algorithms for each operations in the tree
 - *pipeline* or *materialization* strategy for the expression
- E.g. Fig.16.2

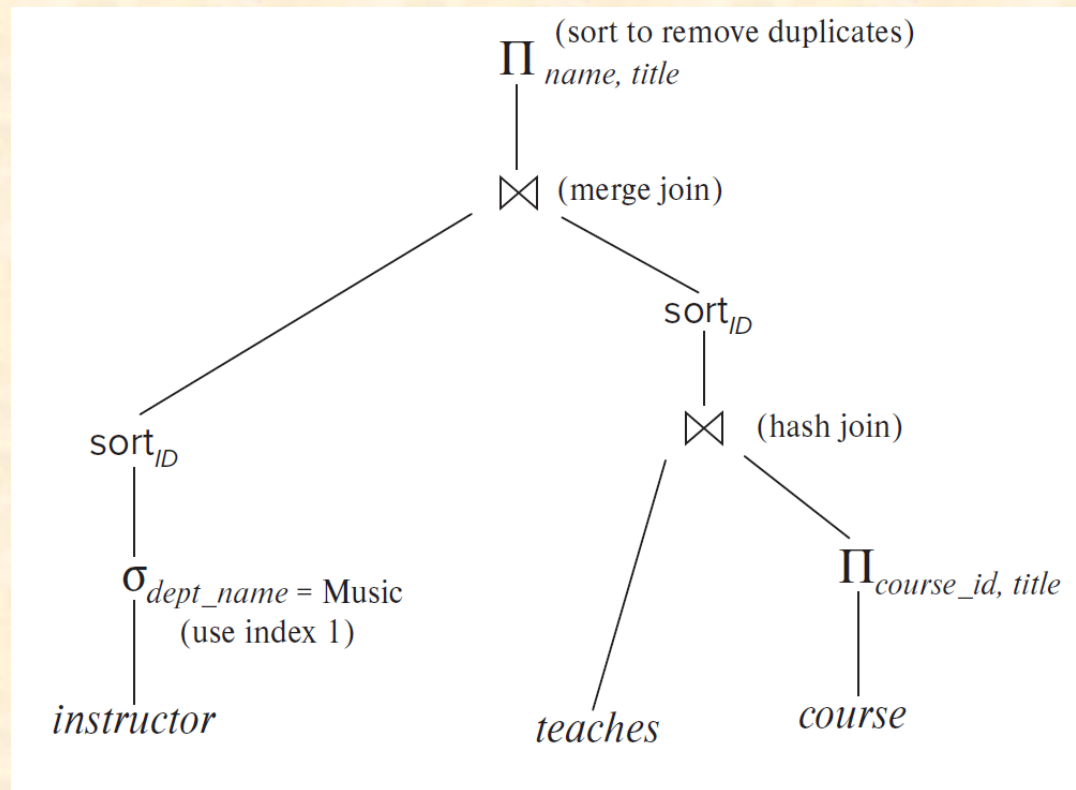


Fig.16.2 An evaluation plan



Choice of Evaluation Plans

- Must consider the interaction of evaluation techniques when choosing evaluation plans
 - choosing the cheapest algorithm for each operation independently may not yield best overall algorithm. E.g.
 - merge-join may be costlier than hash-join, but may provide a sorted output which reduces the cost for an outer level aggregation.
 - nested-loop join may provide opportunity for pipelining



Choice of Evaluation Plans

- **Query optimization**

- choosing the evaluation plan with *lowest* or the *lower cost*

- Practical query optimizers incorporate elements of the following two broad approaches:

1. *Cost-based* (16.4.1, 16.4.2)

Search all the plans and choose the best plan in a cost-based fashion.

2. *heuristic optimization* (16.4.3)

Uses heuristics to choose a plan

- Practical query optimizers incorporate elements of these two approaches



16.4.1 Cost-Based Join Order Selection

- The *join* operator is the basis for multiple table query, and its cost is expensive
- For a relational algebra expression with several *join* operation, the join order influences the query cost of the expression.
- Consider finding the best join-order for $r_1 \bowtie r_2 \bowtie \dots r_n$.
- There are $(2(n-1))!/(n-1)!$ different join orders for above expression. With $n = 7$, the number is 665280, with $n = 10$, the number is greater than 176 billion!
- No need to generate all the join orders. Using dynamic programming, the least-cost join order for any subset of $\{r_1, r_2, \dots, r_n\}$ is computed only once and stored for future use.

Join Order Selection by Dynamic Programming

- No need to generate all the join orders. Using dynamic programming, the least-cost join order for any subset of $\{r_1, r_2, \dots, r_n\}$ is computed only once and stored for future use
- 方法:
P768, Fig.16.7;
- 类似于《算法设计与分析》中的矩阵连乘积

$$\begin{array}{lll} (A((BC)D)) & (A(B(CD))) & ((AB)(CD)) \\ (((AB)C)D) & ((A(BC))D) & \end{array}$$

16.4.2 Cost-based Optimization with Equivalence Rules

- Cost-based optimization
 - search *all??* the possible plans and choose *the best* plan
 - a *optimal* plan with the minimum costs can be obtained
- Cost-based optimization is expensive
- **Physical equivalence rules** allow logical query plan to be converted to physical query plan specifying what algorithms are used for each operation.

Cost-based Optimization with Equivalence Rules

- Efficient optimizer based on equivalent rules depends on
 - A space efficient representation of expressions which avoids making multiple copies of subexpressions
 - Efficient techniques for detecting duplicate derivations of expressions
 - A form of dynamic programming based on **memoization**, which stores the best plan for a subexpression the first time it is optimized, and reuses in on repeated optimization calls on same subexpression
 - Cost-based pruning techniques that avoid generating all plans
- Pioneered by the Volcano project and implemented in the **SQL Server** optimizer



16.4.3 Heuristic Optimization

■ Heuristic optimization

- uses *heuristics* (or *heuristic rules*) to choose *a better* plan
- to reduce the number of choices that must be made in a cost-based fashion
- a *suboptimal* plan with lower costs can be obtained

■ Principles

transforms the query-tree by heuristics to reduce cost

- perform *selection* early to reduce the number of tuples
- perform *projection* early to reduces the number of attributes
- substitute *Cartesian product* and *selection* with *join*





Heuristic Optimization (cont.)

- perform *most restrictive selection* and *join* operations before other similar operations
 - the *most restrictive* operations generate resulting relations with smallest size
- Heuristic optimization used in some versions of **Oracle**

Steps in Heuristic Optimization

- Step1 使用rule1(▶), 将conjunctive selection 分解为多个单独的**选择操作**, 以使单个选择操作尽可能沿查询树下移 (尽早执行选择操作, 以减少中间计算结果)
- Step2 根据**选择操作**的交换率和分配率, 利用rule2, rule7.a, rule7.b, rule11, 将查询树上的每个选择操作尽可能移向叶节点, 以便尽早执行选择操作
- **Step3** 根据**连接操作**的结合律和交换率, 使用rule6(▶), 重新安排查询树中的叶结点, 使得具有***restrictive selection***特征的叶结点先执行
 - ***restrictive selection***: 执行此操作后, 产生的结果关系最小 (所含元组最少)

Steps in Heuristic Optimization (cont.)

- Step4 利用rule4.a(), 以 **连接操作** 代替相邻的选择和笛卡尔乘积操作
- Step5 利用rule3, 8.a, 8.b,12, 将查询树上的 **投影操作** 尽可能下移, 以便尽早执行投影操作, 减少中间计算结果
- Step6 将最后的查询树分解为多个子树, 使子树中的各操作可以采用 **流水线方式** 执行, 以减少对外设的访问次数
 - e.g. Fig.16.2 

重点: 前5步 !!!

Example One

- Given

- S(S#, sname, age, sex)
- C(C#, cname, teacher)
- SC(S#, C#, grade)

- Find all the *students*, who are *male*, and get *grades* more than 90 when they learn some one *course*

- Step1. SQL statement

- ```
SELECT DISTINCT sname
FROM S, SC
WHERE S.s# = SC.s# AND sex=M AND grade > 90
```

selection conditions

join condition



## Example One (cont.)

---

### ■ Principles

select  $A1, A2, \dots, An$   
from  $r_1, r_2, \dots, r_n$   
where  $P$

corresponds to

$$\Pi_{A1, A2, \dots, An} (\sigma_P (r_1 \times r_2 \times \dots \times r_n))$$



## Example One (cont.)

- Step2. Initial Relational algebra expression

- E1 =

$$\Pi_{\text{sname}}(\sigma_{\text{s.s\#}=\text{sc.s\#} \wedge \text{sex}=\text{M} \wedge \text{grade}>90} (S \times SC))$$


或:  $\Pi_{\text{sname}}(\sigma_{\text{sex}=\text{M} \wedge \text{grade}>90} (S \bowtie SC))$

- Step3. Initial query tree

- Fig. 16.0.5 (a), E1

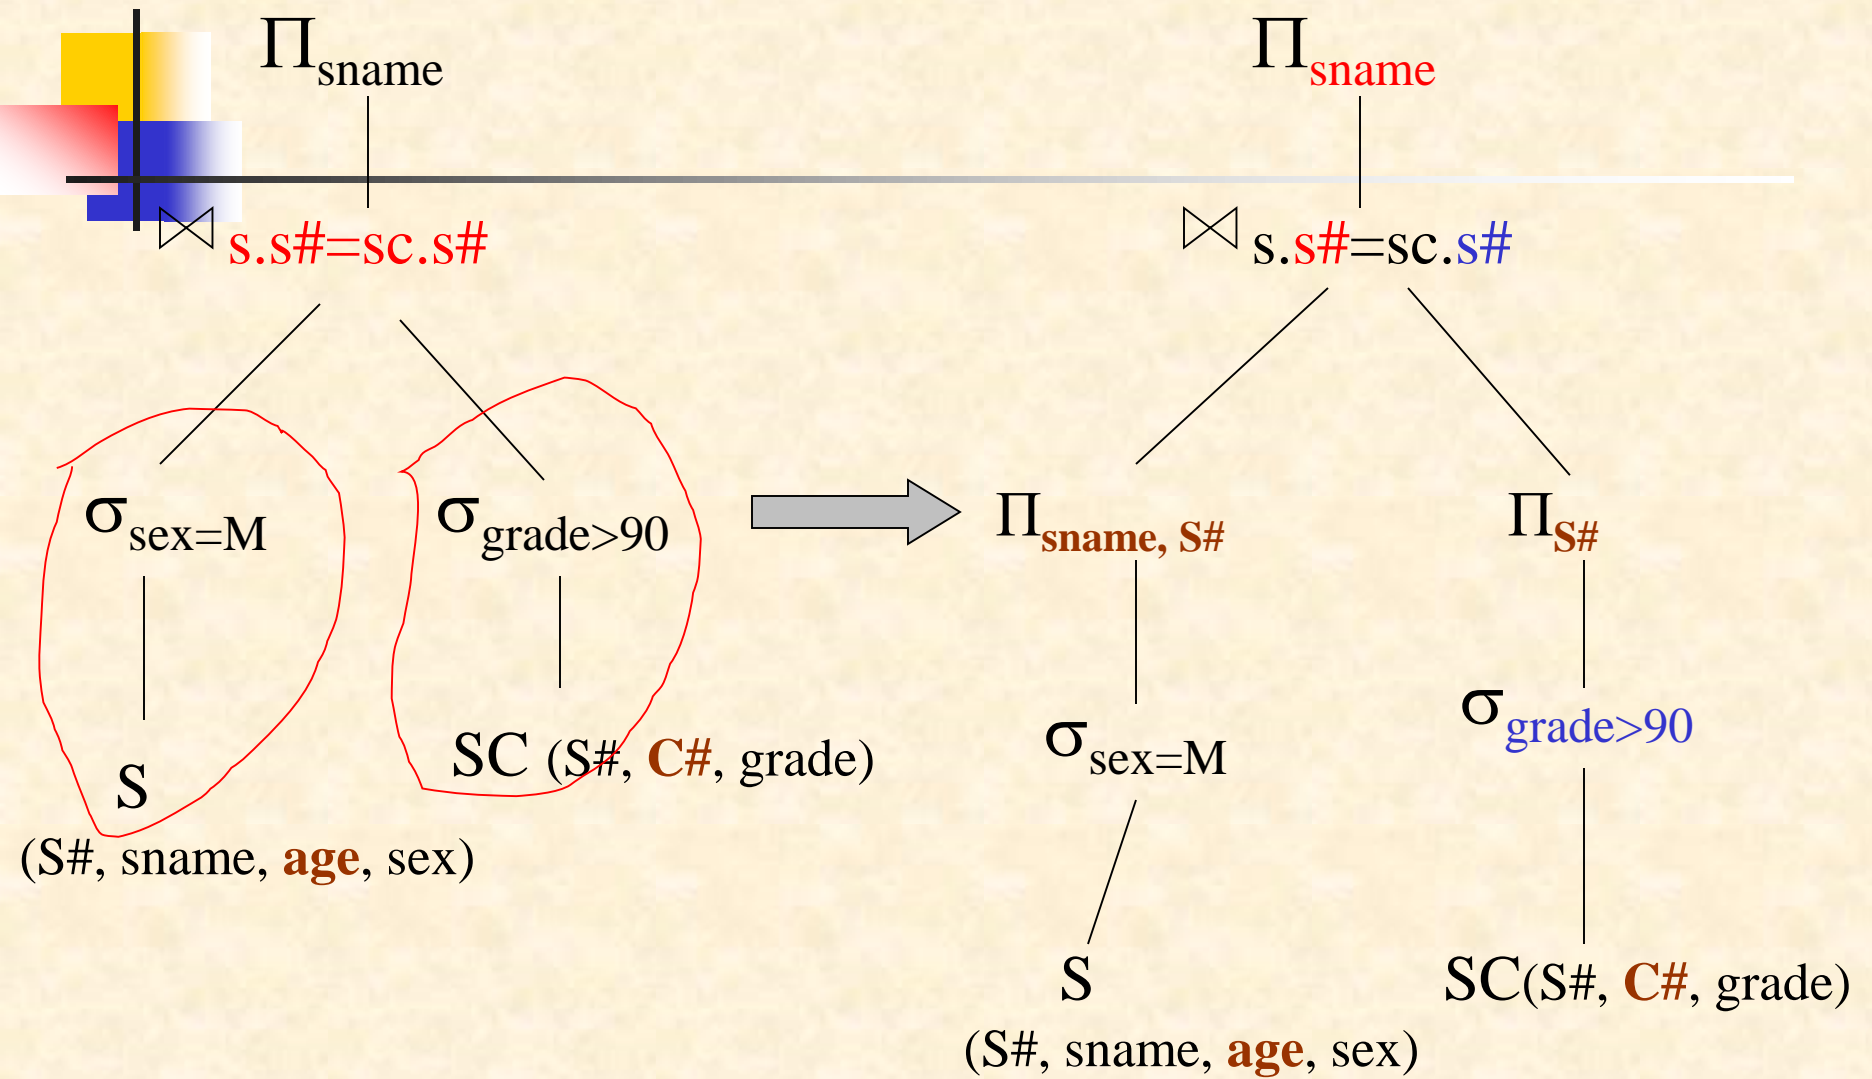
- Step4. By means of *Rule1*, *Rule7b*, distribute  $\sigma$  operations over relation  $S$  and  $SC$

- Fig. 16.0.5 (b), E2



## Example One (cont.)

- Step5. By means of *Rule4.a*, replace *selection* and *Cartesian product* with *natural join*
  - Fig. 16.0.5 (c), E3
- Step6. !! By means of *Rule8b*, distribute  $\Pi$  operations over relation *S* and *SC*
  - Fig. 16.0.5 (d),
- The final optimized expression E4 that is equivalent to initial expression E1 is
  - $$\Pi_{\text{sname}} \{ \Pi_{\text{sname}, S\#} (\sigma_{\text{sex}=\text{M}} (S)) \bowtie \sigma_{\text{grade}>90} (\Pi_{\text{grade}, S\#} (SC)) \}$$



(c) E3

(d) E4



## Example Two

---

- Consider the following insurance database, where the primary keys are underlined,
  - ***Person***(driver-id, name, address)
  - ***Car***(license, model, year)
  - ***Accident***(report-number, date, location)
  - ***Participated***(driver-id , license, report-number, damage-amount)



## Example Two (cont.)

---

### ■ Problems

- Give a SQL statement to find out the ***driver name***, ***license***, ***report-number*** of the ***accidents*** which happened in *Beijing* and *before May 3, 2008*
- Give the initial query tree for this query, and construct an optimized and equivalent *relational algebra expression* for it by means of heuristic optimization



## Example Two (cont.)

- Step1. SQL statement

- **Select** *name, license, Accident.report-number*

- From** *Person, Participate, Accident*

- Where** *Person.driver\_id = Participate.driver\_id AND*

- Accident.report\_number = Participate.report\_number*


- AND** *Location=Beijing AND date < 5/3/2008*

- Step2. Initial expression

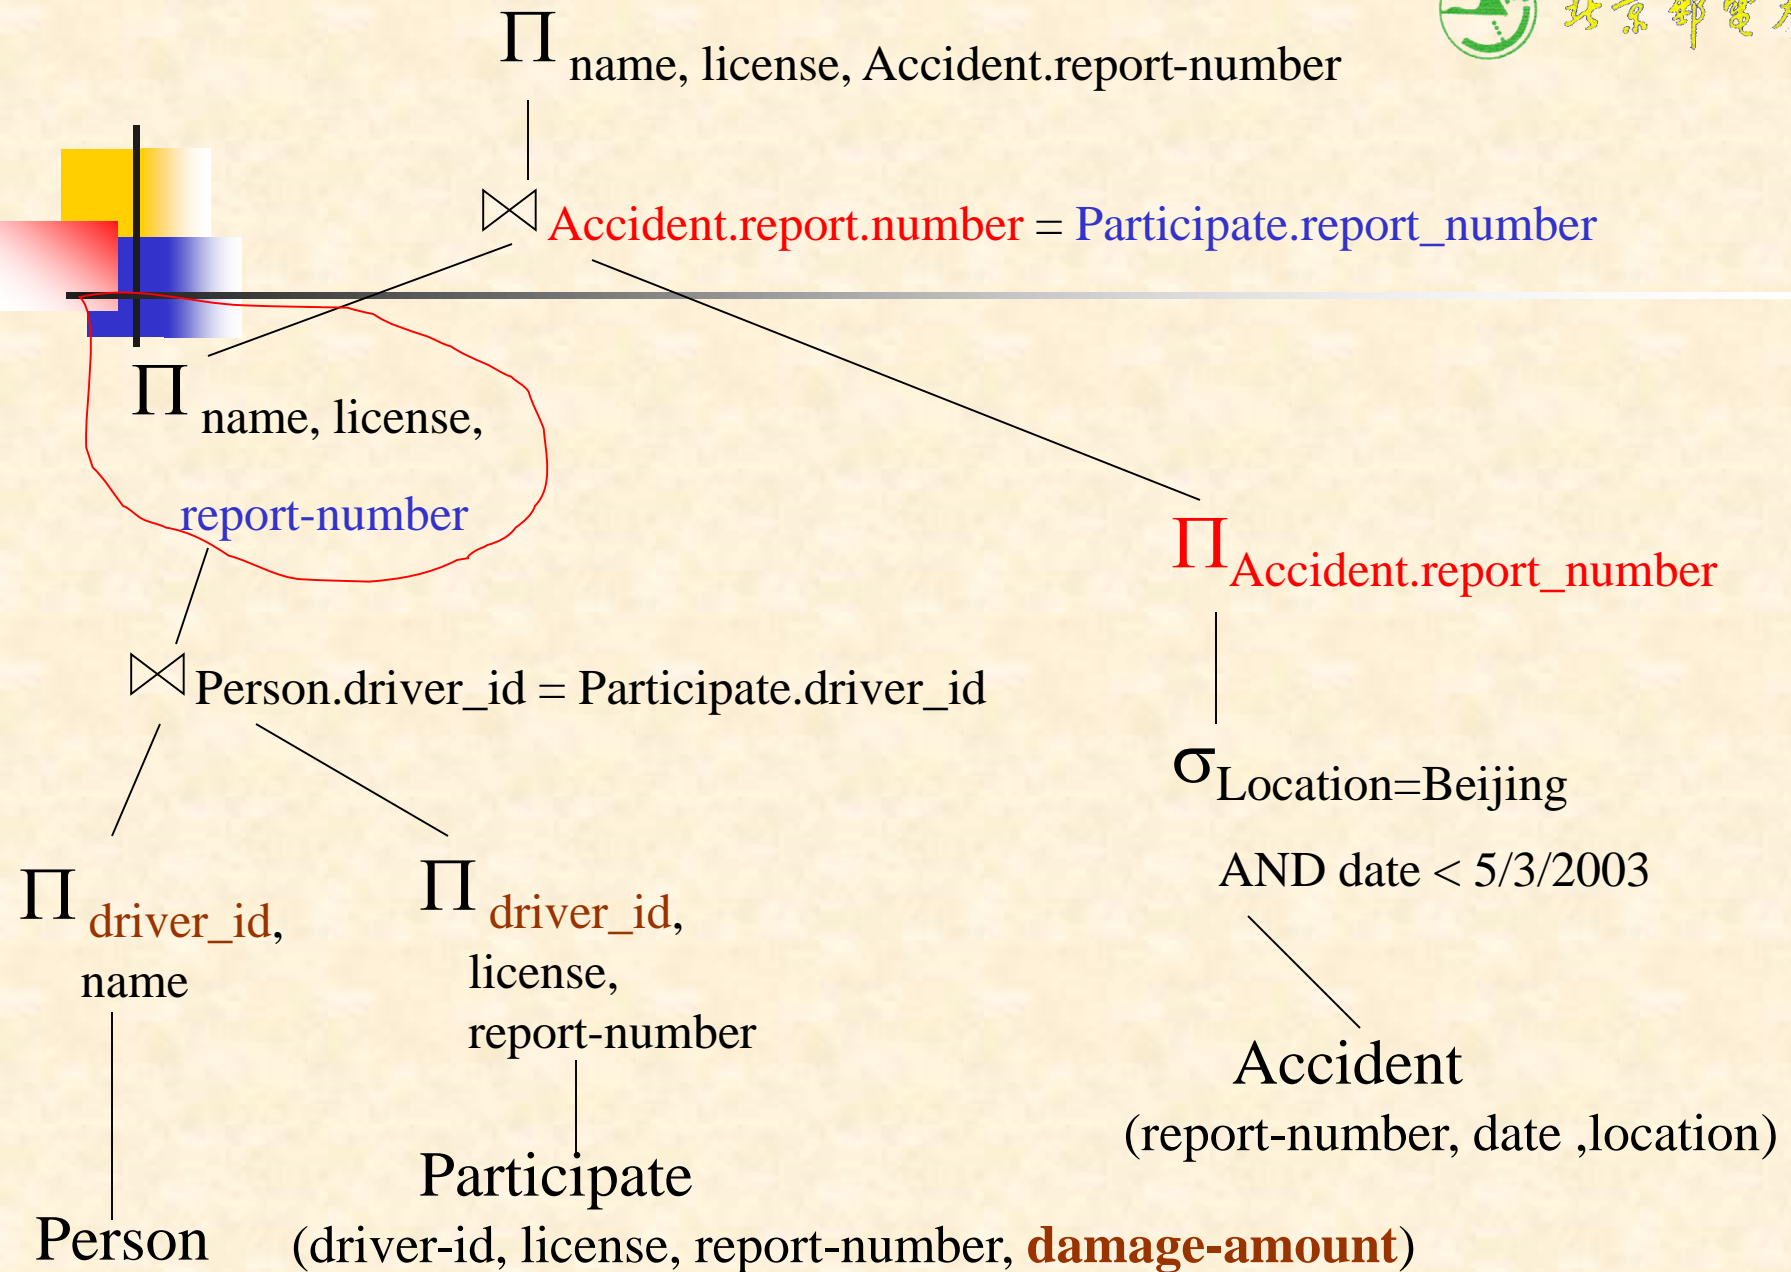
$\Pi_{\text{name, license, Accident.report-number}}$

$(\sigma_{\text{Person.driver\_id = Participate.driver\_id AND Accident.report\_number = Participate.report\_number AND Location=Beijing AND date < 5/3/2003}}$

$(\text{Person} \times \text{Participate} \times \text{Accident})$


$$\Pi_{\text{name, license, Accident.report-number}}$$
$$\sigma_{\text{Person.driver\_id = Participate.driver\_id AND}} \\ \text{Accident.report\_number = Participate.report\_number AND} \\ \text{Location=Beijing AND date < 5/3/2003}$$


(a) initial query tree



(b) optimal query tree



## Example Three

---

- Consider the following relations in banking enterprise database, where the primary keys are underlined
  - *branch* (*branch-name*, *branch-city*, *assets*),
  - *loan* ( *loan-number*, *branch-name* , *amount*)
  - *borrower*( *customer-name*, *loan-number*, *borrow-date*)
  - *customer* (*customer-name*, *customer-street*, *customer-city*)
  - *account* (*account-number*, *branch-name*, *balance*)
  - *depositor* (*customer-name*, *account-number* , *deposit-date*)



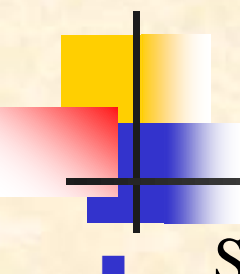
## Example Three (cont.)

---

- For the query “ Find the *names* of all *customers* who have an *loan* at any *branch* that is located in *Brooklyn* and have *assets* more than \$100,000, requiring that *loan-amount* is less than \$1000”
  - give an SQL statement for this query
  - given a initial query tree for the query, and convert it into an optimized query tree by means of heuristic optimization

## Example Three (cont.)



---



SQL

```
select customer-name
from borrower, loan, branch
where loan.loan-number=borrower.loan-number
and branch.branch-name=loan.branch-name
and branch-city= "Brooklyn"
and assets>100000 and amount<1000
```

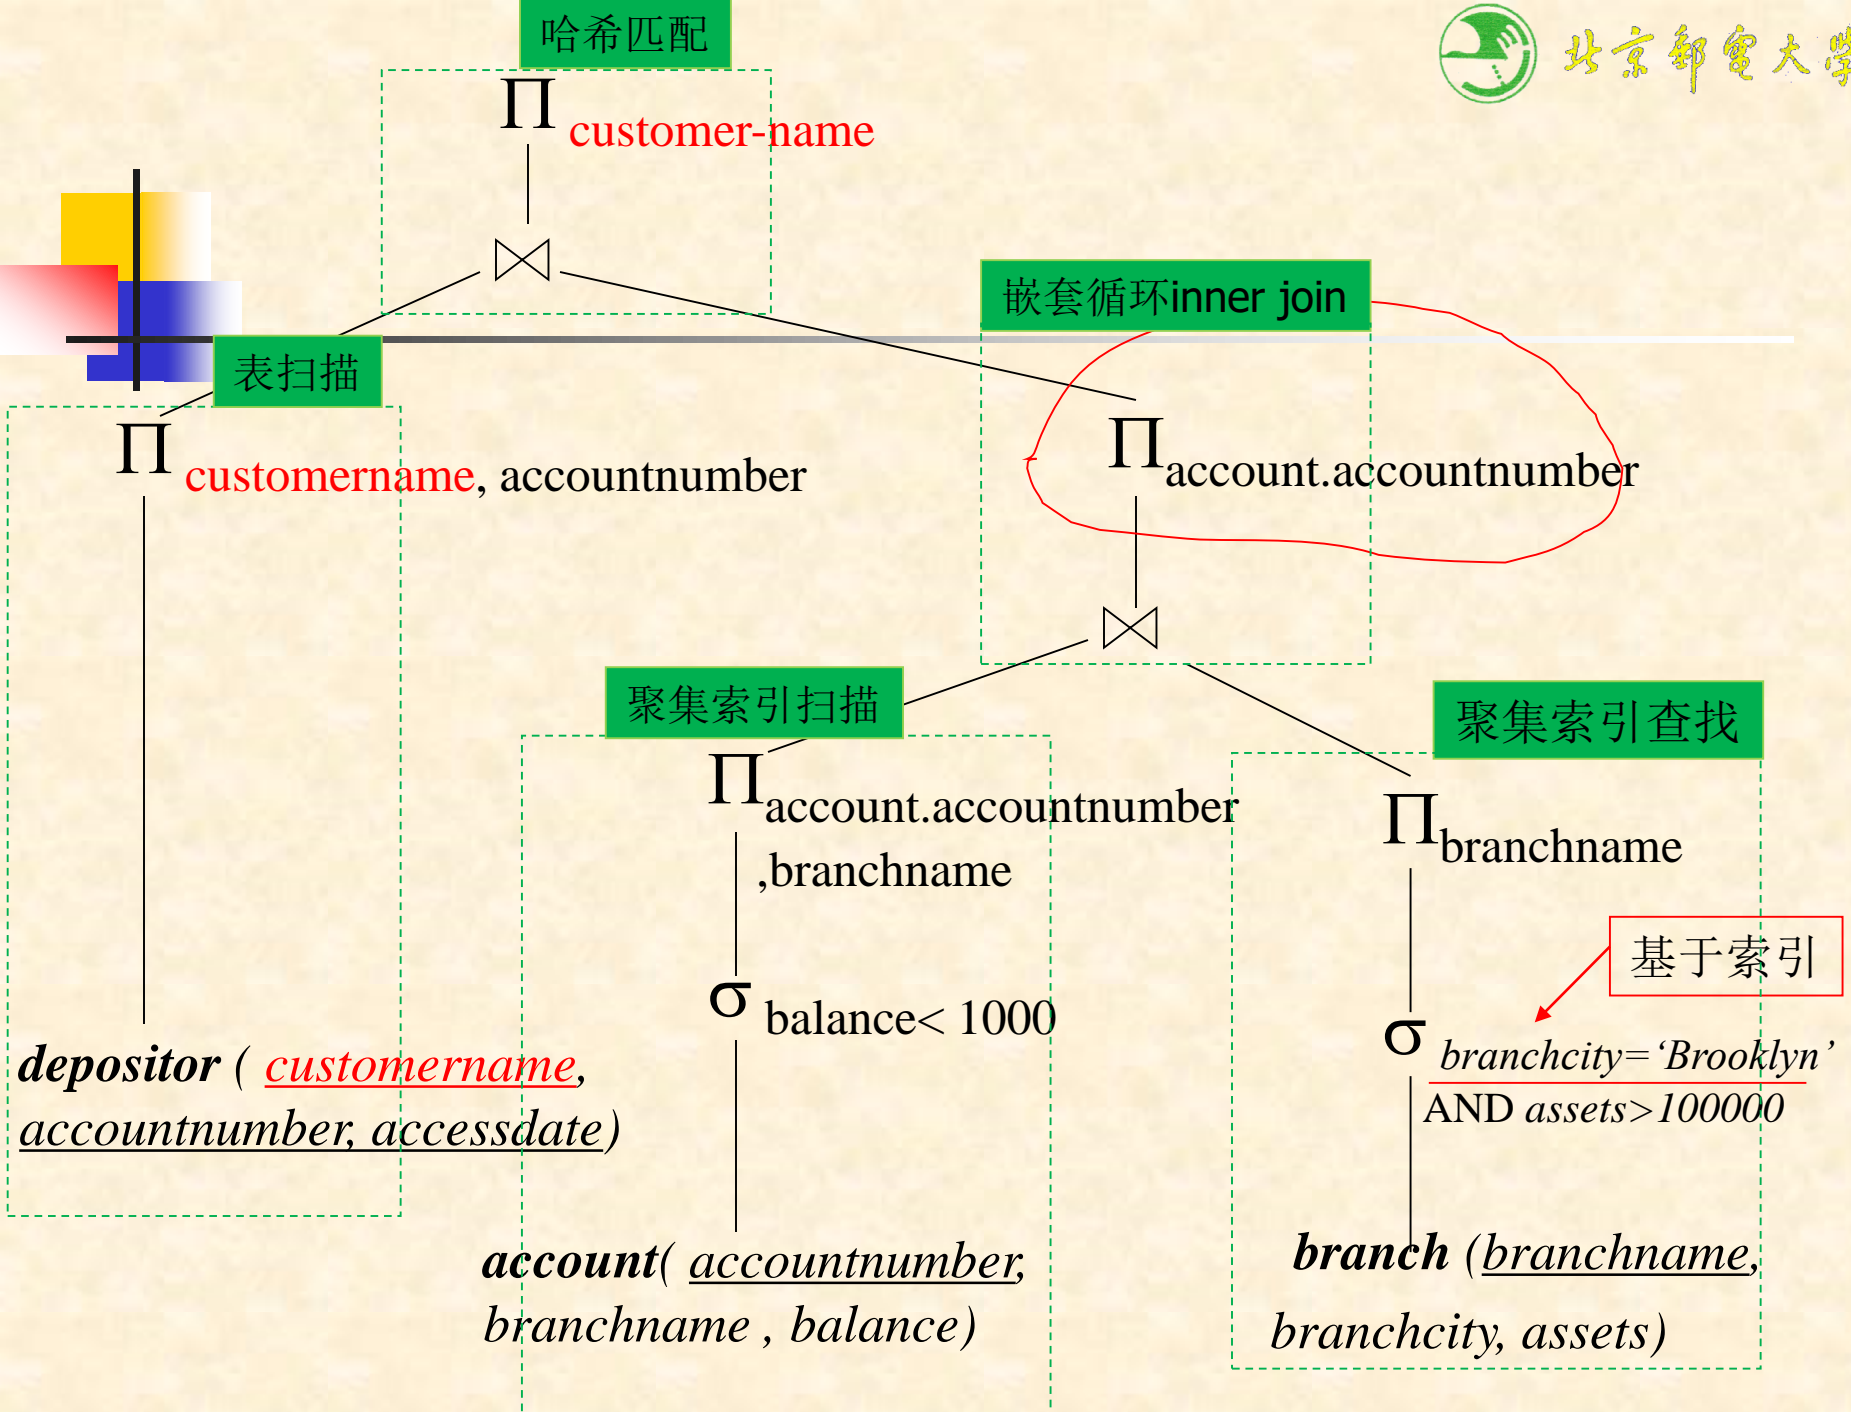




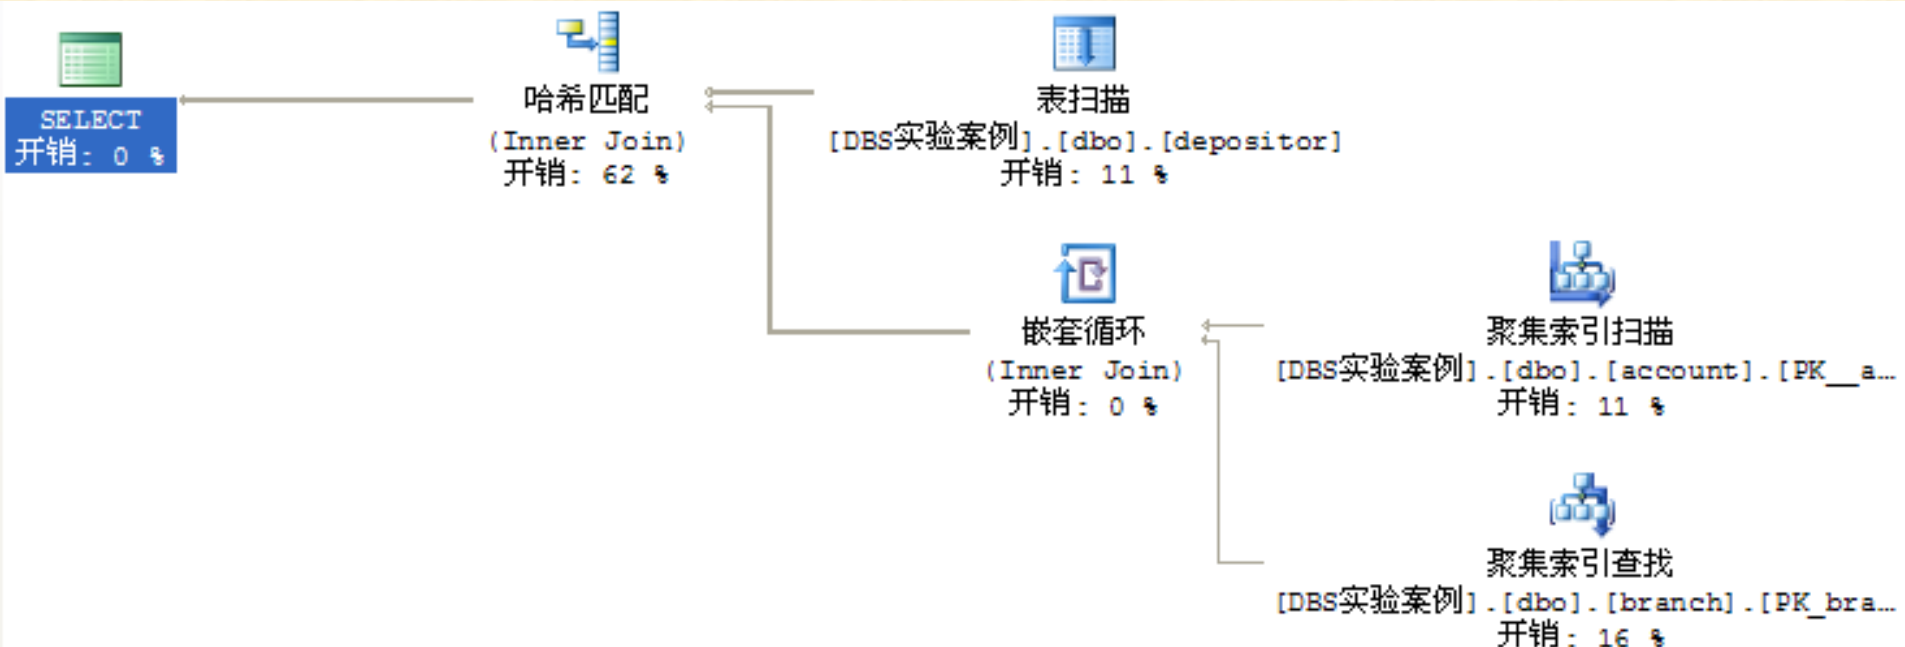


## 实际DBS中的查询优化

- 教科书上的选择、投影操作下移的实现方式  
实际平台下（如**SQL Server**）与后面的连接操作结合在一起做，避免多次关系代数操作：
  1. 采用**pipeline**，在对参与连接运算的关系扫描过程中完成符合条件的元组—选择
  2. 再选出后续步骤需要的属性—投影



```
select customername
from account, depositor, branch
where account.accountnumber=depositor.accountnumber
and branch.branchname=account.branchname
and branchcity='Brooklyn'
and assets>100000 and balance<1000
```



where depositor.accountnumber  
= account.branchname  
'yn'  
and balance<1000

100%

account, depositor, branch where account.accountnumber =

branch.branchname

配  
(join)  
2 %

表扫描  
[DBS实验案例].[dbo].[depositor]  
开销: 11 %

嵌套循环  
(Inner Join)  
开销: 0 %

聚集索引扫描  
[DBS实验案例].[dbo].[account]  
开销: 0 %

谓词  
[DBS实验案例].[dbo].[account].[balance]<  
(1000.)  
对象  
[DBS实验案例].[dbo].[account].  
[PK\_\_account\_\_07F6335A]  
输出列表  
[DBS实验案例].[dbo].  
[account].accountnumber, [DBS实验案例].  
[dbo].[account].branchname

### 聚集索引扫描

整体扫描聚集索引或只扫描一定范围。

#### 物理运算

#### 聚集索引扫描

| Logical Operation | Clustered Index Scan |
|-------------------|----------------------|
| 估计 I/O 开销         | 0.003125             |
| 估计 CPU 开销         | 0.000168             |
| 估计运算符开销           | 0.003293 (11%)       |
| 估计子树大小            | 0.003293             |
| 估计行数              | 10                   |
| 估计行大小             | 37 字节                |
| 已排序               | False                |
| 节点 ID             | 3                    |

#### 谓词

[DBS实验案例].[dbo].[account].[balance]<  
(1000.)

#### 对象

[DBS实验案例].[dbo].[account].  
[PK\_\_account\_\_07F6335A]

#### 输出列表

[DBS实验案例].[dbo].  
[account].accountnumber, [DBS实验案例].  
[dbo].[account].branchname

选择下  
移

投影下  
移

r, branch where account.accountnum



表扫描

实验案例].[dbo].[depositor]

开销: 11 %



嵌套循环

(Inner Join)

开销: 0 %

[DBS实验案例].[dbo].[a

聚集索引查

开销: 11



聚集索引查

[DBS实验案例].[dbo].[B

开销: 16

## 聚集索引查找

扫描聚集索引中特定范围的行。

### 物理运算

### 聚集索引查找

| Logical Operation | Clustered Index Seek |
|-------------------|----------------------|
| 估计 I/O 开销         | 0.003125             |
| 估计 CPU 开销         | 0.0001581            |
| 估计运算符开销           | 0.004706 (16%)       |
| 估计子树大小            | 0.004706             |
| 估计行数              | 1                    |
| 估计行大小             | 29 字节                |
| 已排序               | True                 |
| 节点 ID             | 4                    |

### 谓词

[DBS实验案例].[dbo].[branch].[assets]>  
(100000.) AND [DBS实验案例].[dbo].  
[branch].[branchcity]='Brooklyn'

### 对象

[DBS实验案例].[dbo].[branch].[PK\_branch]

### 输出列表

[DBS实验案例].[dbo].[branch].branchcity,  
[DBS实验案例].[dbo].[branch].assets

### Seek 谓词

前缀: [DBS实验案例].[dbo].  
[branch].branchname=[DBS实验案例].  
[dbo].[account].[branchname]


选择下  
移

投影下  
移



查询开销: 100%

m account, depositor, branch where account.accountnumber=depositor.accountnumber and



哈希匹配  
Inner Join)  
开销: 62 %




表  
[DBS实验案例]. [dbo].  
开销

**嵌套循环**

对于顶部(外部)输入的每一行, 扫描底部(内部)输入, 然后输出匹配的行。

| 物理运算              | 嵌套循环           |
|-------------------|----------------|
| Logical Operation | Inner Join     |
| 估计 I/O 开销         | 0              |
| 估计 CPU 开销         | 0.0000418      |
| 估计运算符开销           | 0.0000554 (0%) |
| 估计子树大小            | 0.0080544      |
| 估计行数              | 10             |
| 估计行大小             | 16 字节          |
| 节点 ID             | 2              |

**输出列表**

[DBS实验案例]. [dbo].  
[account]. accountnumber

**外部引用**

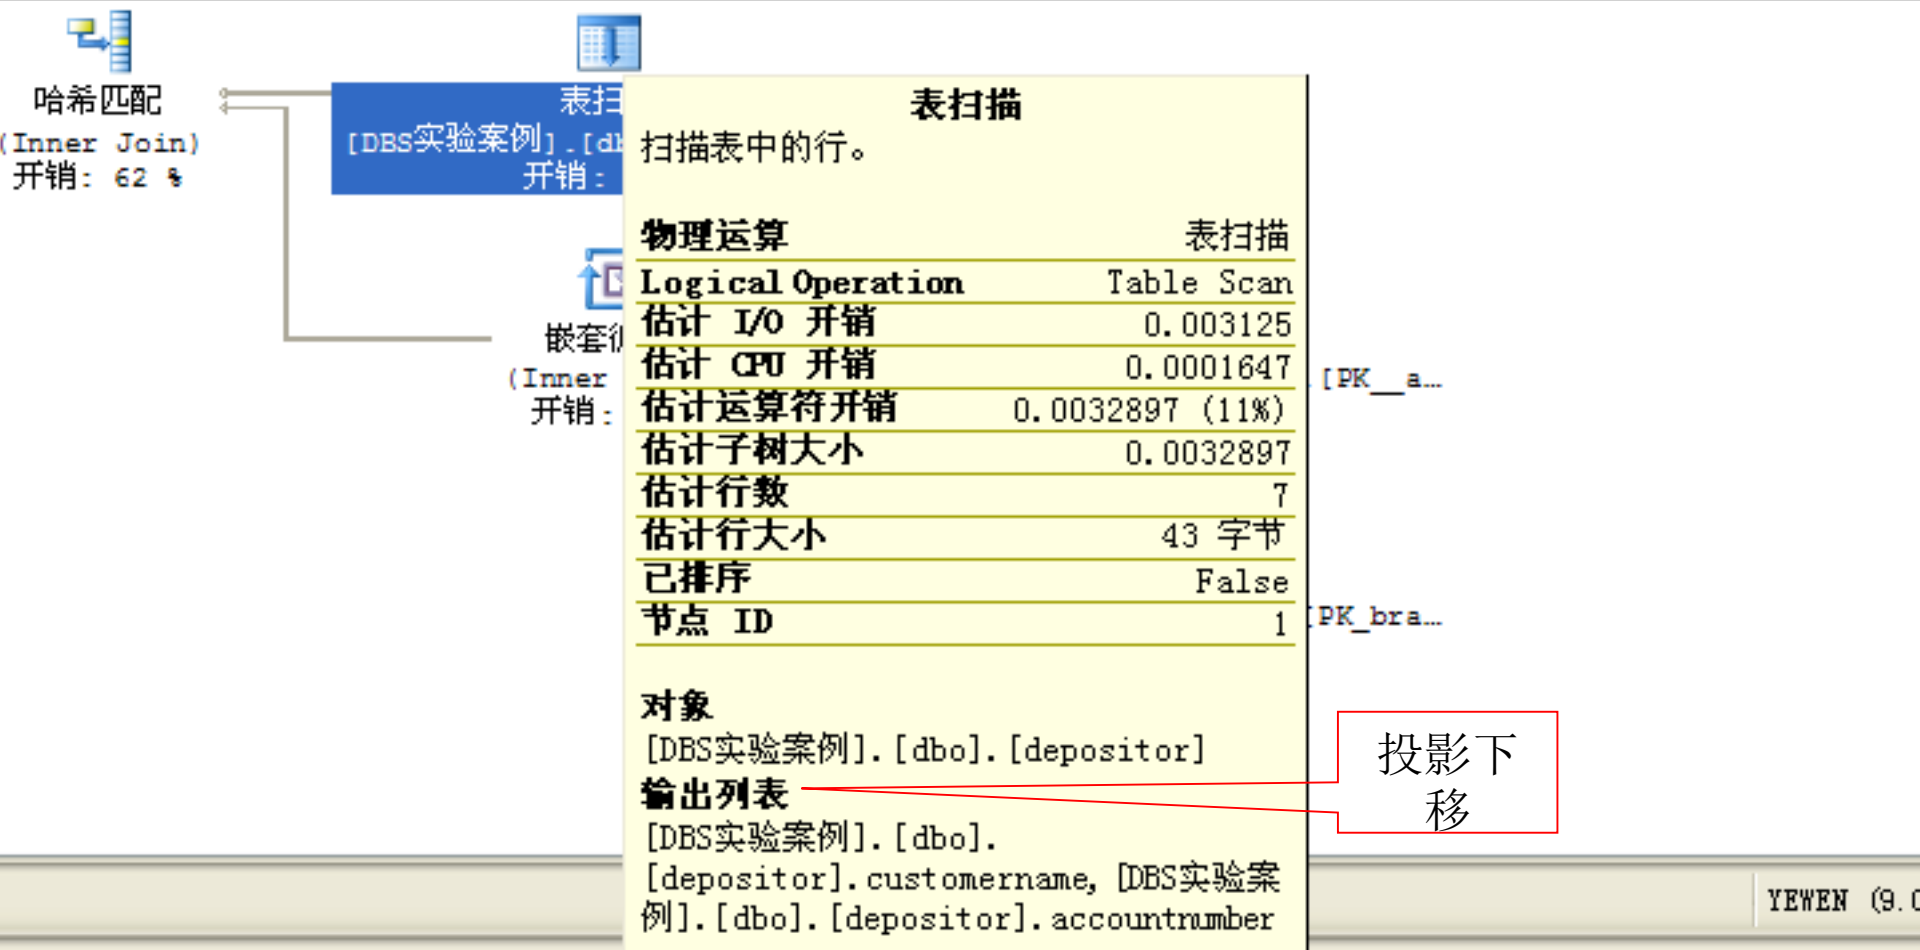
[DBS实验案例]. [dbo].  
[account]. branchname

连接实现方式

投影下移

查询开销: 100%

from account, depositor, branch where account.accountnumber=depositor.accountnumber and br



执行计划

(与该批有关的) 查询开销:

customername from acco

哈希匹配  
(Inner Jo  
开销: 62

连接+投影

### 哈希匹配

使用来自顶部输入的每一行生成哈希表，  
使用来自底部输入的每一行探测该哈希  
表，然后输出所有匹配的行。

| 物理运算              | 哈希匹配            |
|-------------------|-----------------|
| Logical Operation | Inner Join      |
| 估计 I/O 开销         | 0               |
| 估计 CPU 开销         | 0.0181461       |
| 估计运算符开销           | 0.0181491 (62%) |
| 估计子树大小            | 0.0294932       |
| 估计行数              | 7               |
| 估计行大小             | 36 字节           |
| 节点 ID             | 0               |

### 输出列表

[DBS实验案例].[dbo].  
[depositor].customername

### 探测残留

[DBS实验案例].[dbo].[account].  
[accountnumber]=[DBS实验案例].[dbo].  
[depositor].[accountnumber]

### 哈希键探测

[DBS实验案例].[dbo].  
[account].accountnumber

count.accountnumber=depositor.account

er]

BS实验案例].[dbo].[account].[PK\_a...  
开销: 11 %

BS实验案例].[dbo].[branch].[PK\_bra...  
开销: 16 %

## Example Four (cont.)

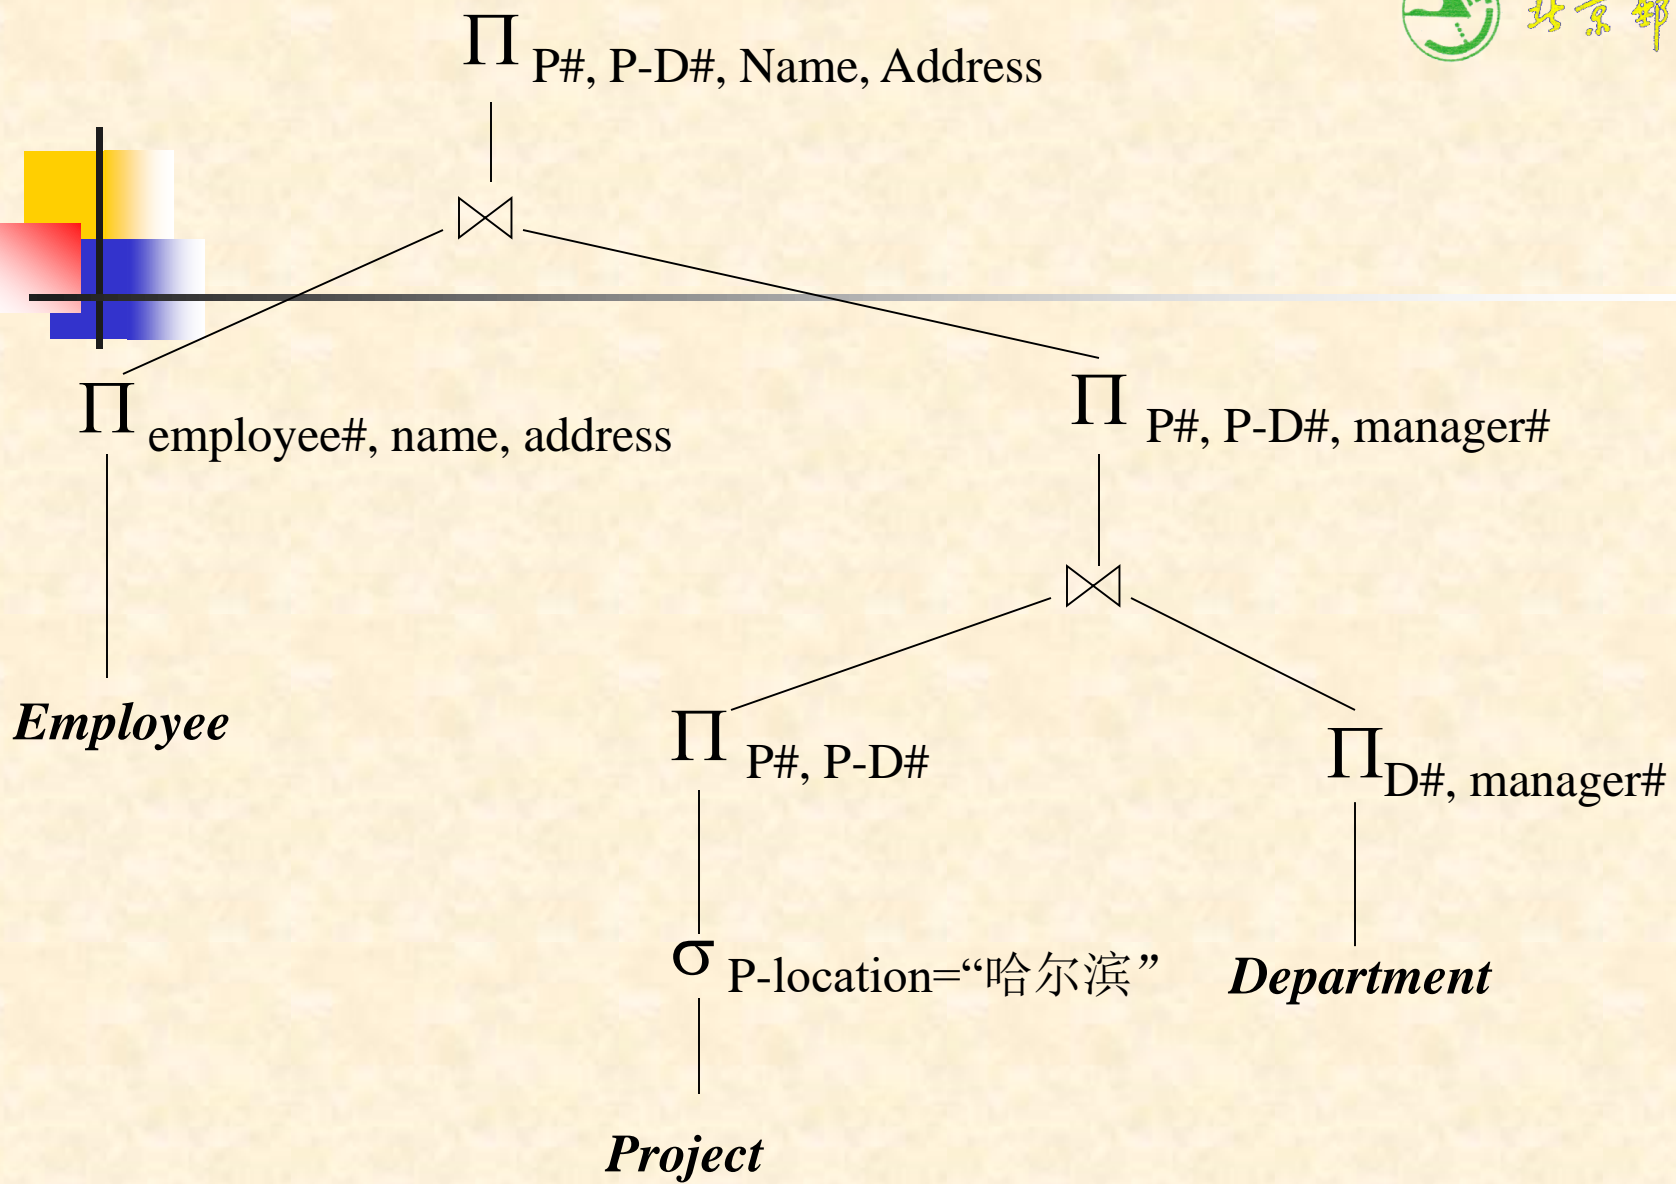

给定如下关系数据库

Employee(Employee#, Name, Address, Super-E#, E-D#)

Department(D#, Dname, manager#, depart-location)

Project(P#, Pname, Plocation, P-D#)

- 查询：对于每个在“哈尔滨”进行的工程项目，列出其工程项目号P#，工程所属部门号P-D#，该部门领导的姓名Name和地址Address
- 要求：写出该查询的SQL语句；转换为关系代数表达式并利用启发式方法进行查询优化；给出优化后的关系代数表达式。
- Select P#, P-D#, Name, Address  
From Project, Department, Employee  
Where Plocation = 哈尔滨 AND P-D# = D#  
AND manager# = Employee#





# 作业1

---

- Consider the following schema, where the primary keys are underlined ,
- Suppliers (supplier-id, supplier-name, city, telephone, address)
- Parts ( parts-id, parts-name, parts-color)
- Catalog (supplier-id, parts-id, price )
- .
- (1) Give an SQL statement to find out the *name* and *telephone* of the suppliers who supply a red *part* whose *price* is below \$2000.
- (2) Translate this SQL statement into an initial query tree, and give an optimized query *tree* for it, by means of heuristic query optimization.



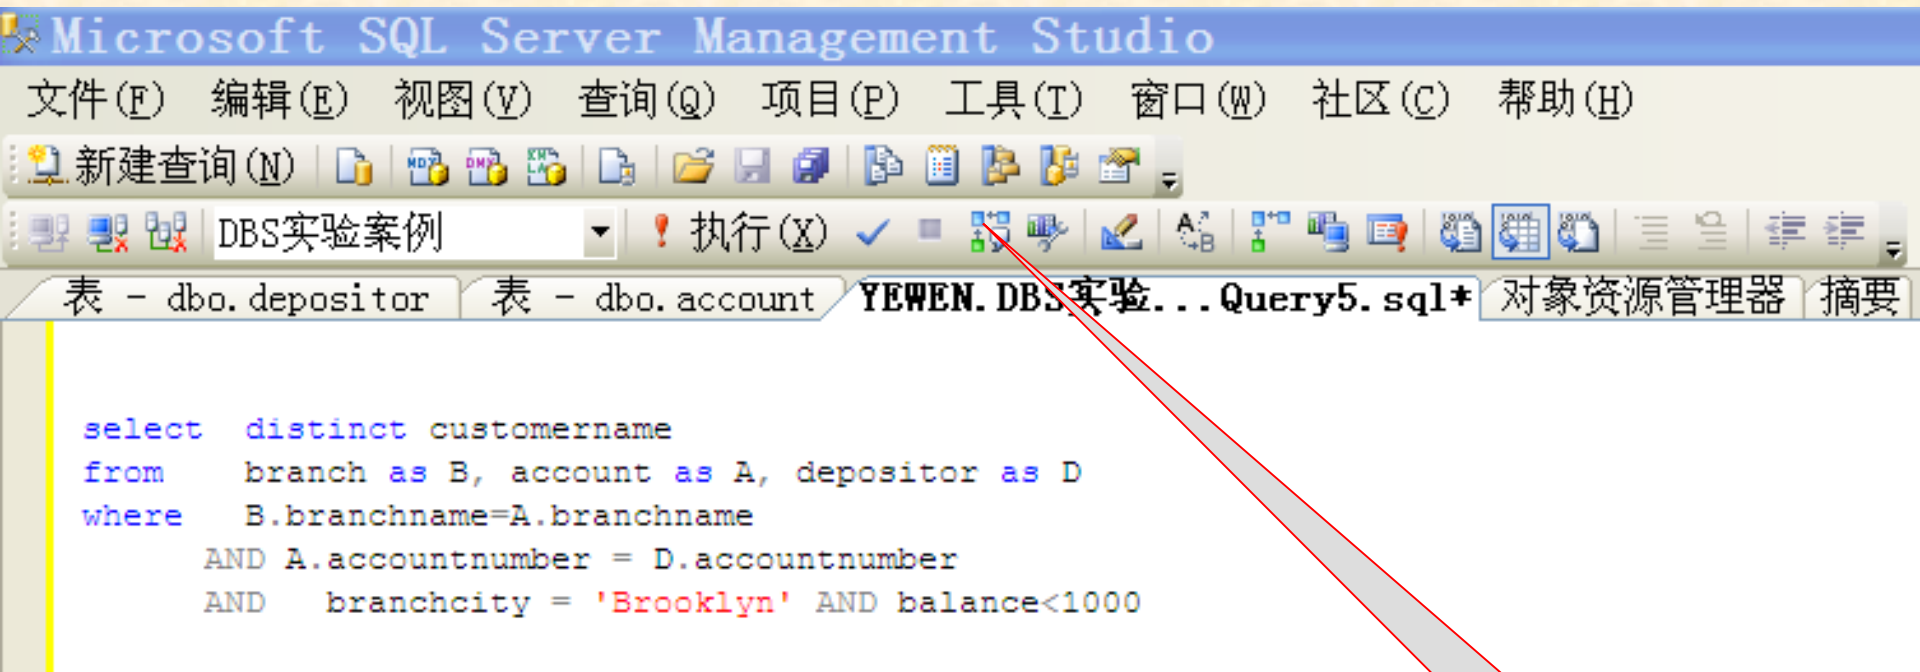


## 作业2

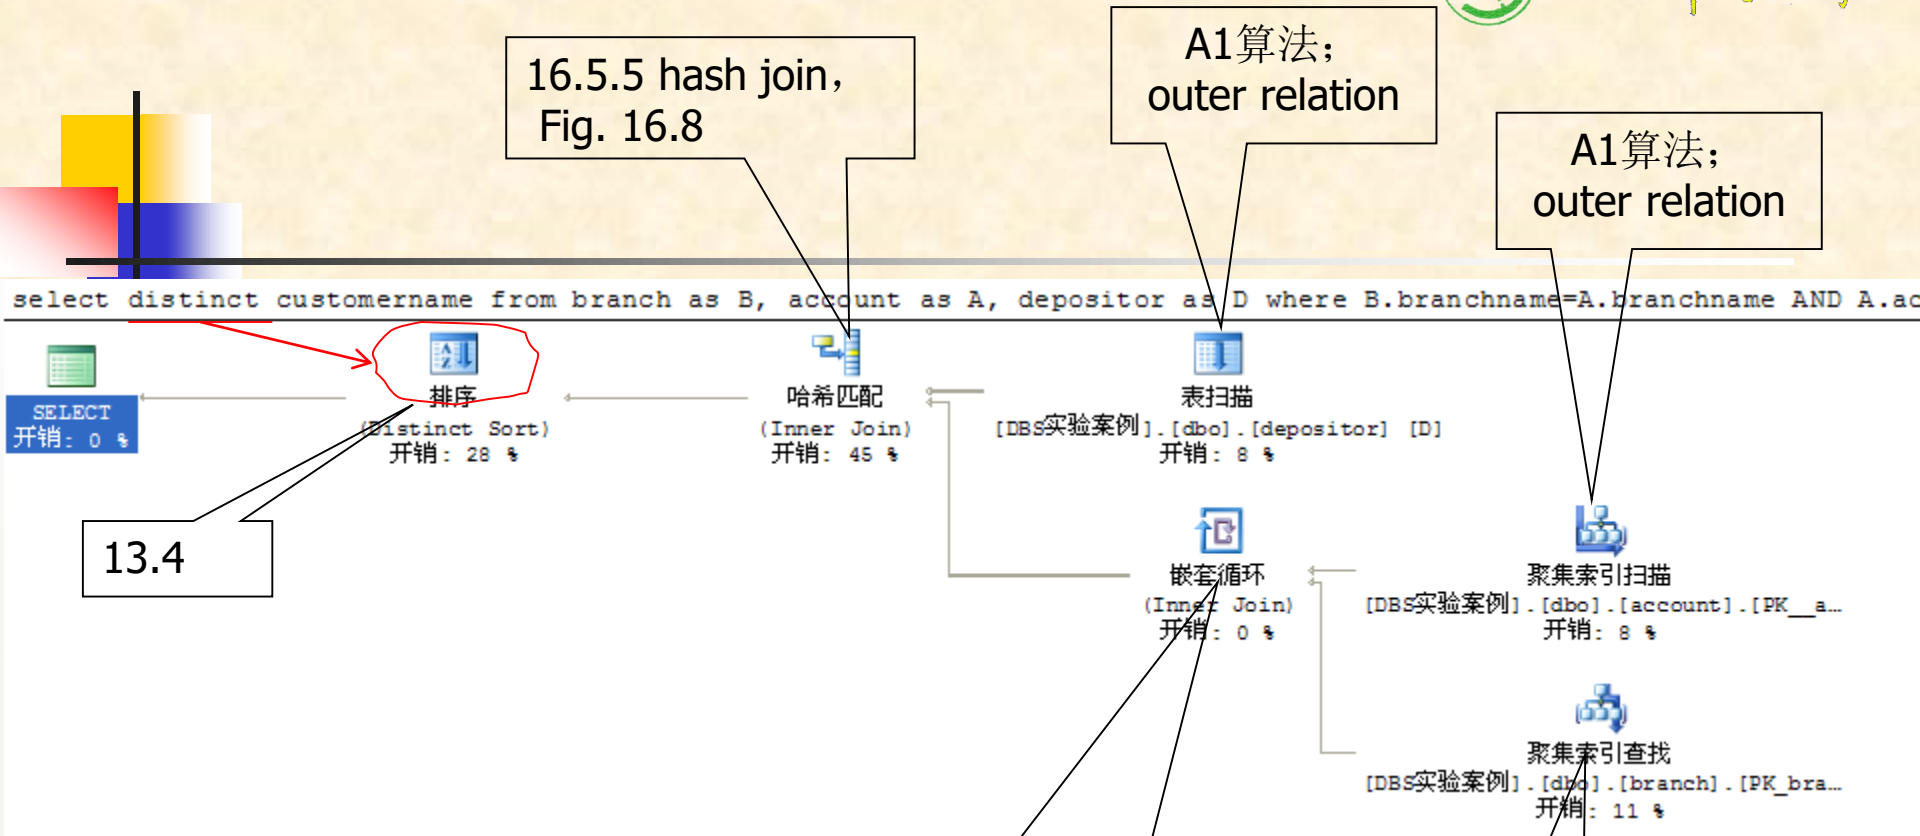
---

- Consider the database University given in the textbook,
- (1) Give a SQL statement to find some students and list their names and departments that they belong to. It is required that their total credits (presented by tot\_cred) are more than 40, their departments are located in Building 3, and they take the course identified by '2016CSDBS'.
- (2) For the SQL statement in (1), give an optimized query tree.

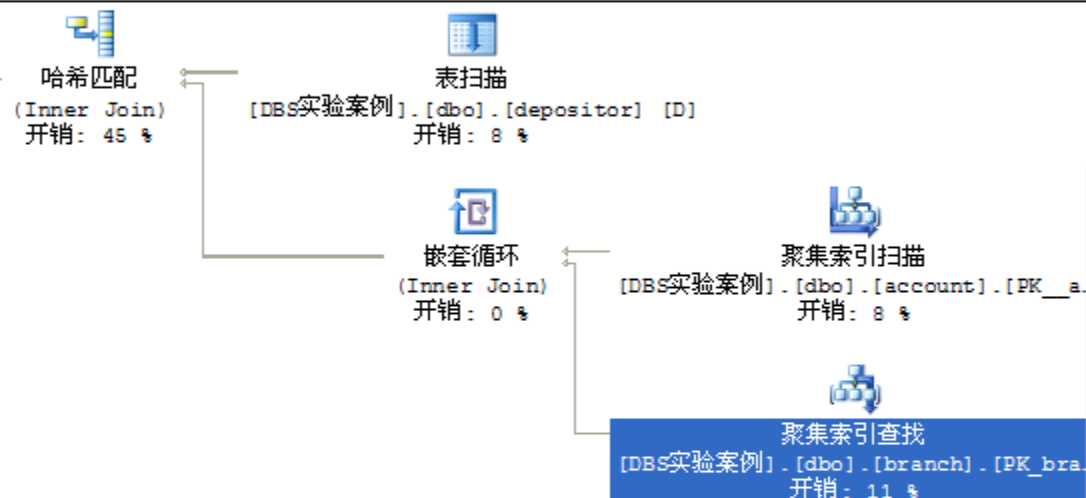
# Appendix A SQL Server平台下 观察比较SQL语句查询执行计划



显示估计的查  
询执行计划



ch as B, account as A, depositor as D where B.branchname=A.branchname AND A.accountnumber = D.accountnumber AND...



**聚集索引查找**  
扫描聚集索引中特定范围的行。

| 物理运算              | 聚集索引查找               |
|-------------------|----------------------|
| Logical Operation | Clustered Index Seek |
| 估计 I/O 开销         | 0.003125             |
| 估计 CPU 开销         | 0.0001581            |
| 估计运算符开销           | 0.0043898 (11%)      |
| 估计子树大小            | 0.0043898            |
| 估计行数              | 1                    |
| 估计行大小             | 20 字节                |
| 已排序               | True                 |
| 节点 ID             | 5                    |

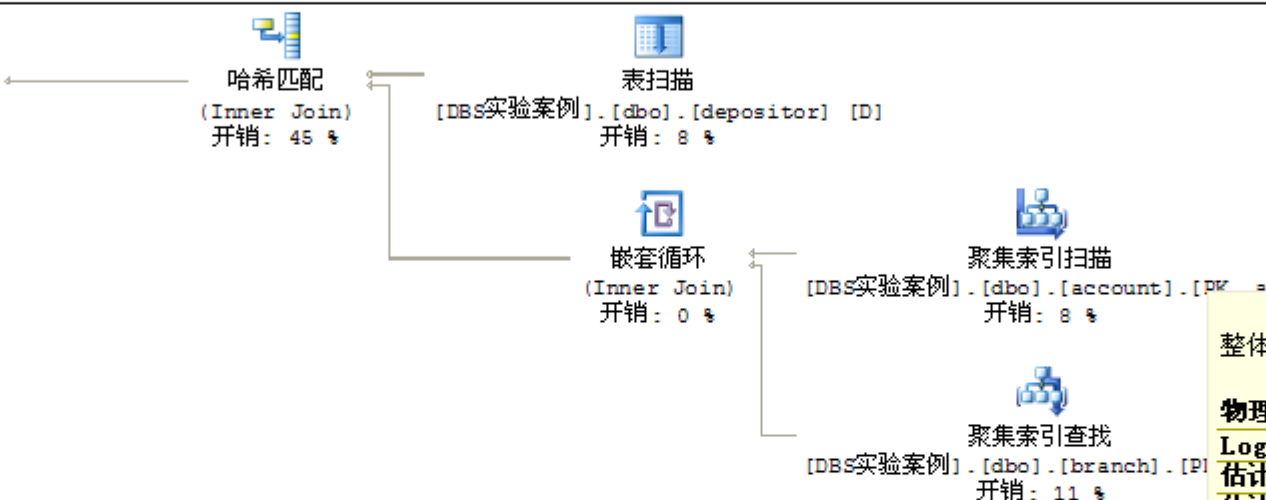
**谓词**  
[DBS实验案例].[dbo].[branch].[branchcity]  
as [B].[branchcity]='Brooklyn'

**对象**  
[DBS实验案例].[dbo].[branch].[PK\_branch]  
[B]

**输出列表**  
[DBS实验案例].[dbo].[branch].branchcity

**Seek 谓词**  
前缀: [DBS实验案例].[dbo].  
[branch].branchname= [DBS实验案例].  
[dbo].[account].[branchname] as [A].  
[branchname]

branch as B, account as A, depositor as D where B.branchname=A.branchname AND A.accountnumber = D.accountnumber AND...



**聚集索引扫描**  
整体扫描聚集索引或只扫描一定范围。

| 物理运算              | 聚集索引扫描               |
|-------------------|----------------------|
| Logical Operation | Clustered Index Scan |
| 估计 I/O 开销         | 0.003125             |
| 估计 CPU 开销         | 0.0001658            |
| 估计运算符开销           | 0.0032908 (8%)       |
| 估计子树大小            | 0.0032908            |
| 估计行数              | 8                    |
| 估计行大小             | 33 字节                |
| 已排序               | False                |
| 节点 ID             | 4                    |

**谓词**  
[DBS实验案例].[dbo].[account].[balance]  
as [A].[balance]<(1000.)

**对象**  
[DBS实验案例].[dbo].[account].  
[PK\_\_account\_\_07F6335A] [A]

**输出列表**  
[DBS实验案例].[dbo].  
[account].accountnumber, [DBS实验案例].  
[dbo].[account].branchname



例

! 执行(X)

or 表 - dbo.account YEWEN.DBS实验... Query5. sql\* 对象资源管理器 摘要

```
customername from branch as B, account as A, depositor as D where B.branchname=A.branchname AND A.accountnumber = D.ac
```



排序

(Distinct Sort)  
开销: 28 %



哈希匹配

(Inner Join)  
开销: 45 %



表扫描

[DBS实验案例].[dbo].[depositor] [D]  
开销: 8 %



嵌套循环

(Inner)

开销:

对于顶部(外部)输入的每一行, 扫描底部(内部)输入, 然后输出匹配的行。

## 嵌套循环

## 物理运算

## 嵌套循环

## Logical Operation

## Inner Join

估计 I/O 开销

0

估计 CPU 开销

0.0000334

估计运算符开销

0.0000411 (0%)

估计子树大小

0.0077217

估计行数

8

估计行大小

12 字节

节点 ID

3

## 输出列表

[DBS实验案例].[dbo].  
[account].accountnumber

## 外部引用

[DBS实验案例].[dbo].  
[account].branchname



YEWEN.DB实验...Query5. sql\*

对象资源管理器

摘要

行计划

inct customername from branch as B, account as A, depositor as D where B.branchname=A.branchname AND A.accountnumber

排序  
(Distinct Sort)  
开销: 28 %

哈希匹配  
(Inner Join)  
开销: 45 %

表扫描  
扫描表中的行。

物理运算

Logical Operation

估计 I/O 开销

估计 CPU 开销

估计运算符开销

估计子树大小

估计行数

估计行大小

已排序

节点 ID

对象

输出列表

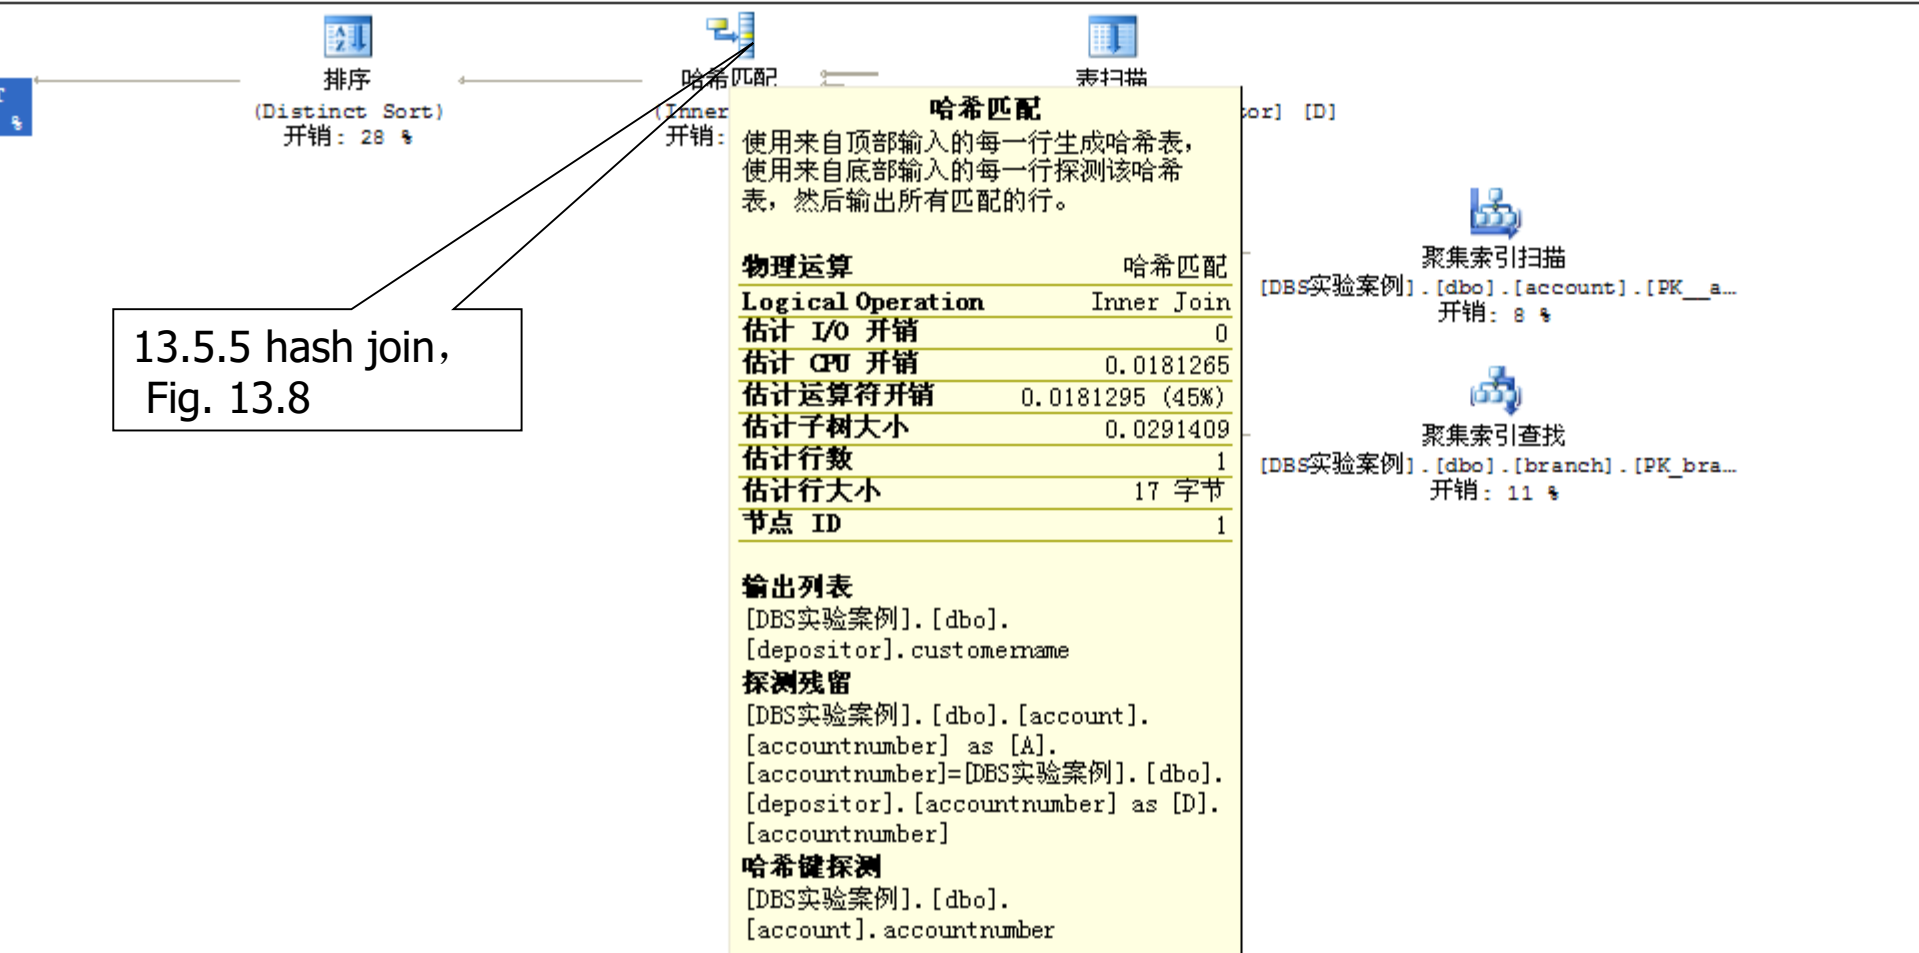
YEWEN (9.0 RTM)

YEWEN\yewenbupt

中文(中国)

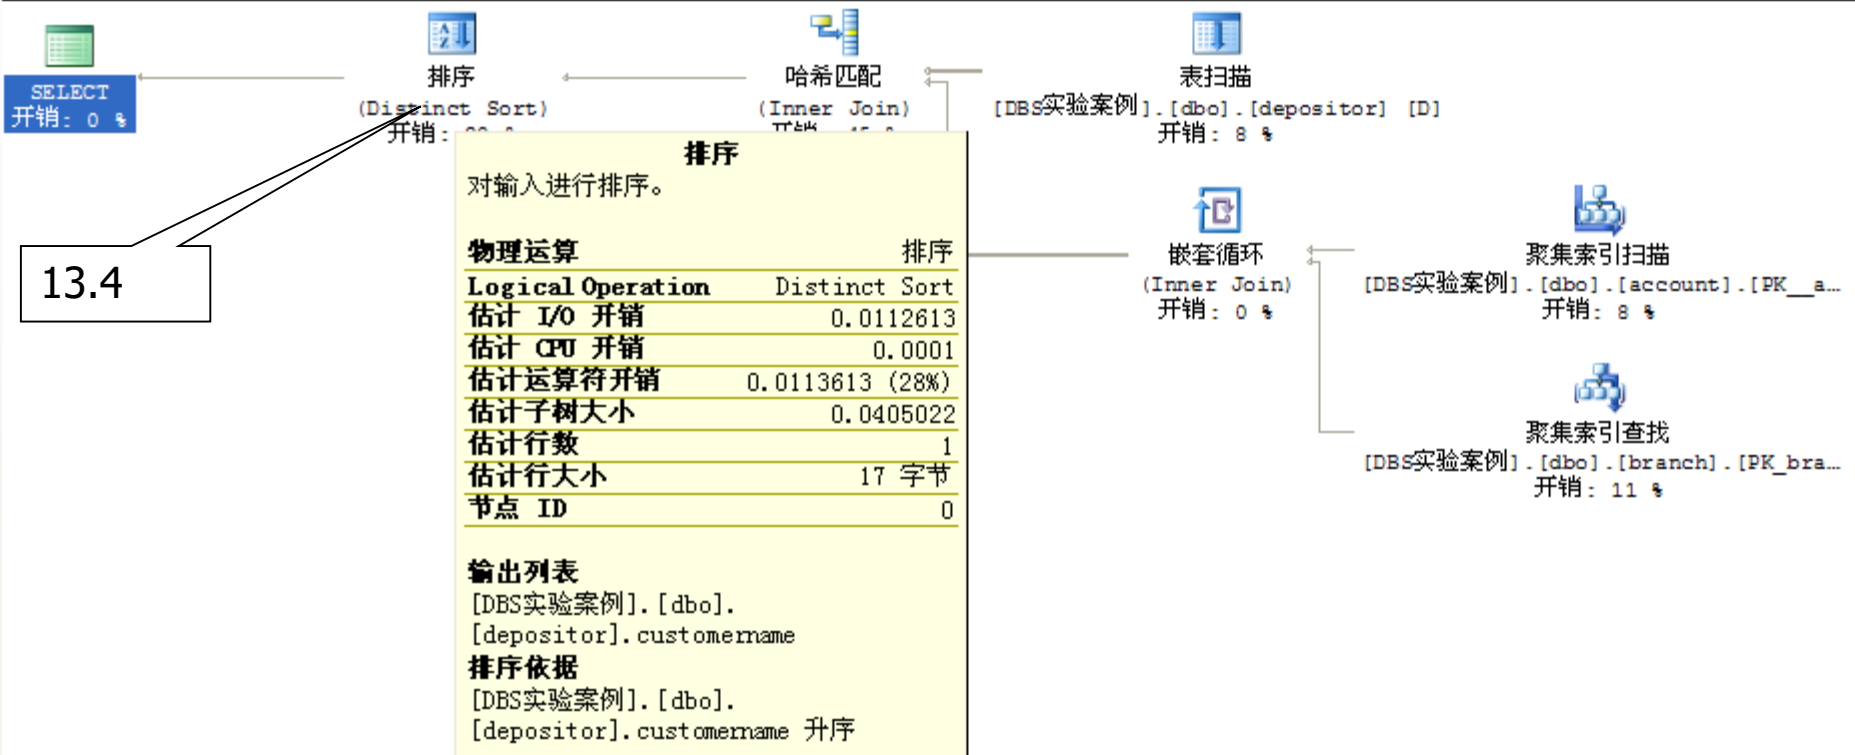
76%

执行计划  
distinct customername from branch as B, account as A, depositor as D where B.branchname=A.branchname AND A.accountnumber=D.accountnumber



13.5.5 hash join,  
Fig. 13.8

```
select distinct customername from branch as B, account as A, depositor as D where B.branchname=A.branchname AND
```



13.4

## 比较2条SQL语句查询成本

同时提交1批（2条）查询语句：

```
select distinct customername
from branch as B, account as A, depositor as D
where B.branchname=A.branchname
 AND A.accountnumber = D.accountnumber
 AND branchcity = 'Brooklyn' AND balance<1000
```

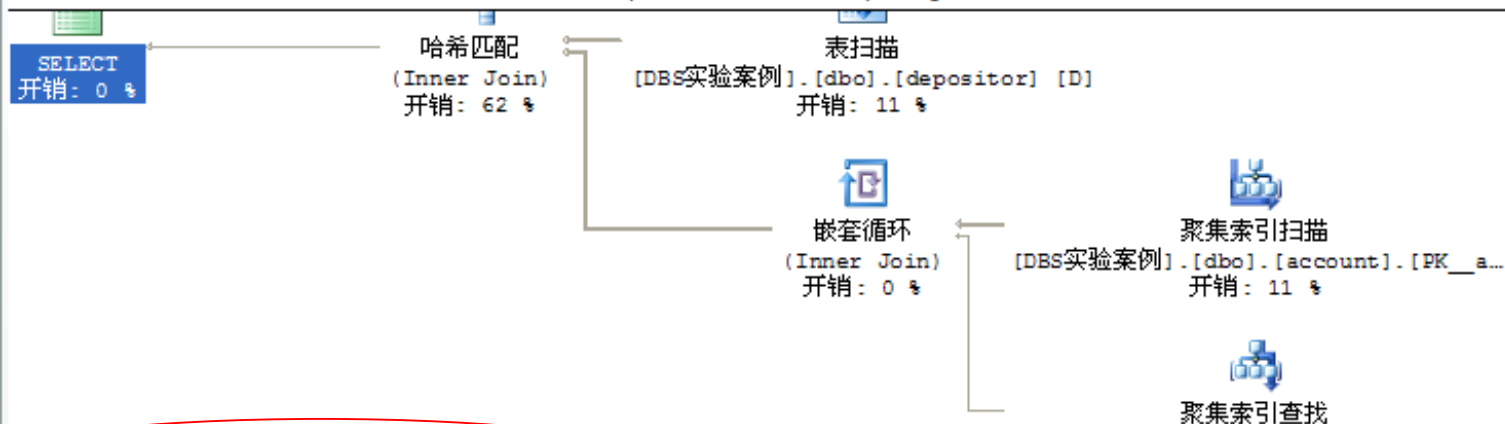
```
select customername
from branch as B, account as A, depositor as D
where A.accountnumber = D.accountnumber
 AND B.branchname=A.branchname
 AND branchcity = 'Brooklyn' AND balance<1000
```



该批次中2条查询语句的成本之比： 42% vs. 58%

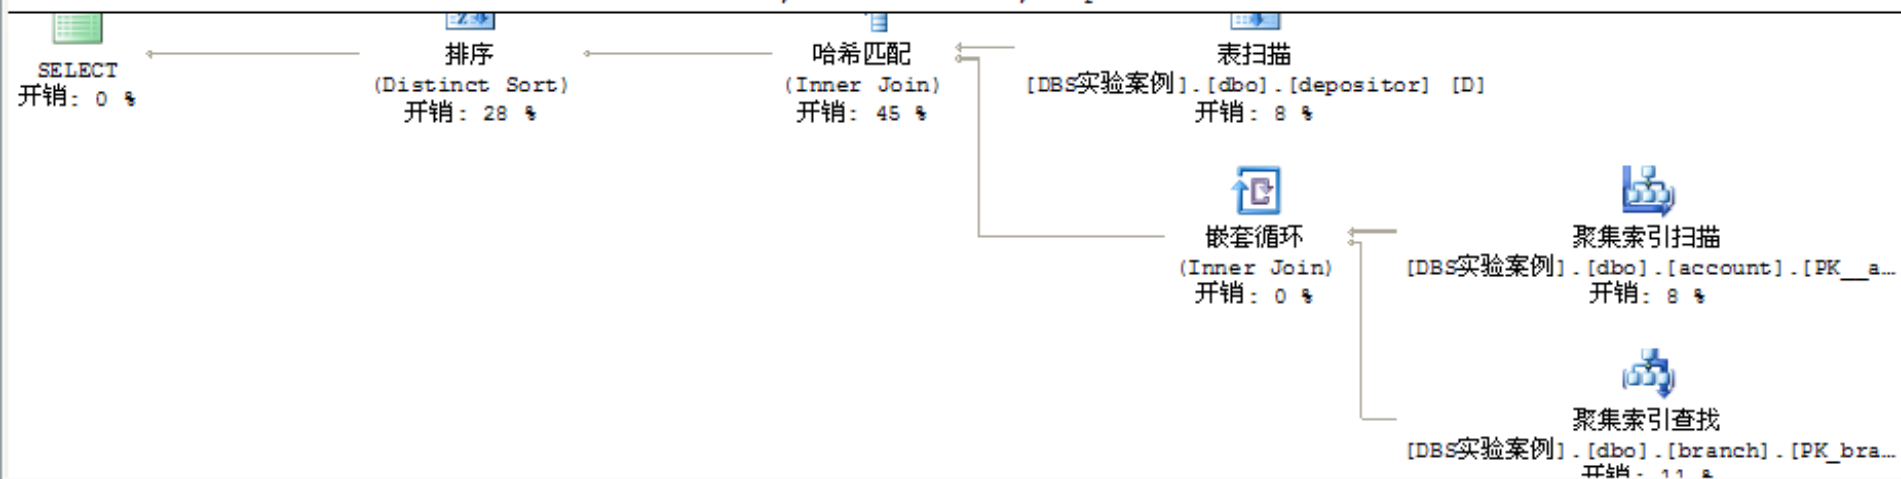
查询 1: (与该批有关的) 查询开销: 42%

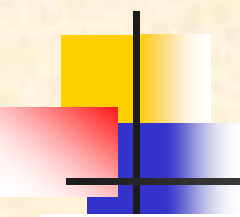
select customername from branch as B, account as A, depositor as D where B.branchname=A.branchname AND A.accountnumber=D.accountnumber



查询 2: (与该批有关的) 查询开销: 58%

select distinct customername from branch as B, account as A, depositor as D where B.branchname=A.branchname AND A.accountnumber=D.accountnumber





```
select *
from account
where balance >= 0 and balance <= 50
```


```
select *
from account
where balance in (
 select balance
 from account
 where balance >= 0 and balance <= 50)
```



消息 执行计划

查询 1: (与该批有关的) 查询开销: 33%

```
select * from account where balance >= 0 and balance <= 50
```



查询 2: (与该批有关的) 查询开销: 67%

```
select * from account where balance in (select balance from account where balance >= 0 and balance <
```

