



数据库系统原理

Database System Principle

邵莹侠

Email: shaoyx@bupt.edu.cn

北京邮电大学计算机学院

计算机应用技术中心




PART 7

TRANSACTION MANAGEMENT

C, Pascal programs

程序编译 / 编译器



process / thread

程序执行 / OS

进程管理
并发控制
进程调度
死锁处理


Chapter.
15, 16

Chapter.
17, 18, 19

Fig. 15.0.1

query

Query processing / DBMS



transaction

事务处理 / DBMS

事务管理 (§ 17)
并发控制 (§ 18)
事务调度 (§ 18)
死锁处理
恢复技术 (§ 19)



Introduction to Part 7

- DBS中，从DBMS的角度，用户对DB的访问体现为DBS中1个或多个**事务**的执行
- 事务是DBS中应用程序/数据库应用系统的基本逻辑单位，也是DBMS管理DBS运行的基本单位，§ 17.1/17.2
- 多用户DBS中，多个用户对DB的并发访问体现为DBS中多个**数据库事务**的并发执行
- 事务是动态的，具有一定生命周期
 - 事务具有多种状态，事务的执行体现为各状态间的转换过程，§ 17.4, Fig.17.1

Introduction to Part 7 (cont.)

- 在DBS中，事务执行时，特别是当多个事务对**共享数据**进行并发访问时，为维护DBS系统中数据的正确性(integrities)，事务必须满足一定的约束条件，表现为事务的4个基本特征(**ACID**)，§ 17.1/17.2/17.4/17/5
 - **原子性**(atomicity)，**一致性**(consistency)，**独立性/隔离性**(isolation)，**永久性/持续性/操作结果永久保持性**(durability)
- DBMS中的**recovery-management component**负责保障事务的原子性(§ 19)
- **事务设计者/DBS应用程序编程者**负责保障事务的一致性



Introduction to Part 7 (cont.)

- DBMS中的concurrency-control component负责保障事务的独立性(§ 18)
- DBMS中的recovery-management component负责保障事务的永久性(§ 19)
- 当DBS中存在多个并发事务时，DBMS通过事务调度对各事务进行并发控制(§ 18.4)，以保证并发事务的运行结果正确性
 - DBMS中的concurrency-control component依据事务调度可串行性 (Serializability)的基本原理，对事务进行并发控制和调度, § 18.6



Introduction to Part 5 (cont.)

- 具体实现技术包括：
 - 基于锁的并发控制技术(§ 18.1) ， 多粒度控制技术(§ 18.3)；
 - 基于时间戳的并发控制技术(§ 18.4)、基于验证的并发控制技术(§ 18.5)、多版本控制(§ 18.6)
- 此外， DBMS的concurrency-control component还需要对多个并发执行的事务进行死锁处理(§ 18.2)



Chapter 17

TRANSACTIONS



§ 17.1 Transaction Concept

- Definition and composition
- *Read* and *write*
- *ACID* properties



Transaction Concept (cont.)

- The transaction is
 - *a unit of* DBS application *program executing* that accesses and possibly updates various data items in DB
- A transaction is *initiated* by a user application program written in
 - data manipulating language such as SQL, for example *select-from-where*, or
 - programming languages, e.g. C++ or Java, with embedded DB accesses in JDBC or ODBC



Transaction Concept (cont.)

- A transaction consists of
 - *a collection of operations* that form a single logic unit of application works, delimited by the statements of the form *begin transaction* and *end transaction*
 - operations: *read* or *write* on DB, or other operations
 - begin-transaction
 - $op_1;$
 - $op_2;$
 - .
 - $op_n;$
 - end-transaction (e.g. *commit* or *abort/rollback*)

Transaction Concept (cont.)

- Conceptually, a transaction can be defined by *read* and *write* operations, which are independent of DBS query languages (e.g. SQL) and DBMS (e.g. SQL Server)
- E.g.1. *Account-transfer* T1 :

transfer \$50 from account A to account B

T1 : **begin-transaction**

read (account_A);

account_A := account_A -50;

write (account_A);

read (account_ B);

account_ B := account_ B + 50;

write (account_ B);

end-transaction

begin-transaction

```
read ( amount_A );  
amount_A := amount_A - 50;  
if amount_A < 0 then  
  then begin  
    print('insufficient funds')  
    rollback T2  
  end  
else begin  
  write (account_A );  
  read (amount_B );  
  amount_B := amount_B + 50;  
  write (amount_B );  
  commit T2  
end
```

- abort/ **rollback** T2: 撤销之前对loan_A 的修改, 即恢复loan_A 的原值, 并立即结束T2, 不再执行其后的各操作

Fig.17.0.1 *Account_transfer* T2 with **rollback**

```
DECLARE @transfer_name varchar(10) /*事务变量定义*/
SET @transfer_name = 'I-transfer-from-A-to-B' /*事务命名*/
BEGIN TRANSACTION @transfer_name /*事务开始*/
USE ACCOUNT /*打开数据库ACCOUNT*/
GO /*将上述批SQL语句提交SQL Server*/
UPDATE account_A /* 修改A帐户*/
    SET balance = balance - 50
    WHERE branch_name= 'Brooklyn'
UPDATE account_B /* 修改B帐户*/
    SET balance = balance + 50
    WHERE branch_name= 'Brooklyn'
GO
COMMIT TRANSACTION @transfer_name /*事务提交*/
GO
```

Fig.17.0.2 *Account_transfer* T3 in T-SQL



ACID Properties (ref. to 17.2)


To preserve the integrity of data the database system must ensure:

- **Atomicity.** Either all operations of the transaction are properly reflected in the database or none are.
- **Consistency.** Execution of a transaction in isolation preserves the consistency of the database.

ACID Properties (cont.)

- **Isolation.** Although multiple transactions may execute concurrently, each transaction must be unaware of other concurrently executing transactions. Intermediate transaction results must be hidden from other concurrently executed transactions.
 - That is, for every pair of transactions T_i and T_j , it appears to T_i that either T_j finished execution before T_i started, or T_j started execution after T_i finished——好像两者串行执行
- **Durability.** After a transaction completes successfully, the changes it has made to the database persist, even if there are system failures.

17.2 A Simple Transaction Model (*Read* and *Write*)

- Each data access in transactions, such as *select*, *Update* in SQL, or *API* in ODBC, is translated/decomposed by DBMS into one or more *read* and *write* operations
- Refer to Fig. 17.0.3 
- DBMS executes these *read* and *write* operations to fulfill data access in transactions
- For each transaction, DBMS allocates a *local buffer* in main memory as the working area for this transaction
- The database permanently resides on disk, but some portion of it is temporarily residing in the *disk buffer* in main memory

Read and Write Operations (cont.)

■ *read(X)*

- transfer the data item *X* from DB files on disk or the *disk/system buffer* to a variable, also called *X*, in the *local buffer* belonging to the transaction that executes the *read* operation

■ *write(X)*

- **conceptually**, transfer the value in the variable *X* in the *local buffer* of the transaction that execute the *write* operation back to the data item *X* in the DB files on disk
- In real DBS, the *write* operation issued by transactions does not necessarily result in the immediate update of the data on the disks, this operation may be temporarily stored in the *system buffers* and executed on the disk later
 - refer to Fig. 17.0.3

Transaction

data accesses on data item **x** and **y** issued by T_i ,
e.g. *select, insert, delete, update, ...*

DBMS

逻辑读写

read(x)

write(y)

物理读写

input(X)

output(Y) / *reflect*

DB file on disk

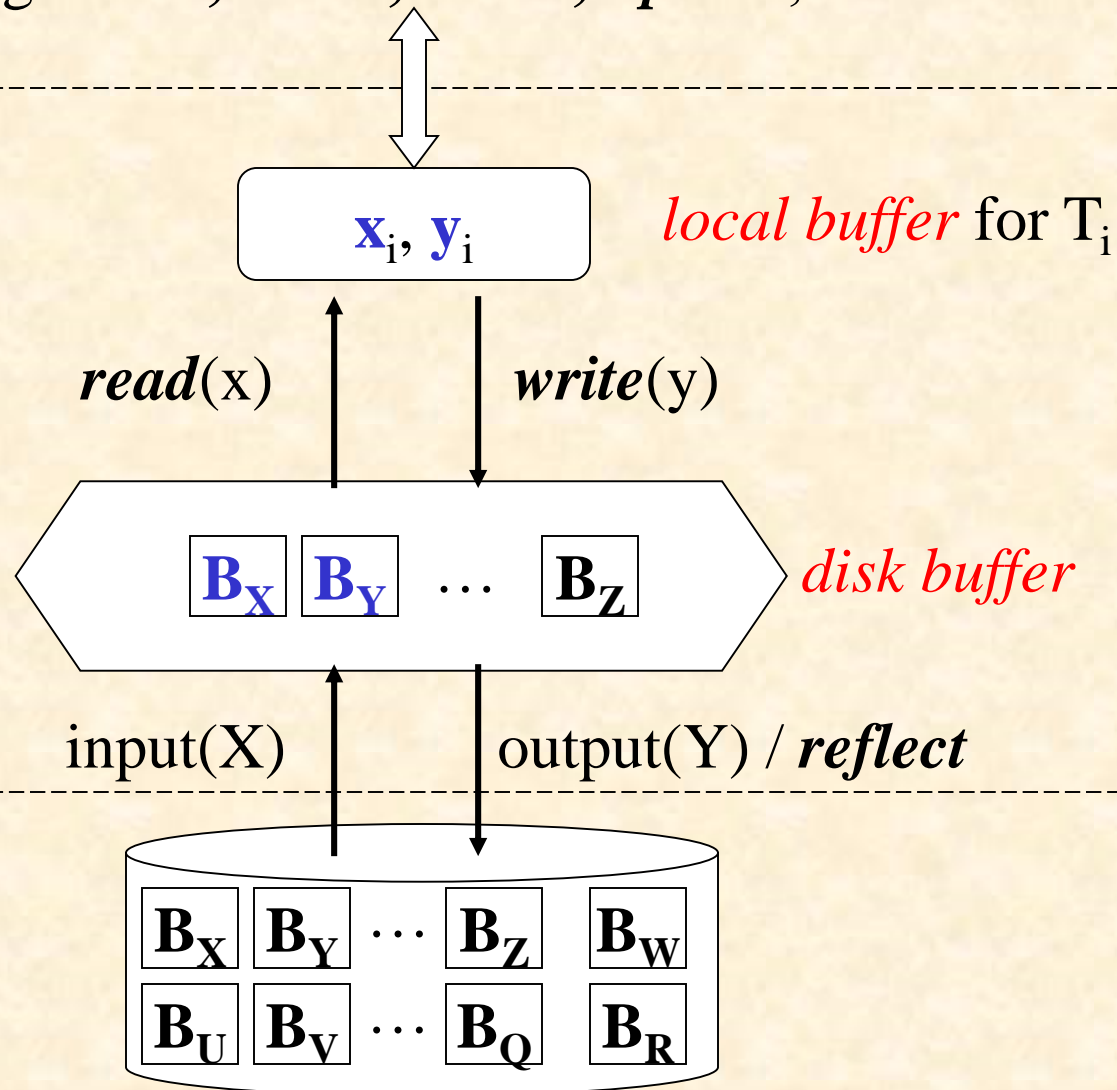


Fig.17.0.3 *read* and *write* operations





Transaction Properties

- To guarantee integrities of DB, *ACID Transaction Properties* should be maintained
 - Atomicity, Consistency, Isolation, Durability

- Consistency
 - execution of a transaction in isolation (that is, with no other transaction executing concurrently) preserves the consistency of the DB, i.e. **before** and **after** the transaction execution, DB is in a **correct** DB state
 - /*1个事务的正确执行使得DB从一个正确状态转移到另一个正确状态
 - DB state: contents of tables in DB, or DB instance

- consistency/correctness

integrity constraints, specified by the application logic, are not violated

- during transaction execution the database may be temporarily inconsistent.

/*事务执行过程中，完整性约束有可能被破坏

- ensuring consistency of an individual transaction is the responsibility of the application programmers who codes the transaction


- E.g. The sum of A and B is unchanged by the execution of the transaction

A=1000, B=2000

1. *read* (A);
2. A := A -50;
3. *write* (aA);
4. *read* (B)
5. B := 50;
6. *write* (accB);

A=950, B=2050

Atomicity

- Either all operations of the transaction are *reflected* () properly in the database, or none are
 - /*事务中所有操作或者全部成功完成，并且这些操作结果 被写入到DB；或者事务中的所有操作一个都不作，该事务对数据库和其他事务没有任何影响
 - /*将事务所有操作作为1个不可分割的整体
- E.g. If the transaction fails after step 3 and before step 6, money will be “lost” leading to an inconsistent database state
 - failure could be due to software or hardware

A=1000, B=2000

1. *read* (A);
2. A := A -50;
3. *write* (aA);

A=950, B=2000

4. *read* (B)
5. B := 50;
6. *write* (B);



Atomicity (cont.)

- The log-based recovery component in DBMS is responsible for ensuring atomicity



Isolation

- Even though multiple transactions may execute concurrently, the system guarantee that all transactions seem to execute **serially**, so the consistent states of DB can be preserved
 - each transaction is unaware of other transactions executing concurrently in the system
 - /*对多个并发执行的事务，1个事务的执行不能被其他事务干扰，即1个事物的内部操作和数据对其它事务是隔离的，并发执行的事务间不能相互干扰，从执行结果来看，相当于事务串行执行
- **Isolation** guarantees the correctness of concurrent execution of more than one transactions
- The concurrency control component in DBMS ensures the isolation properties

- **E.g.** if between steps 3 and 6 (of the fund transfer transaction) , another transaction **T2** is allowed to access the partially updated database, it will see an inconsistent database (the sum $A + B$ will be less than it should be).

A=1000,

B=2000

T1

T2

1. **read(A)**
2. $A := A - 50$
3. **write(A)**

A=950,

B=2000


read(A), read(B), print(A+B)

4. **read(B)**
5. $B := B + 50$
6. **write(B)**


- Isolation can be ensured trivially by running transactions **serially**
 - That is, one after the other.



Durability

- After a transaction completes successfully (i.e. the transaction is *committed*), the changes it has made to the DB persist, even if there are **system failure**
 - /*事务一旦提交（成功完成），它对数据库中数据的改变就应该是永久性的，这些改变不随其后的数据库系统错误而丢失
 - ensured by recovery management component
- **E.g.1** the transfer of the \$50 has taken place, the updates to the database by the transaction must persist even if there are software or hardware failures.
- **E.g.2** the modifications on *Y* in Fig.17.0.3 
 - the modified results in *local buffer* should be *reflected* to the *DB file* on disk

Transaction Properties (cont.)

e.g.2. for T3 in Fig.17.0.2 

begin-trans.;












update(A) ;

update(B) ;

commit

local buffer
for T3

disk buffer





Transaction Properties (cont.)

■ ACID mechanisms in DBS

- A : transaction management / recovery management component
- C: programmer/transaction designer, integrity constraints testing mechanism in DBMS
- I: concurrency control component
- D: recovery management component

§ 17.4 Transaction Atomicity and Durability - States

- Committed (提交) transaction
 - successfully complete all its operations

- Aborted (撤销/夭折) transaction
 - not finish all its operations successfully
 - *rollback*

- During its execution, a transaction can stay in several states, as shown in Fig.17.1
 - taking Fig.17.0.5, Fig.17.0.6 as examples

Transaction States (cont.)

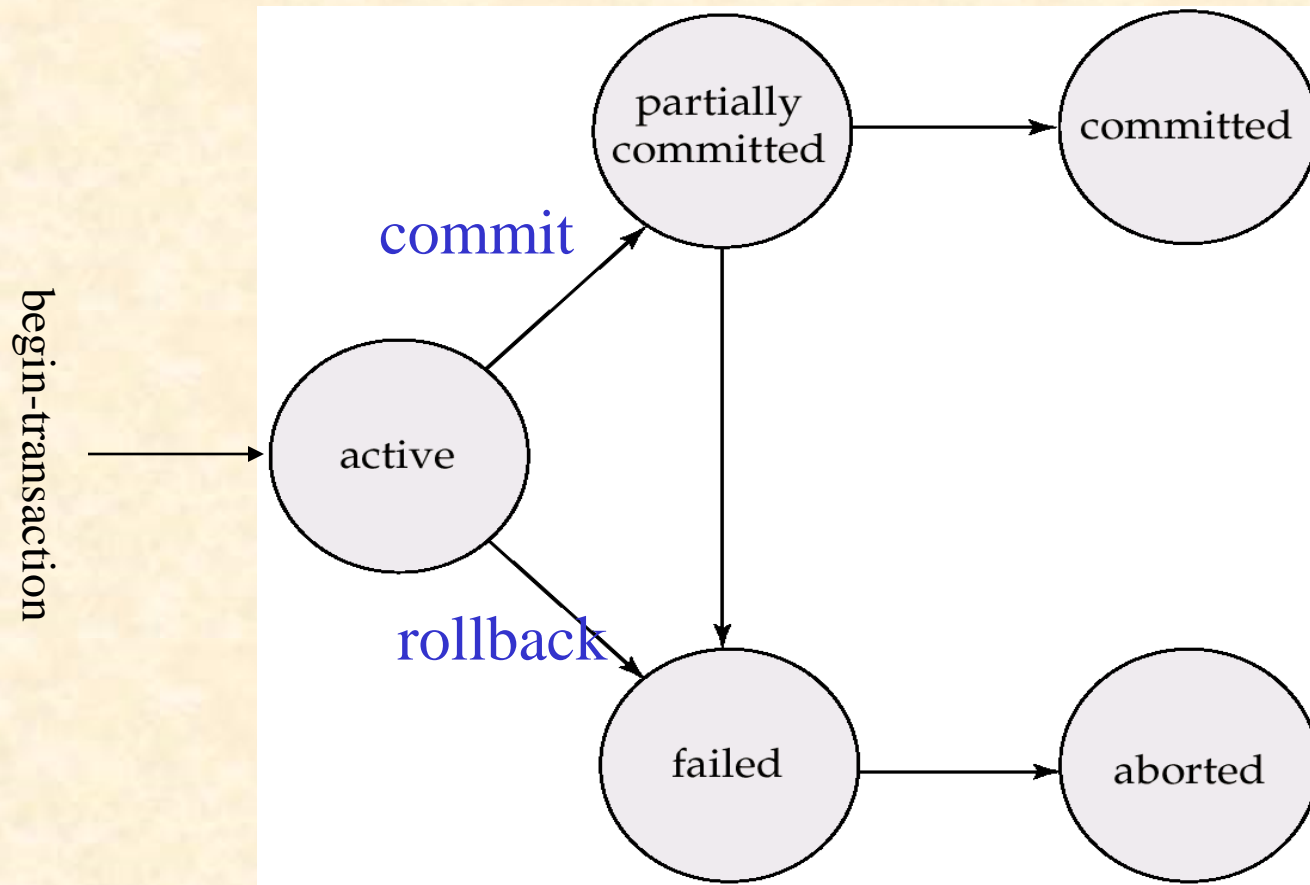


Fig.17.1. State diagram of a transaction

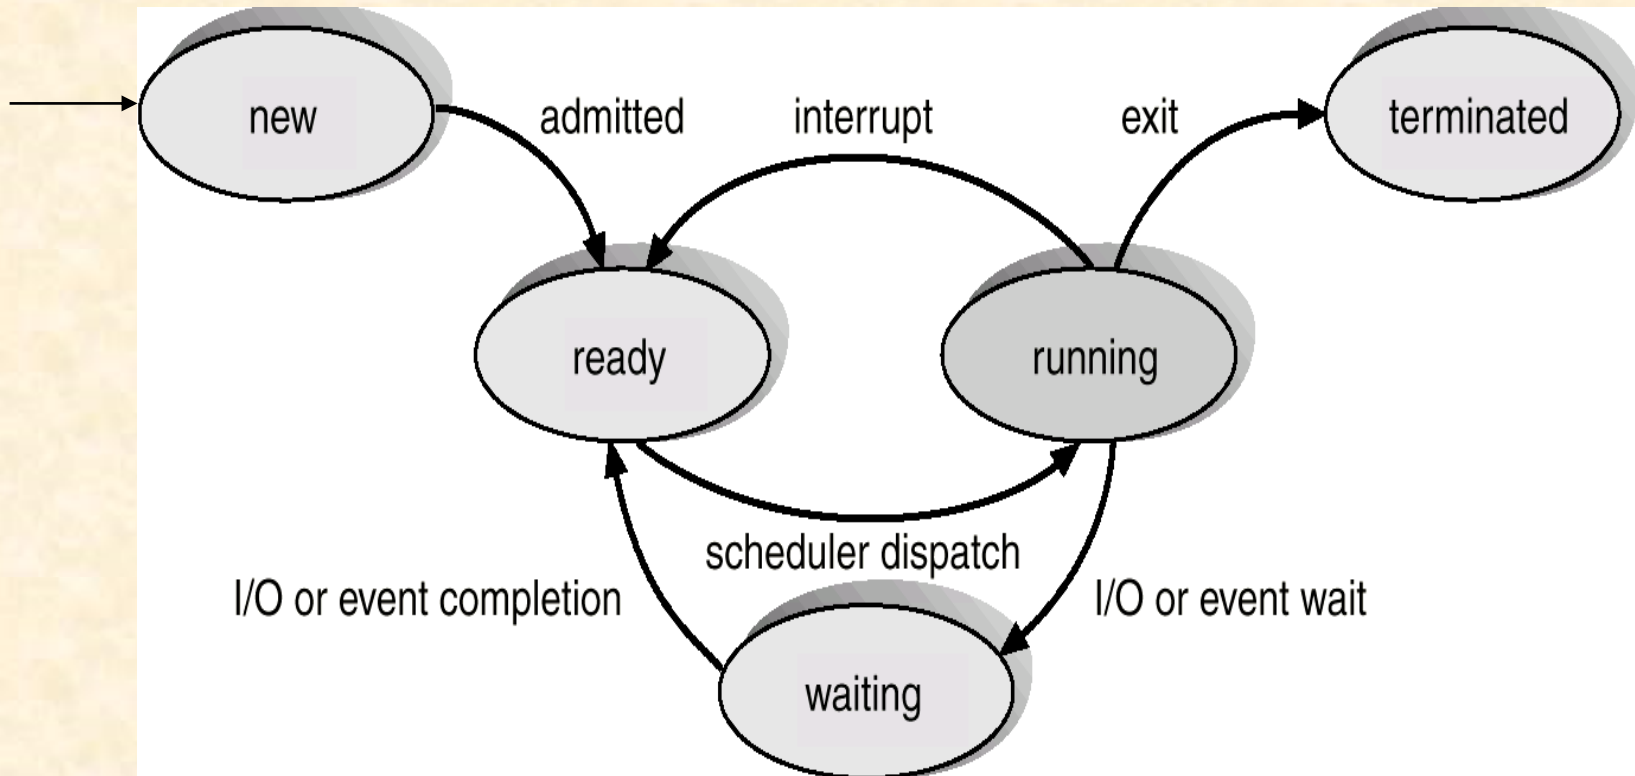


Fig.17.0.4 Process states

Transaction

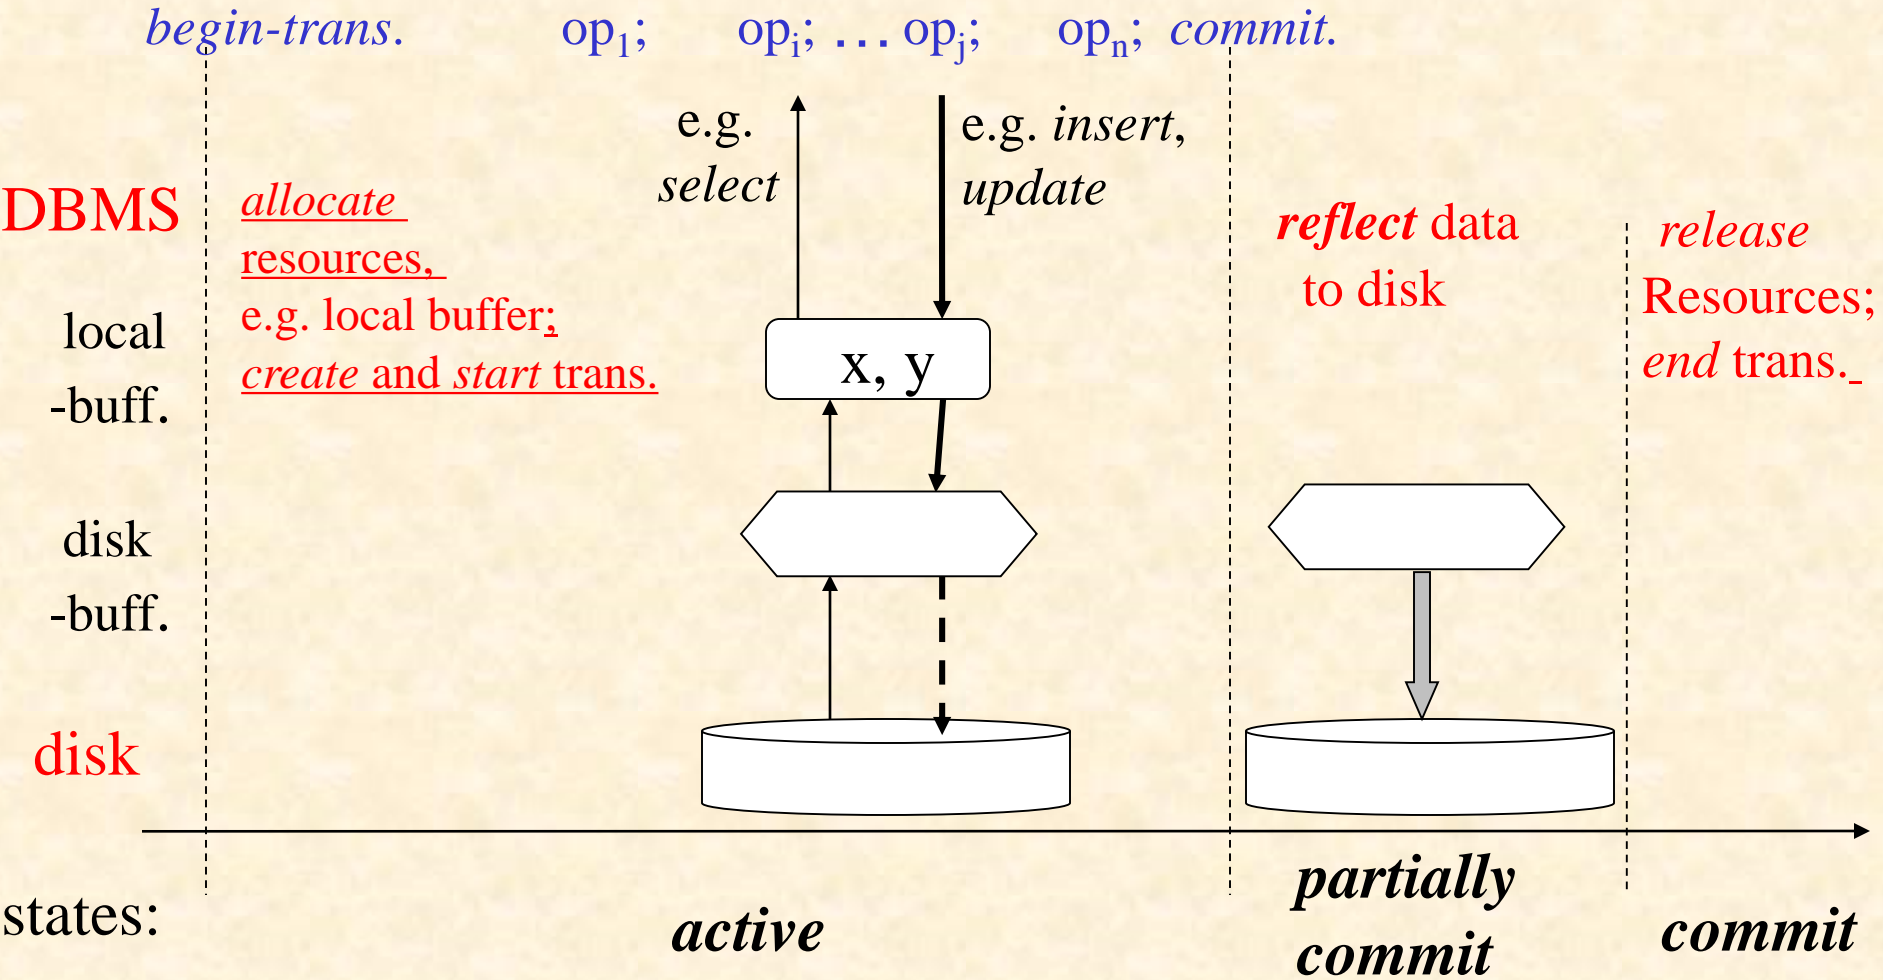


Fig.17.0.5 Life-cycle of a successful/*committed* transaction

Transaction

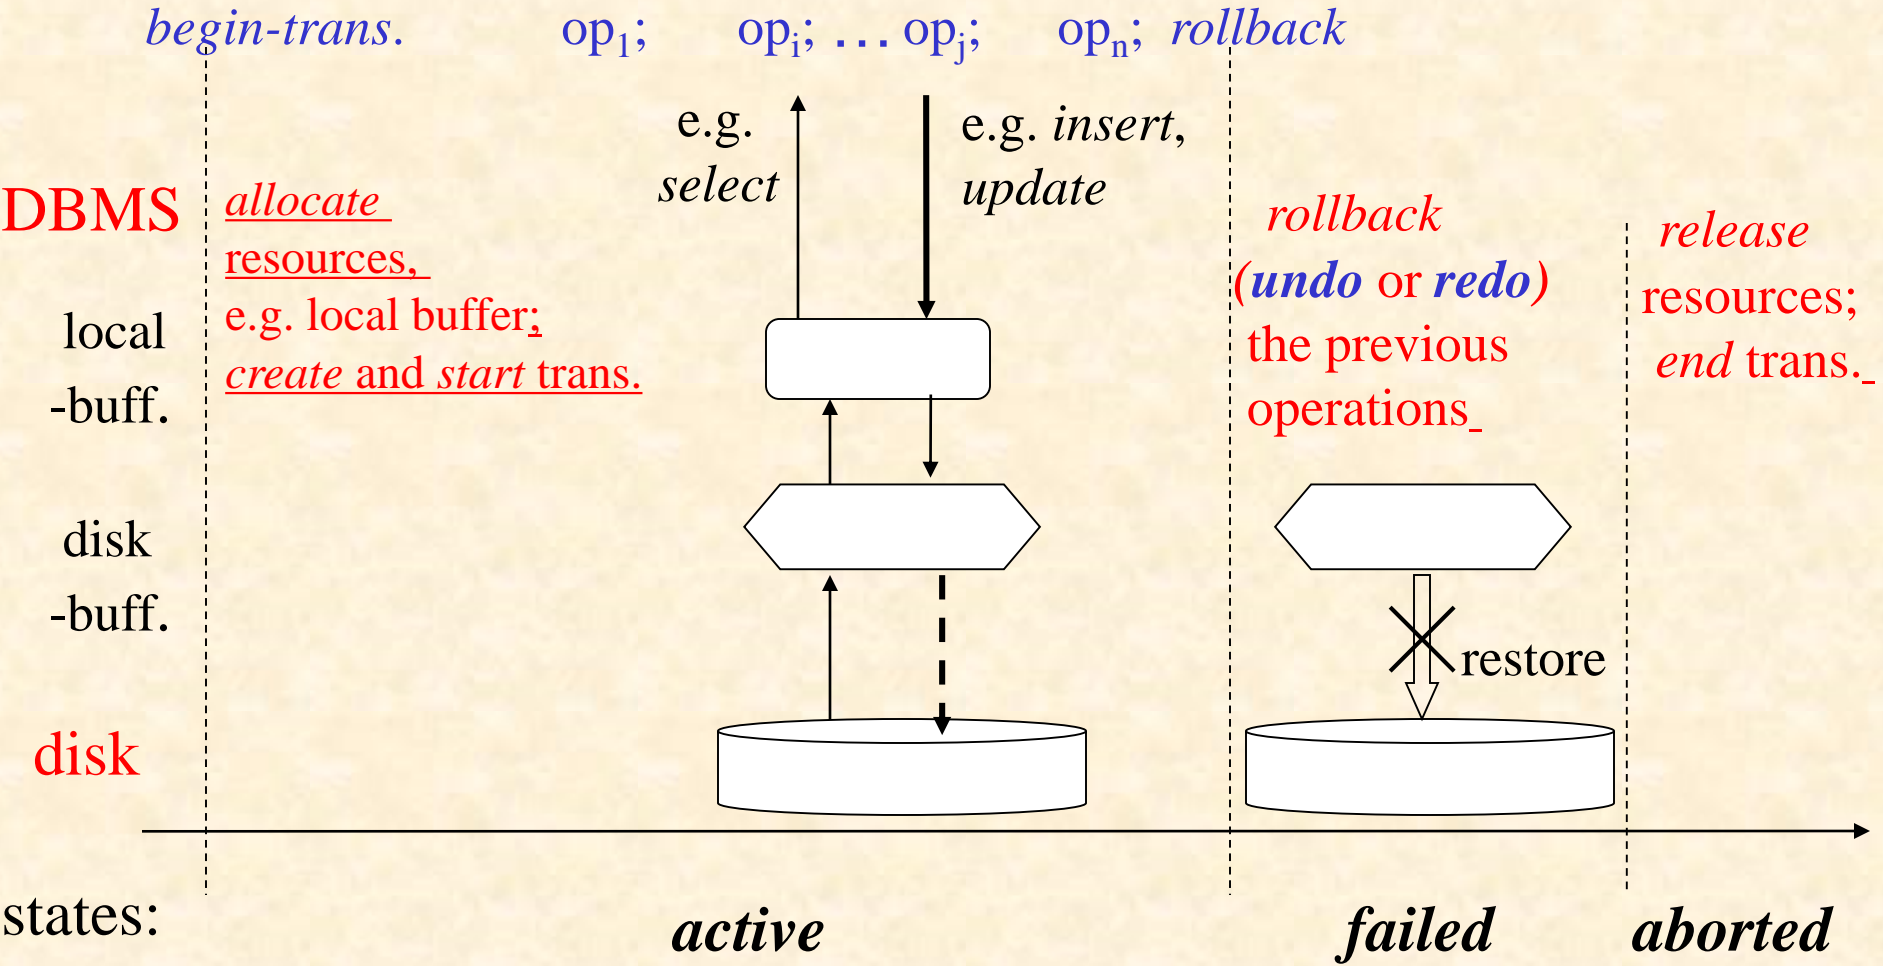



Fig.17.0.6 Life-cycle of a *rollback* transaction

Transaction States (cont.)



■ Active

- the initial state, the transaction stays in this state while it is executing from *begin-transaction*
- in this state, the transaction is created (after **begin-transaction** is submitted), resources, e.g. local buffer are allocated to the transaction
 - e.g. 

■ Note

- transaction operations are executed, but the modifications on DB issued by operations **may** only stored temporally in the disk buffer, not *reflected* to DB file on disks immediately

Transaction States (cont.)

- Partially committed (部分提交) :


- after the final statement (i.e. *Commit*) has been submitted, the transaction enters this state
- in this state, the influences that the transaction's operations have made on DB are reflected from the disk buffer to DB by DBMS

/*DBMS将disk缓冲区中的DB修改结果写入磁盘DB文件中

- e.g. 



Transaction States (cont.)

■ Committed

- after the transaction has successfully completed all its operations, and all the results of these operations have been reflected to the DB in *partial commit* state, it enters this state
- in this state, the transaction
 - releases the resources occupied
 - and then terminated (quits DBS),
the cycle-life of the transaction is ended
- e.g. 

Transaction States (cont.)

■ Failed


- if it is discovered that normal execution can no longer proceed, the transaction enters this state
 - e.g. *Account_transfer* T2 in Fig.17.0.1 
- in this state, the transaction is *rolled back* to restore DB to the state prior to the start of the transactions (§ 19)
/*在此状态下进行失败处理工作
- e.g. Fig.17.0.6 

Transaction States (cont.)

- Aborted (夭折/异常结束/中止状态)

- after the transaction has been rolled back in *failed* state, and the DB has been restored to its state prior to the start of the transactions, it enters this state to end the transaction

/*表示DBMS已完成失败处理工作, 事务将要退出系统

- in this state, the transaction
 - releases the resources occupied
 - reports to its user that the transaction is *aborted*
 - and then terminated (quits DBS), the cycle-life of the transaction is ended
- e.g. Fig.17.0.6 

An Example of *Roll Back* in Transaction

- Roll back (回滚、滚回)

- 撤销到目前为止事务已经对数据库做的所有修改，使数据库状态恢复到事务开始前的状态

- E.g. define table *SC* as

```
create table SC  
(s#      integer,  
  c#      integer,  
  grade integer,  
  check (grade between 0 AND 100 )  
)
```

define the *update*
transaction **update-tran**
as:

```
update SC  
set grade = grade + 5  
where ....
```

update on some tuples

restore/rollback the
tuples updated

<i>s#</i>	<i>c#</i>	<i>grade</i>
101	1	85
102	2	55
201	2	75
210	5	100
401	3	77
...
510	2	60

(a) the initial
SC instance

<i>s#</i>	<i>c#</i>	<i>grade</i>
101	1	90
102	2	60
201	2	80
210	5	105
401	3	77
...
510	2	60

(b) *SC* instance when
failure occurs during
execution of
transaction, i.e.
integrity is violated

<i>s#</i>	<i>c#</i>	<i>grade</i>
101	1	85
102	2	55
201	2	75
210	5	100
401	3	77
...
510	2	60

(c) *SC* instance after
roll back, i.e. canceling
all the updates on *SC*
having been made,
equal to the *initial*
instance

An Example of *Roll Back* in Transaction (cont.)

- When tuple 1, 2, 3 have been successfully updated and tuple 4 is updated, the values of $t4[grade]=105 > 100$, the integrity is violated
- DBMS *rollbacks* the transaction **update-tran**
 - restore the values of $t1[grade]$, $t2[grade]$, $t3[grade]$ and $t4[grade]$ to its initial values
- Another Example
 - importing data into DB in SQL Server, where existing some error data (e.g. unmatched data type) in the input files

§ 17.5 Transaction Durability (for Concurrent Executions)

- Concurrent executions of a set of transactions allows
 - high throughput and resource utilization, reduced waiting time
 - e.g. the multi-user ticket-ordering system
- Demerits for allowing transaction concurrency
 - for a set of transactions which contain operations on *shared data*, the different schedules (serial, concurrent) may result in different final DB states after all these transactions terminate
- Concurrency-control schemes and transaction schedules are needed



Concurrent Executions (cont.)

- A *schedule* on a set of concurrent transactions, arranged by DBMS concurrency-control component, specifies the chronological order in which instructions/operations of concurrent transactions are executed
- a schedule for a set of transactions must consist of all instructions of those transactions
- must preserve the order in which the instructions appear in each individual transaction.
- A transaction that successfully completes its execution will have a commit instructions as the last statement
- A transaction that fails to successfully complete its execution will have an abort instruction as the last statement



An Example

- T1: transfer \$50 funds from account-A to account-B
- T2: transfer 10% of the balance from A to B
- initial value of A : \$1000
initial value of B : \$2000
initial (A + B) : \$3000

T_1	T_2
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	<div>← A=950, B=1050</div> read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit
	Final results: A=\$855, B=\$2145 $(A + B) = \$3000$

Fig.17.2 Serial S1
i.e. T1; T2

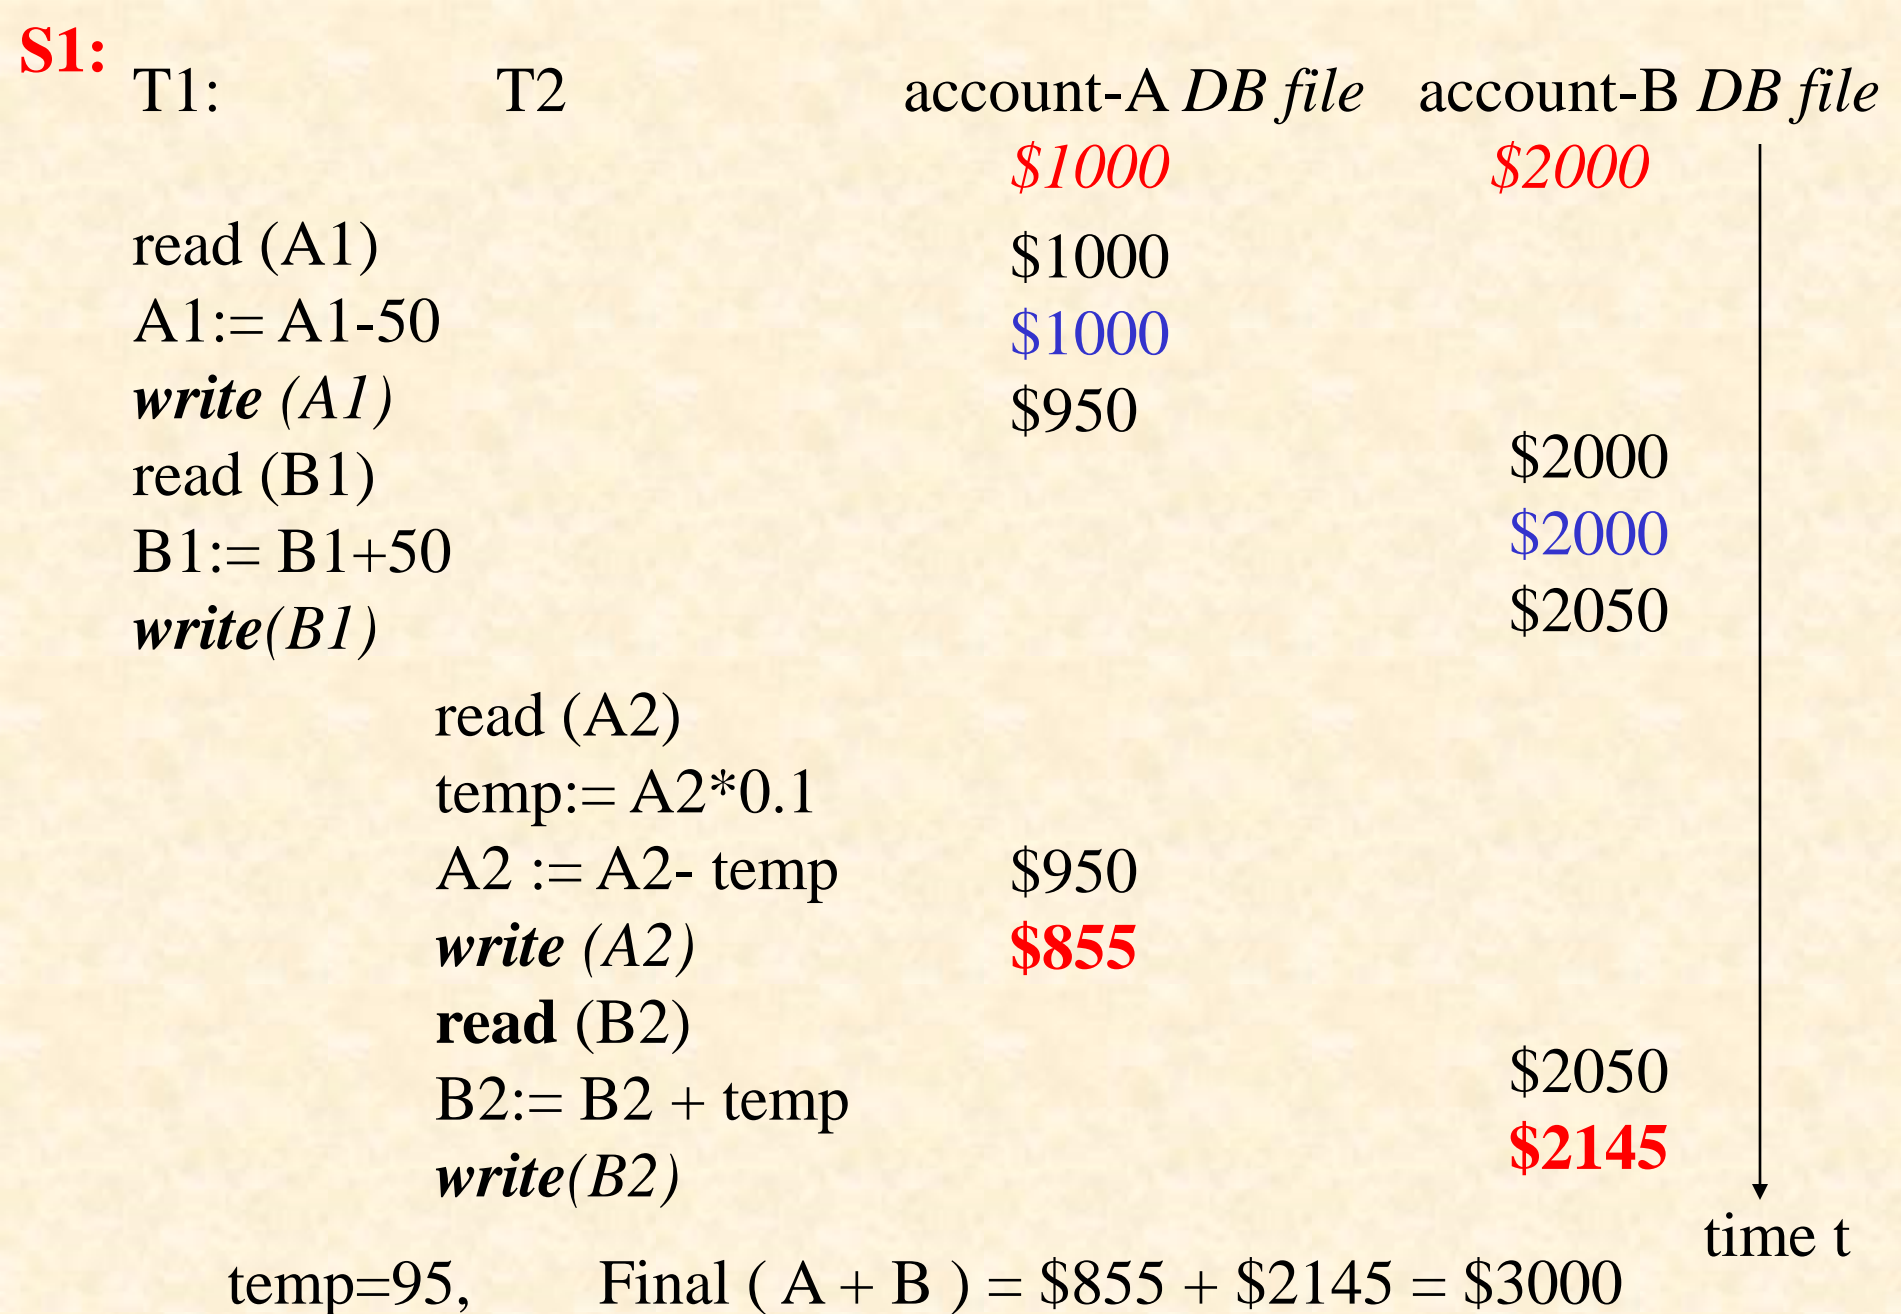


Fig.17.0.7 Analysis of serial schedule S1

An Example (cont.)

- After T_1 and T_2 are finished under the constraints of the schedule S_2 ,
 - the sum of A and B is preserved as 3000
 - the final value of *account-A* is **850**
 - the final value of *account-B* is **2150**
- Serial schedule S_1 and S_2 result in different final values of shared data items A and B ,
 - though they are all correct schedules

T_1	T_2
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit

Fig.17.3 Serial S_2
i.e. $T_2; T_1$

T_1	T_2
read (A) $A := A - 50$ write (A)	read (A) $temp := A * 0.1$ $A := A - temp$ write (A)
read (B) $B := B + 50$ write (B) commit	read (B) $B := B + temp$ write (B) commit

Fig.17.4 S3

并发调度S3将针对A的2次写操作、针对B的2次写操作安排在相近时间执行，利于DBMS在checkpoint时刻一次性地将A、B的修改结果写回数据库文件，减少DB file的写/output操作次数；

类似地，将针对A的2次读操作、针对B的2次读操作安排在相近时间执行，有利于减少对DB file 的读/input操作次数

■Final results:
 $A = \$855, B = \2145
 $(A + B) = \$3000$

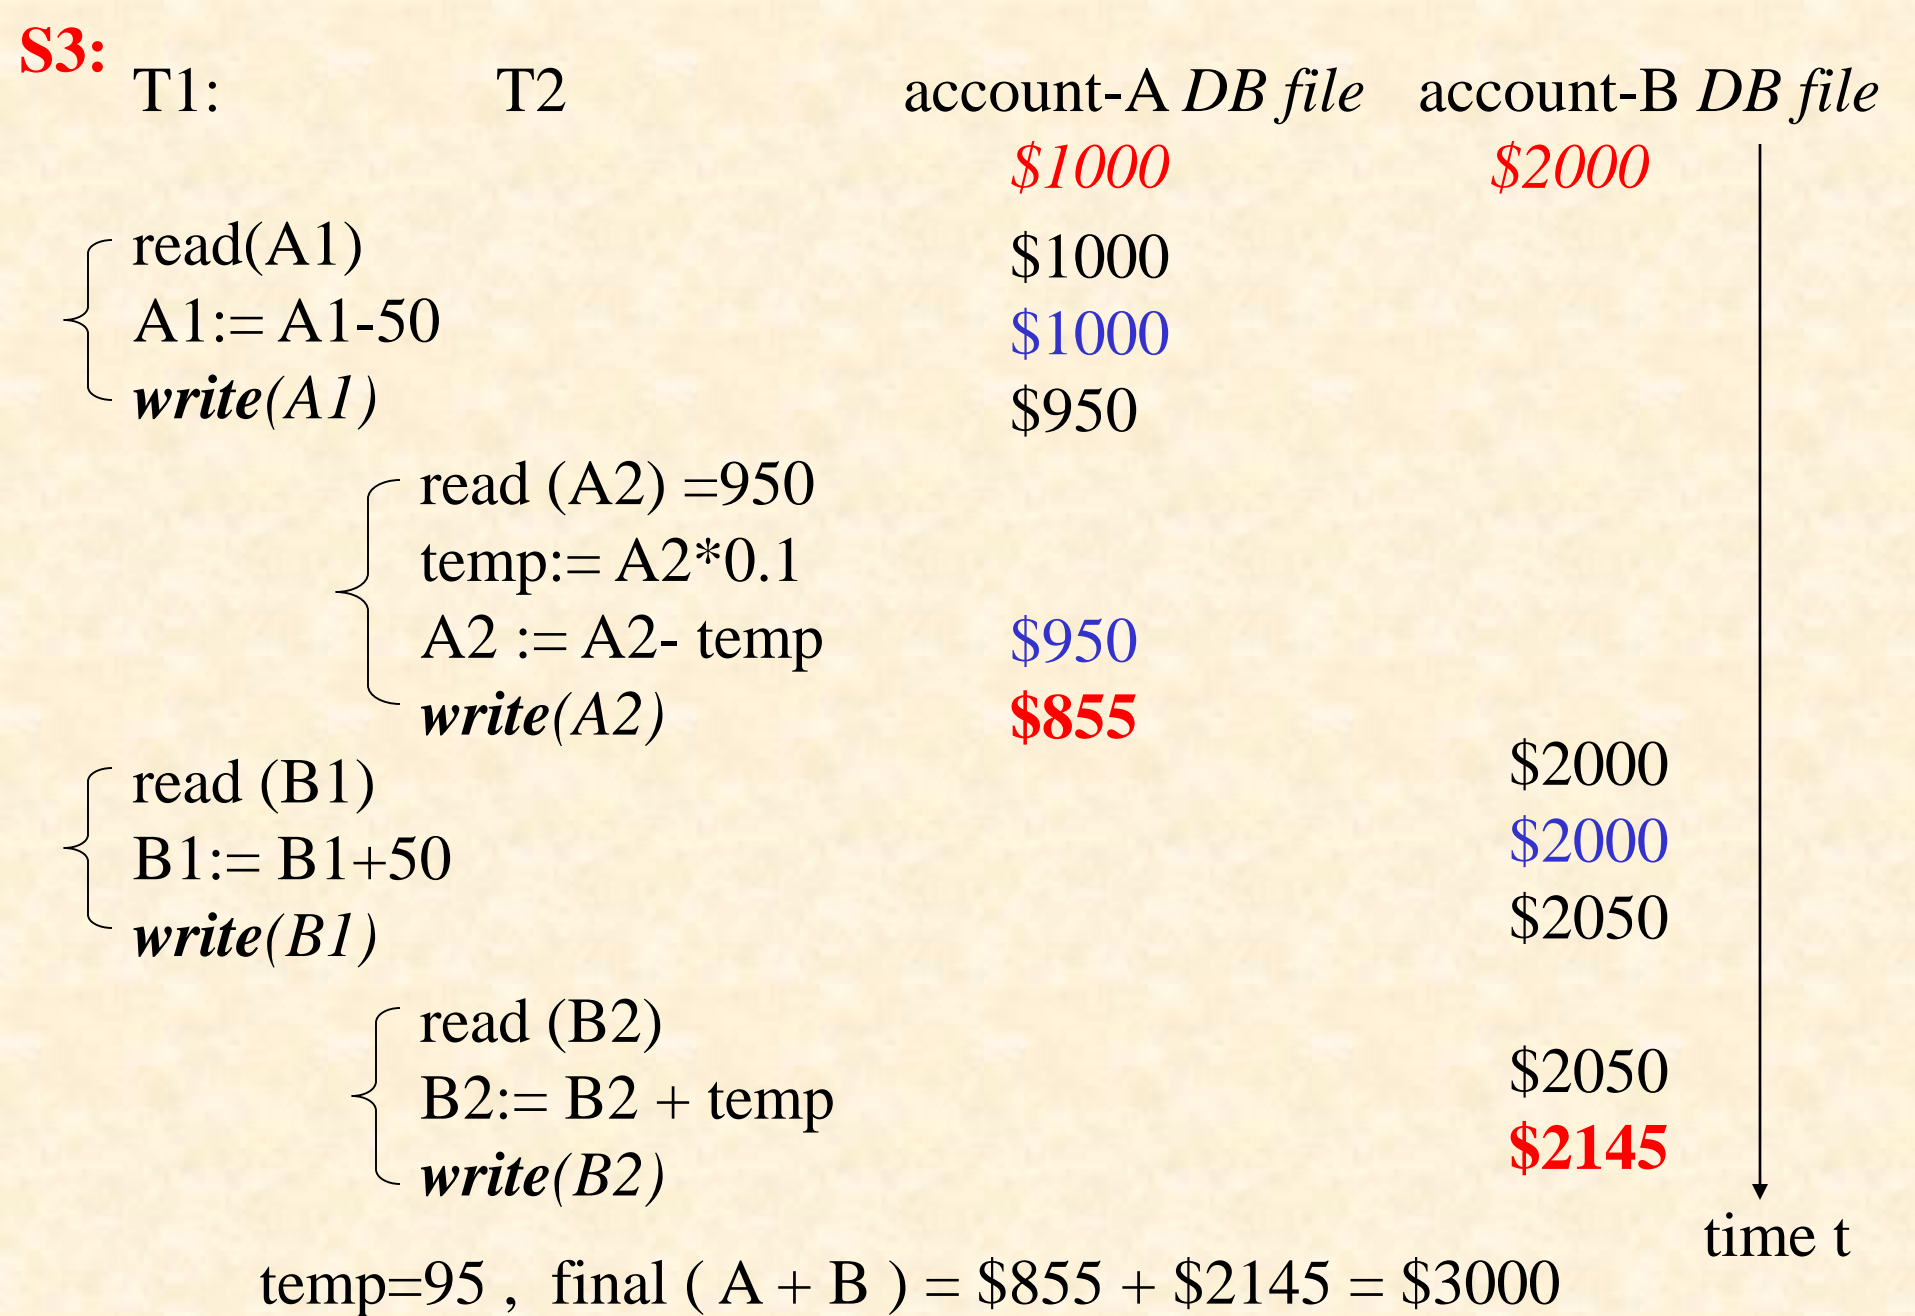


Fig.17.0.8 Analysis of concurrent schedule S3

T_1

read (A)
 $A := A - 50$

write (A)
read (B)
 $B := B + 50$
write (B)
commit

 T_2

read (A)
 $temp := A * 0.1$
 $A := A - temp$
write (A)
read (B)

$B := B + temp$
write (B)
commit

■ Note:

assuming in Fig.17.5,
read fetches data directly
from DB file on disk

Final results:

A=\$950, B=\$2100

(A + B) = **\$3050**

Fig.17.5 Concurrent S4

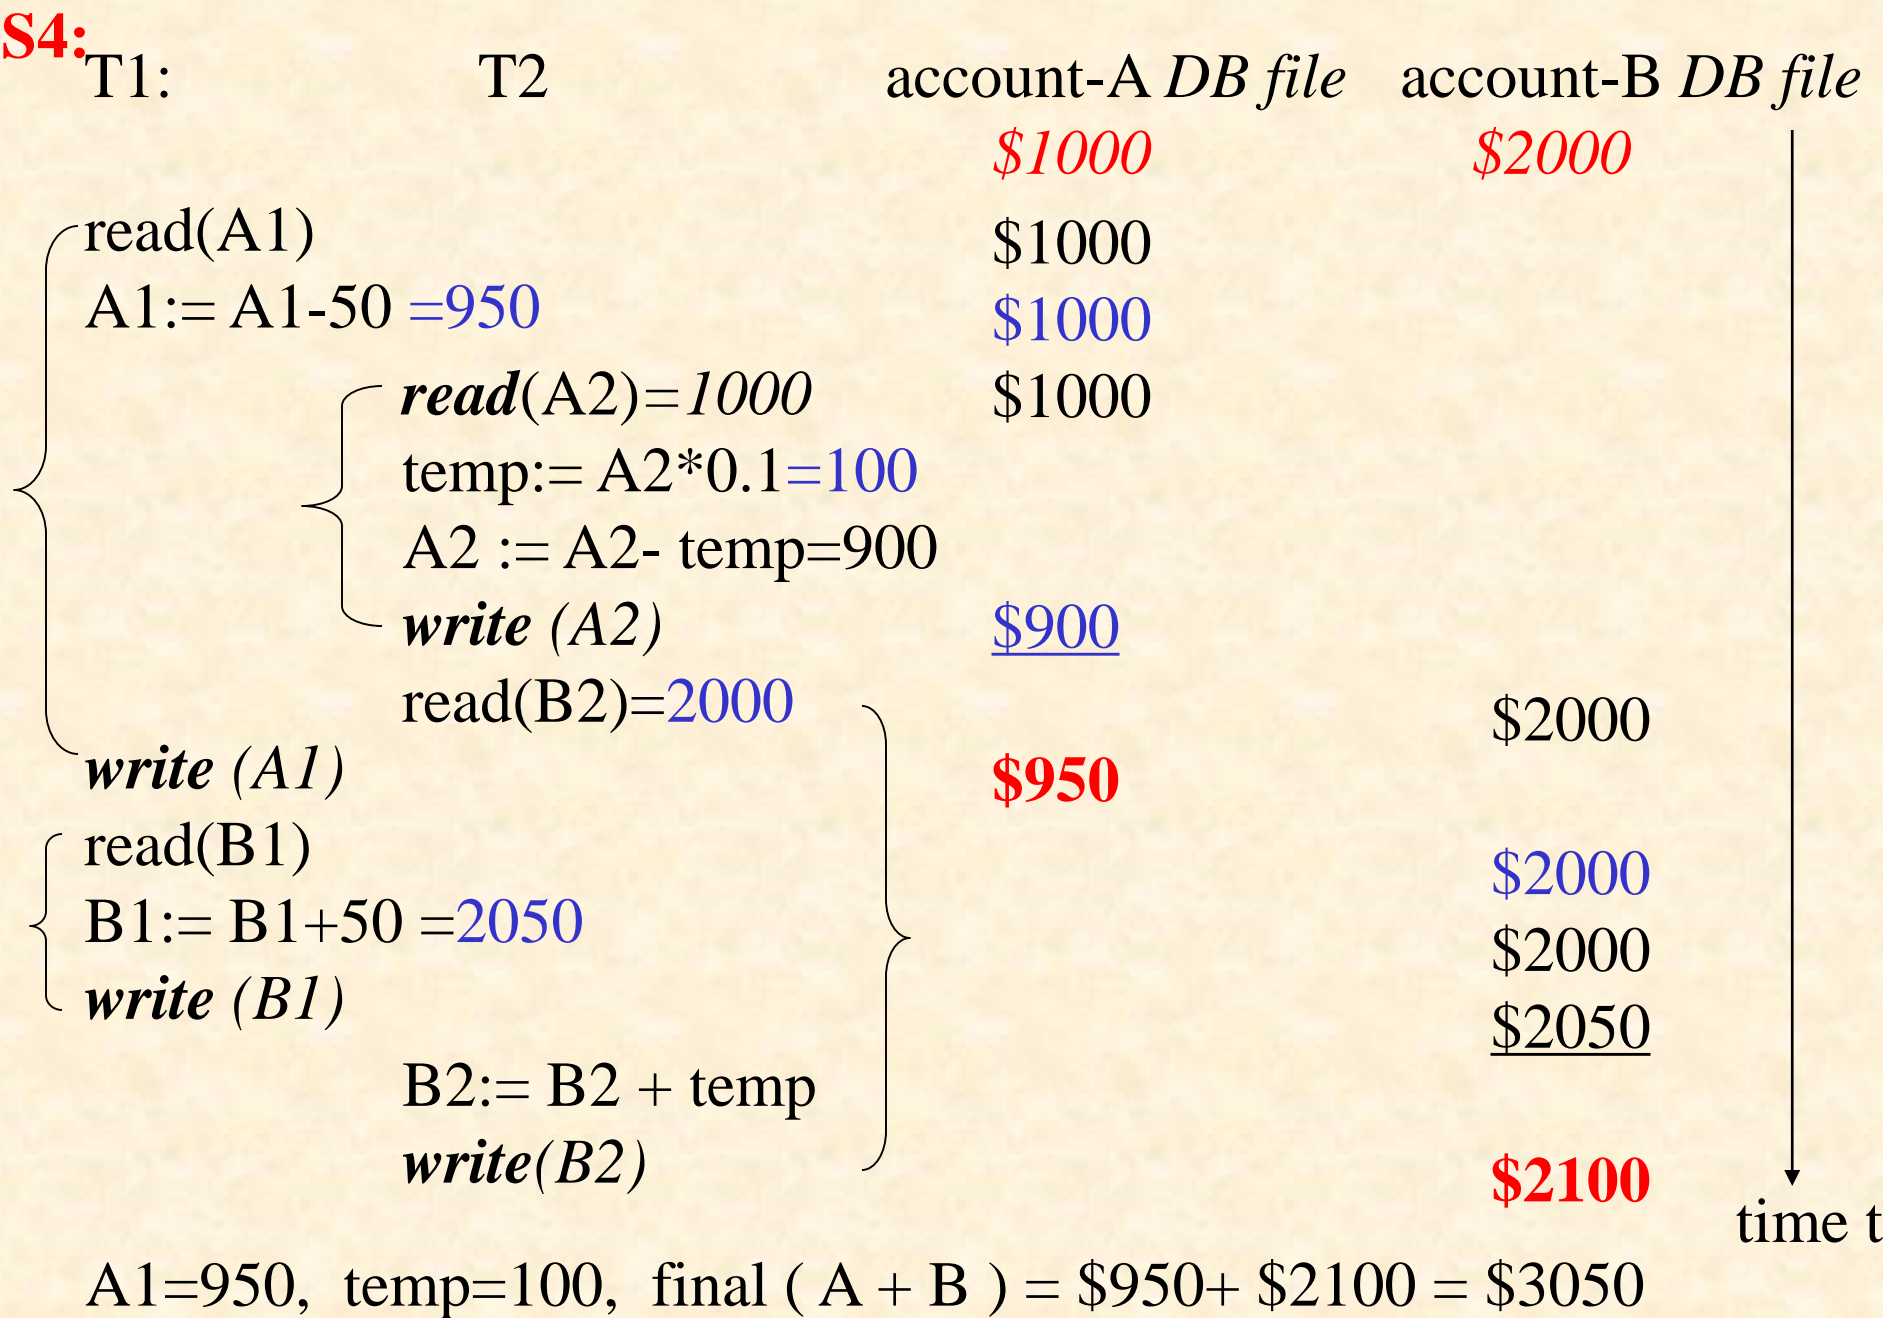




Fig.17.0.9 Analysis 1 of concurrent schedule S4

Why S3 gives a correct result while S4 leaves a wrong state ?

- In S3, T1 and T2 access shared data *A* and *B* **serially**
 - with respect to shared data A and B, S3 has the same executing results as serial *S1*, and S3 is **equivalent to serial S1**
- In S4
 - T1 and T2 operate on shared data *A* and *B* in **interweaving** (交错, 交织) ways
 - S4 has not the same executing results as serial S1, and S4 is not **equivalent to serial S1**;
 - S4 leaves the database in an **inconsistent** state, and isolation and consistency are violated

§ 17.6 Serializability

- /*数据库系统中，多个事务的**串行执行**可以保证事务的ACID特性（如S1, S2），但执行效率低
- /*多个事务**并发执行**时，如果事务操作的调度顺序不当（如S4），将影响DBS正确性
 - e.g. S4 in 
- /*事务并发控制基本原理
 - 事务调度可串行化：
相互冲突的操作串行执行，非冲突操作并行/交织执行
 - e.g. S3 in 

Serializability (cont.)

- How to identify whether or not a *concurrent schedule* S on a set of transactions is right ?
 - S is **equivalent** to a *serial* one S' ?
- /*将对 S 正确性的判断转换为对 S 可串行性的判断
- Schedule equivalence
 - conflict equivalence
 - view equivalence



Serializability

- **Basic Assumption** – Each transaction preserves database consistency.
- Thus, serial execution of a set of transactions preserves database consistency.
- A (possibly concurrent) schedule is serializable if it is equivalent to a serial schedule.
- Different forms of schedule equivalence give rise to the notions of:
 1. **conflict serializability**
 2. **view serializability**

17.5.1 Conflict Serializability

- Given two transactions $T_i = \{ I_i \}$, $T_j = \{ I_j \}$, and a schedule S on T_i and T_j
 - I_i of transactions T_i and I_j of T_j are **conflict**, if and only if
 - there exists some item Q accessed by both I_i and I_j , and at least one of these two instructions is ***write***(Q)
 - *i.e.*
 - $I_i = \mathbf{read}(Q)$, $I_j = \mathbf{write}(Q)$
 - $I_i = \mathbf{write}(Q)$, $I_j = \mathbf{read}(Q)$
 - $I_i = \mathbf{write}(Q)$, $I_j = \mathbf{write}(Q)$
- $l_i = \mathbf{read}(Q)$, $l_j = \mathbf{read}(Q)$,
 l_i and l_j don't conflict.

Conflict Serializability (cont.)

- S and S' are *conflict equivalent*
 - if the schedule S can be transformed into S' by a series of swaps of **non-conflicting** instructions
 - e.g.

T_1	T_2
read(A)	
write(A)	
	read(A)
read(B)	
	write(A)
write(B)	
	read(B)
	write(B)

Fig.17.7 S5

T_1	T_2
read (A)	
write (A)	
read (B)	
write (B)	
	read (A)
	write (A)
	read (B)
	write (B)

Fig.17.8 S6

Conflict Serializability (cont.)



■ Note:

- only swapping of two operations/instructions I_i and I_j that belong to two different T_i and T_j respectively are permitted
- swapping should not change the orders of instruction executing in each transaction
 - e.g. in Fig.17.7, swapping of *read*(B) and *write*(B) in T_1 is not allowed
- swapping should not change the orders of **conflict** instruction executing in S
 - e.g. in Fig.17.7, swapping of *write*(A) in T_1 and *read*(A) in T_2 is not allowed

Conflict Serializability (cont.)

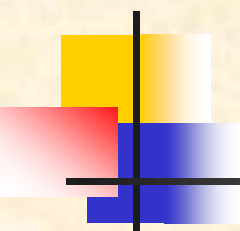
- A *concurrent* schedule S is *conflict serializable* if it is conflict equivalent to a serial schedule S'
 - e.g. concurrent $S5$ in Fig.17.7 is equivalent to $S6$ in Fig.17.8
- There are some schedules of transactions, which are not conflict serializable
 - e.g.

T_3	T_4
read (Q)	
write (Q)	write (Q)



Identification of Conflict Serializability

- Given a **concurrent** schedule S , how to determine whether or not S is conflict serializable (i.e. is right) ?
- Method 1
 - starting from the concurrent S , maintaining executing orders of conflicting operations in S , and swapping executing orders of non-conflicting operations in S
 - observing whether or not a serial S' can be obtained
- Method 2
 - precedence-graph in Fig.17.10, 17.11



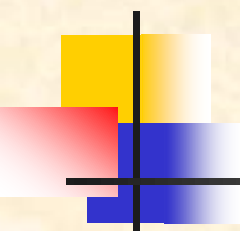
Identification of Conflict Serializability

-An Example

- Given the following concurrent schedule S on transactions T1, T2, T3, and T4
 - determine whether or not S is conflict-serializable ?
 - if S is conflict-serializable, give its equivalent serial schedule; and if it is not, give the reason

Identification of Conflict Serializability

-An Example (cont.)



Time	T1	T2	T3	T4
1			read(X)	
2		write(X)		
3			write(Y)	
4				read(Y)
5	read(Y)			
6				read(X)
7				write(Z)
8	write(X)			

Fig.14.0.10

- Answer: S is equivalent to $\langle T3; T2; T4; T1 \rangle$



Precedence-graph

- Given a schedule S of a set of transactions T_1, T_2, \dots, T_n , the precedence graph (前趋图) $G(S)$ for S is
 - a direct graph $G(V, E)$ where the vertices are the transactions
 - an arc from $T_i \rightarrow T_j$, for which one of three conditions holds
 - T_i executes write(Q) before T_j executes read(Q),
 - T_i executes read(Q) before T_j executes write(Q)
 - T_i executes write(Q) before T_j executes write(Q)

Precedence-graph (cont.)




Fig. 17.0.11 Cases for arc $T_i \rightarrow T_j$ in precedence graphs

Precedence-graph (cont.)

- A schedule is **conflict serializable** if and only if its precedence graph is **acyclic** (无环的)
- E.g. Concurrent schedule S (in the next slide) and its precedence graph







Fig.17.0.12
Precedence graph for S

T_1	T_2	T_3	T_4	T_5
read(Y) read(Z)	read(X)			read(V) read(W) read(W)
	read(Y) write(Y)	<u>write(Z)</u>		
read(U)			read(Y) write(Y) <u>read(Z)</u> write(Z)	
read(U) write(U)				


Fig.17.0.12 Concurrent schedule S

Testing for Conflict Serializability

- Cycle-detection algorithms exist which take order n^2 time, where n is the number of vertices in the graph.
 - (Better algorithms take order $n + e$ where e is the number of edges.)
- If precedence graph is acyclic, the serializability order can be obtained by a *topological sorting* of the graph.
 - That is, a linear order consistent with the partial order of the graph.
 - For example, a serializability order for the schedule (a) would be one of either (b) or (c)



Construction of Conflict Serializable Concurrent Schedule

- Given a **serial** schedule S , how to construct a concurrent S' , which is conflict equivalent to S
 - starting from schedule S , maintaining executing orders of conflicting operations in S , and swapping executing orders of non-conflicting operations in S , resulting in a conflict serializable S'
 - e.g. serial schedule S_6 in Fig.17.8 and concurrent S_5 in Fig.17.7 


§ 17.7 Recoverability

(for transaction isolation and atomicity)

- Recoverability address how to deal with the schedules when transaction failures occur
- Requirements
 - when transaction T_i fails, the operations already executed should be *rolled back/undone/aborted/canceled* to ensure the atomicity
 - when transaction T_i fails, any T_j that is dependent on T_i should also be aborted ,
 - e.g. T_j reads data having been written by T_i

17.7.1 Recoverability (cont.)

- A **recoverable schedule** is a schedule, where for each pair of T_i and T_j
 - if T_j reads a data items previously written by T_i , the **commit** operation of T_i appears before the **commit** operation of T_j .



/*解释:

T_j 依赖于 T_i ; 当 T_i 中的 $write(Q)$ 需要 aborted/rollback 时, 由于 T_j 还没有提交, T_j 可以 rollback, 因此 T_j 中的 $read(Q)$ 操作也随之 aborted

Fig.17.0.18 Definition of recoverable schedule

An example of non-recoverable schedule

- S9 in Fig.17.14 is *not* recoverable if T_9 commits immediately after the $read(A)$
 - if T_8 is aborted *after* T_9 has been committed, the value of the data item **A** will be *rolledback* to its old value before T_8 executes.
 - but T_9 has read the new value of the data item **A** and has been committed

Fig.17.14 S9

T_8	T_9
read (A) write (A)	read (A) commit
read (B)	

17.7.2 Cascadeless Schedules

■ Cascading rollback

- a single transaction failure leads to a series of transaction rollbacks
- E.g. before T_{10} , T_{11} and T_{12} are committed, just after T_{12} reads the item A, when T_{10} rollbacks due to failures, its dependent transactions T_{11} should also rollback, resulting in T_{12} also rollback

T_{10}	T_{11}	T_{12}
read (A) read (B) write (A)	read (A) write (A)	read (A)
abort		

Fig.17.15 S10

Cascadeless Schedules (cont.)

- Cascading rollback is undesirable, because it leads to the undoing of a significant amount of previous works;
- **Cascadeless** schedules is expected for concurrent transactions
- **Cascadeless** schedules are the schedules that cascading rollbacks cannot occur; formally,
 - for each pair of transactions T_i and T_j such that T_j reads a data item Q previously written by T_i , the **commit** operation of T_i appears before the **read** operation of T_j .
 - Fig.17.0.19

Cascadeless Schedules (cont.)

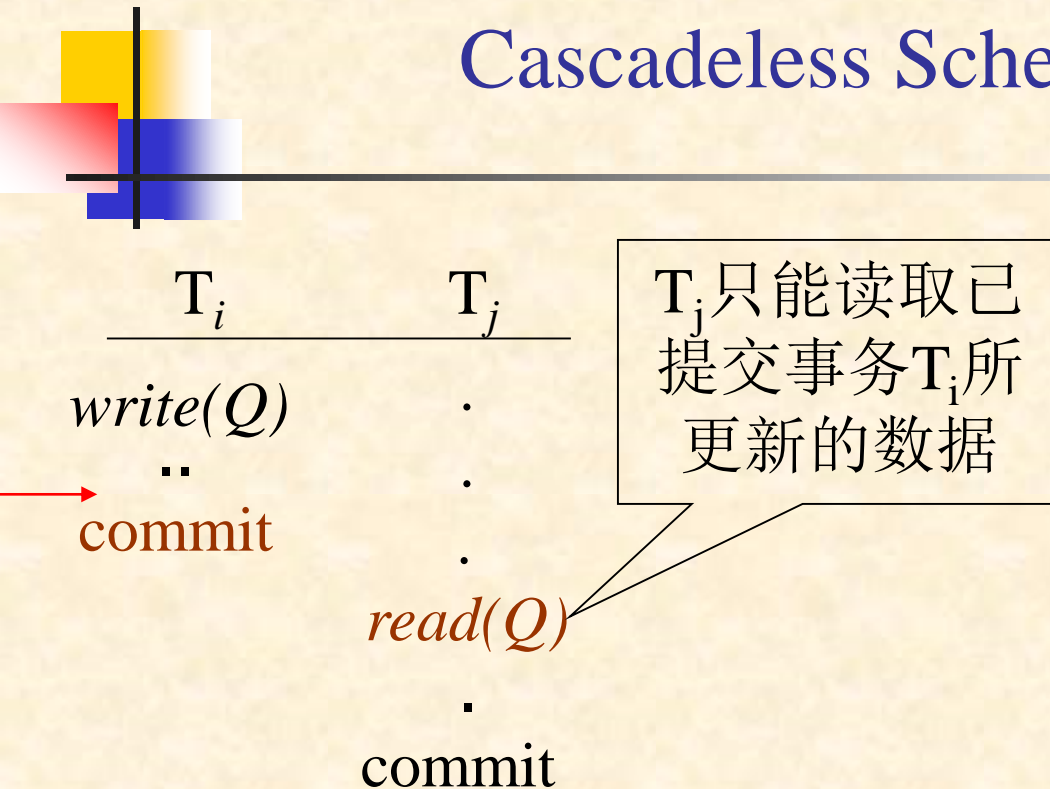


Fig.17.0.19

Definition of cascadeless schedule

/*注意与 recoverable schedule定义的区别

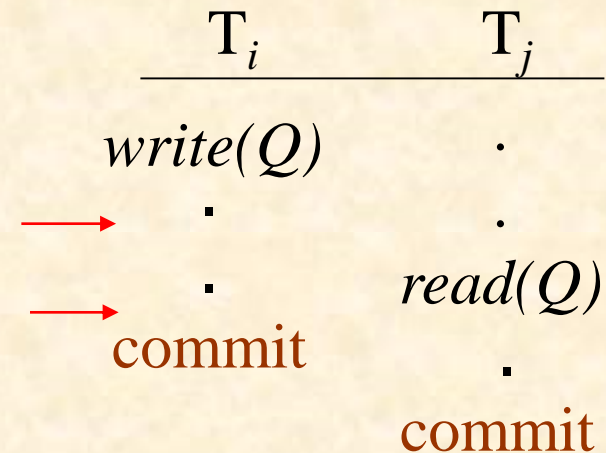
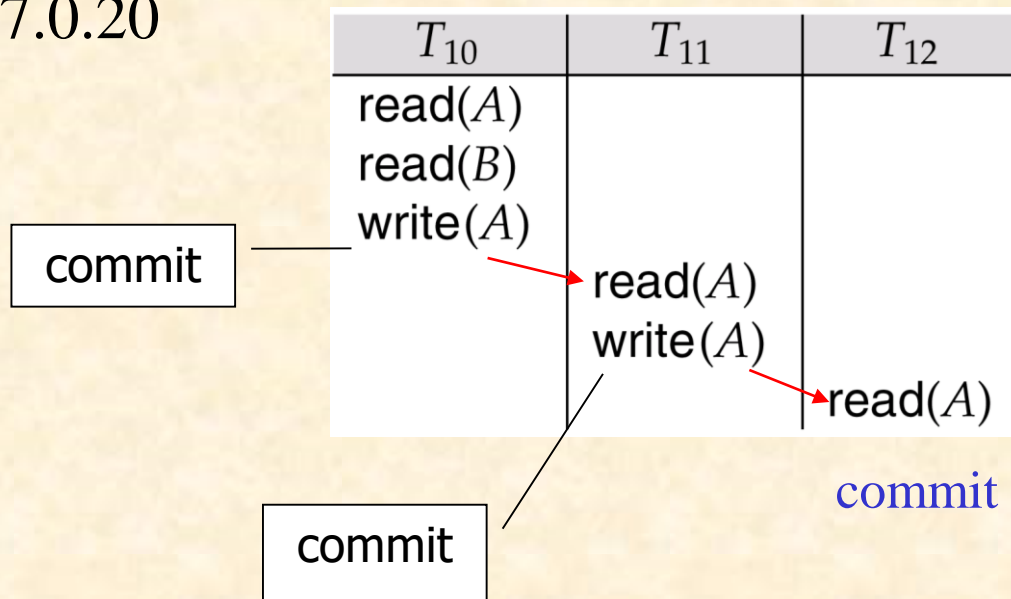


Fig.17.0.18 Definition of recoverable schedule

Cascadeless schedule vs recoverable schedule

- Every cascadeless schedule is also recoverable! !
- But a recoverable may not be a cascadeless schedule
- An example of cascadeless and recoverable schedule in Fig.17.0.20



17.8 Transaction Isolation Level (for Concurrency Control)

- A database must provide a mechanism that will ensure that all possible schedules are both:
 - Conflict serializable.
 - Recoverable and preferably cascadeless
- A policy in which only one transaction can execute at a time generates serial schedules, but provides a poor degree of concurrency
- Concurrency-control schemes tradeoff between the amount of concurrency they allow and the amount of overhead that they incur
- **Goal** – to develop concurrency control protocols that will assure serializability.



Weak Levels of Consistency

- Some applications are willing to live with weak levels of consistency, allowing schedules that are not serializable
 - E.g., a read-only transaction that wants to get an approximate total balance of all accounts
 - E.g., database statistics computed for query optimization can be approximate (why?)
 - Such transactions need not be serializable with respect to other transactions
- Tradeoff accuracy for performance



Levels of Consistency in SQL-92

- **Serializable** — default
- **Repeatable read** — only committed records to be read, repeated reads of same record must return same value.
However, a transaction may not be serializable – it may find some records inserted by a transaction but not find others.
- **Read committed** — only committed records can be read, but successive reads of record may return different (but committed) values.
- **Read uncommitted** — even uncommitted records may be read.



Levels of Consistency in SQL-92

Lower degrees of consistency useful for gathering approximate information about the database

Warning: some database systems do not ensure serializable schedules by default

E.g., Oracle and PostgreSQL by default support a level of consistency called snapshot isolation (not part of the SQL standard)



Transaction Definition in SQL

- Data manipulation language must include a construct for specifying the set of actions that comprise a transaction.
- In SQL, a transaction begins implicitly.
- A transaction in SQL ends by:
 - **Commit work** commits current transaction and begins a new one.
 - **Rollback work** causes current transaction to abort.
- In almost all database systems, by default, every SQL statement also commits implicitly if it executes successfully
 - Implicit commit can be turned off by a database directive
 - E.g. in JDBC, `connection.setAutoCommit(false);`

略 Appendix A

Transactions in SQL Server

- SQL Server系统中有三种**事务执行模式**：显示、自动提交、隐性，用于启动事务执行
- Explicit transactions
 - delimited by *begin-transaction* and *commit* or *rollback*
 - defined by user/application program, *explicitly controlled by user*
 - 工作在SQL Server系统的显示事务执行模式下
 - e.g. *account-transfer transaction* in Fig.15.0.2 

Appendix A

Transactions in SQL Server (cont.)

- Auto-commit transactions

- SQL Server 默认的事务管理模式

- 在此模式下，每条单独的*Transact-SQL*语句均是一个事务

- e.g.

CREATE, ALTER TABLE, SELECT, UPDATE,
INSERT, DELETE, DROP, FETCH, OPEN

- 每条语句完成时，可以成功提交(commit)，也可以失败滚回(rollback)

- e.g. update SC


set Grade=Grade +10

当某个学生Grade已经超过99，再执行本语句将导致事务回滚，SQL Server撤销整个更新语句的执行结果

Appendix A

Transactions in SQL Server (cont.)

- Implicit transactions /隐性事务
 - SQL Server系统利用隐性事务模式设置命令设置系统的事务执行模式
set implicit_transaction on
 - 在此模式下，每个事务无需（用begin-transaction）显示地定义事务的开始，但仍以显示的**commit** or **rollback**结束
 - 当一个事务结束后，系统自动启动下一个事务，无需begin-transaction，形成连续的事务链



set implicit_transaction on

(用户定义)

(单独T-SQL语句, 每条语句对应一个单独隐性事务)

op₁;

op₂;

...

op_m

commit;

事务1

op_{m+1};

op_{m+2};

...

op_k

rollback;

事务2

...

Create;

Drop;

....

Select;

Open;

Drop;

Insert;

Delete;

Update;

...

set implicit_transaction off



SQL Server 批处理中的事务

- 批处理

- 包含1个或多个SQL语句的组，由应用程序从客户端一次性地发送到服务器端的DBMS引擎去执行

- 事务与批之间多对多关系

- 1个事务可以包括多个批，1个批中也可以由多个事务组成

- SQL Server系统中，在大多数情况下，当批中的某些SQL语句执行发生错误时，DBMS将停止执行当前语句及其后面的语句，但发生错误之前已经执行的语句是有效的，不会被回滚(rollback)

Check 语义约束: grade between 0 and 100

begin tran

insert into SC values ('s1', 'c1', 90)

update SC

set GRADE=GRADE + 12

commit tran

select * from SC

■ -----

update 操作被回滚;

但Insert 操作未被回滚, ('s1', 'c1', 90)被成功插入;

不满足事务原子性要求!!!

批处理中的事务 (cont.)

在SQL Server 中，如果需要事务中的全部操作/语句作为1个整体，同时提交或回滚，在提交“批”之前，采用下述语句进行设置：

```
set XACT_ABORT ON
```

```
begin tran
```

```
insert into SC values ('s1', 'c1', 90)
```

```
update SC
```

```
set GRADE=GRADE + 12
```

```
commit tran
```

```
select * from SC
```

- update 操作、insert操作均被回滚，('s1', 'c1', 90)无法插入



略 Appendix B Savepoint

- **SAVEPOINT** 语句定义事务执行过程中的中间点/状态
 - 使用保存点将各组相关语句分开可以在事务内标识重要的状态
- 使用 **ROLLBACK TO SAVEPOINT** 语句可以撤消该点后的所有更改，即将事务的执行退回到以前的状态

```
BEGIN TRANSACTION
USE student-DB
INSERT INTO student
    VALUES ( “ 03402”, “王菲”, “CS”, “1985/05/15”)
SAVE TRAN My-savepoint                /* defining save-point */
DELETE FROM student
    WHERE name= “王菲” or “章立”
ROLLBACK TRAN My-savepoint
COMMIT TRAN
GO
```

Note:

/*rollback将操作滚回到保存点 *My-savepoint*: delete 操作被***rolled back***, 而 insert操作则不被***rolled back***;

DB恢复为delete操作执行前的状态, 新插入的元组 (“03402”, “王菲”, “CS”, “1985/05/15”)并未被删除, 仍然保存在数据库中*/

Fig.17.0.14 An example of savepoint in T-SQL

st-id	st-name	depart-ment	birth-date
03405	章立	CS	1985/01/25
...
03409	李龙	CS	1984/12/20
...
03411	赵新	CS	1985/06/18

(a) DB before the transaction starts

st-id	st-name	depart-ment	birth-date
03402	王菲	CS	1985/05/15
03405	章立	CS	1985/01/25
...
03409	李龙	CS	1984/12/20
...
03411	赵新	CS	1985/06/18

(b) DB after *insert* is issued

Fig.17.0.15-I DB instances at different timepoints

st-id	st-name	depart-ment	birth-date
...
03409	李龙	CS	1984/12/20
...
03411	赵新	CS	1985/06/18

(c) DB after *delete* is issued

st-id	st-name	depart-ment	birth-date
03402	王菲	CS	1985/05/15
03405	章立	CS	1985/01/25
...
03409	李龙	CS	1984/12/20
...
03411	赵新	CS	1985/06/18

(d) DB after *rollback* is issued

Fig.17.0.15-II DB instances at different timepoints

Have a break